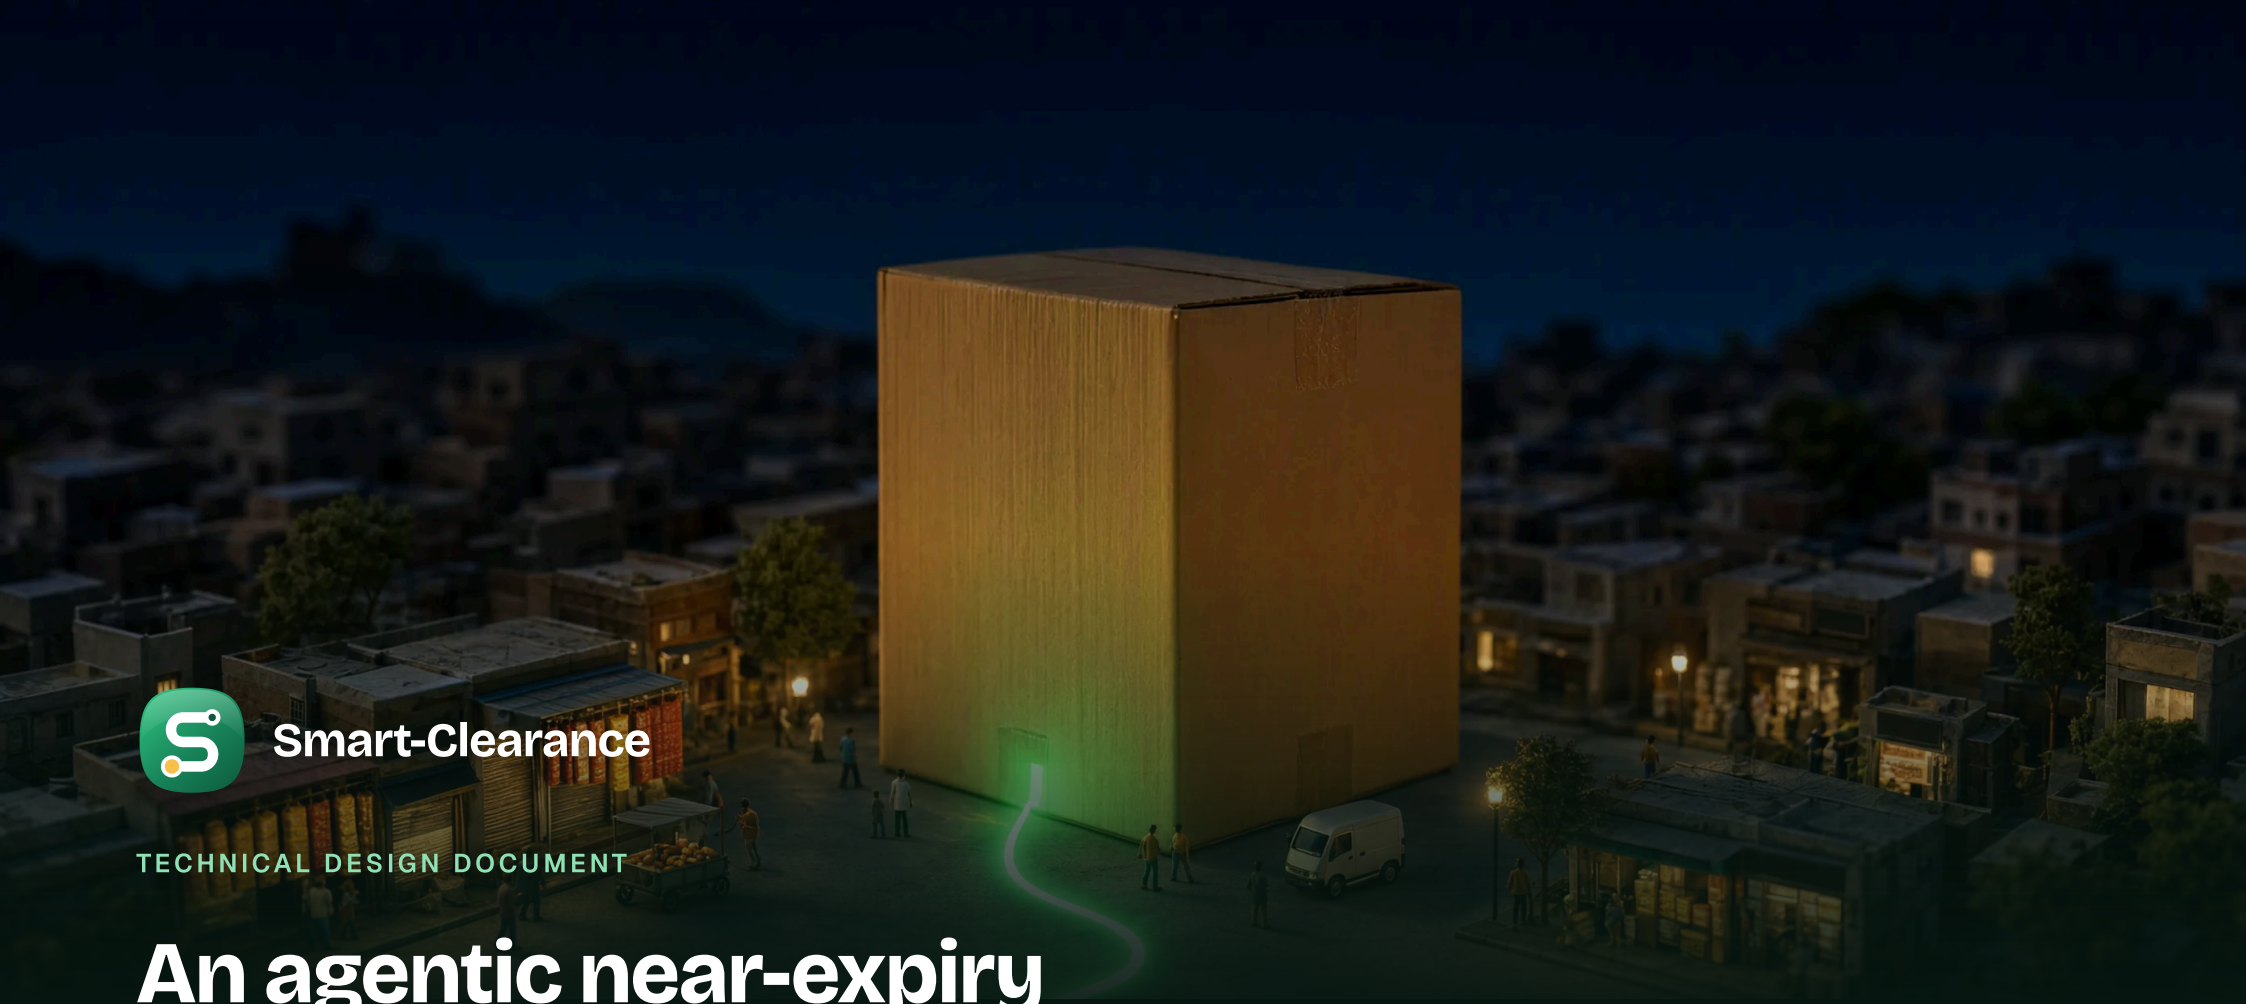


Smart-Clearance

TECHNICAL DESIGN DOCUMENT

An agentic near-expiry router, built on Google Cloud.

How Smart-Clearance is designed and built: Firebase, Cloud Run and Gemini on Agent Platform, the agents on Google's Agent Development Kit, the data, identity, delivery, testing and observability, with the measurements, decisions and known gaps.

REFERENCE	VERSION	DATE	STATUS	PREPARED FOR
SC-TDD-001	1.0	15 October 2026	Issued	Google AI Hackathon 2026 · Retail & Commerce

Document control

Document information

Document	Smart-Clearance · Technical Design Document
Reference	SC-TDD-001
Version and date	1.0, 15 October 2026
Status	Issued, for the Google AI Hackathon 2026, Retail & Commerce track
Classification	Public. The repository is public, and every company, person and figure in Smart-Clearance is fictional.
Owner	The Smart-Clearance team
Source	The architecture page (docs/smart-clearance-architecture.html) and ARCHITECTURE.md, as of 15 October 2026

Revision history

Version	Date	Change
1.0	15 October 2026	First issue: the architecture page as a technical design document, with performance and cost, the design decisions, the known gaps, what comes next, every Lucidchart diagram and a glossary.

The system it describes

Repository	github.com/aibuilder2026/smart-clearance
Google Cloud	Project aibuilder-510213, region asia-south1; Gemini on Agent Platform's global endpoint
Live apps	smartclearance.web.app · smartclearance-console.web.app · munchly-smartclearance.web.app · smartclearance-demo.web.app

Related documents

- **ARCHITECTURE.md**: the technical architecture in full, with every diagram
- **README.md**: setting up, running and operating it
- **TESTING.md**, **DEMO.md**, **PRODUCT.md**: the tests, the demo, the product
- **Lucidchart**: the diagrams' sources, linked from Appendix A

Conventions

- Green: Smart-Clearance's own code
- Blue: a managed Google Cloud service
- Teal: AI, Gemini and the agents
- Amber: a person's yes

Figures are production measurements of 8 to 10 October 2026, priced at Google's INR list prices. Diagrams are Lucidchart exports, drawn with Google Cloud's own icons.

Contents

1 **Introduction** 4
Purpose, scope, audience and the design principles

2 **The system** 6
Everything that runs, in one picture

3 **Firebase, Cloud Run and Gemini** 7
The three platform choices, and why not Firestore

4 **Google Cloud, service by service** 9
What each service does here, and why

5 **The agents** 16
Ten agents along nine stops, and one person's yes

6 **The application flow** 18
A batch's journey, request by request

7 **Identity and security** 30
Every caller signs as someone; nobody holds a key

8 **Delivery** 31
GitHub Actions gates; Cloud Build ships

9 **Testing and quality** 33
From a pixel to a live Gemini reply

10 **Observability** 36
One trace, from a tap to an agent's report

11 **Performance, cost and scale** 37
Measured on production, priced in INR

12 **Design decisions** 39
What was chosen, and why

13 **Known gaps and risks** 41
What is not done, and where it shows

14 **What comes next** 42
The roadmap, area by area

A **Appendix A · The diagrams** 43
Every Lucidchart diagram, a page each

B **Appendix B · Glossary** 67
The domain's words, and the stack's

1 INTRODUCTION

What this document is, and the ideas the design rests on

Smart-Clearance finds an FMCG brand’s stock that will not sell before its date, prices every way out, and waits for one person’s yes; then agents run the plan, draft the papers and post the impact. This document describes how it is built.

1.1 Purpose and scope

It sets out the design of the platform as built: the four web apps, backend-api, the agents, the data, the Google Cloud infrastructure, delivery, testing and observability, with the production measurements, the decisions behind them and the gaps that remain. The product and its market are in PRODUCT.md; running and operating it, in README.md.

1.2 Audience

The hackathon's judges, engineers joining the work, and reviewers of a change. It assumes familiarity with web applications and Google Cloud, and explains the domain's words in Appendix B.

One system of record

backend-api on Cloud SQL owns every write, each in one transaction with its audit line.

Code orchestrates, models write

The journey's state machine decides what runs; Gemini supplies words, never a figure.

A person's yes

The plan's approval and a destruction's second yes gate every step that cannot be undone.

Stateless agents on events

One run a Pub/Sub event, ordered by batch; each report is checked before it is kept.

Declared and keyless

Every resource in Terraform; no service-account key anywhere; CI signs in through Workload Identity Federation.

One trace, end to end

A traceparent from the browser to Gemini; spans, log lines, audit rows and agent runs share it.

Smart-Clearance

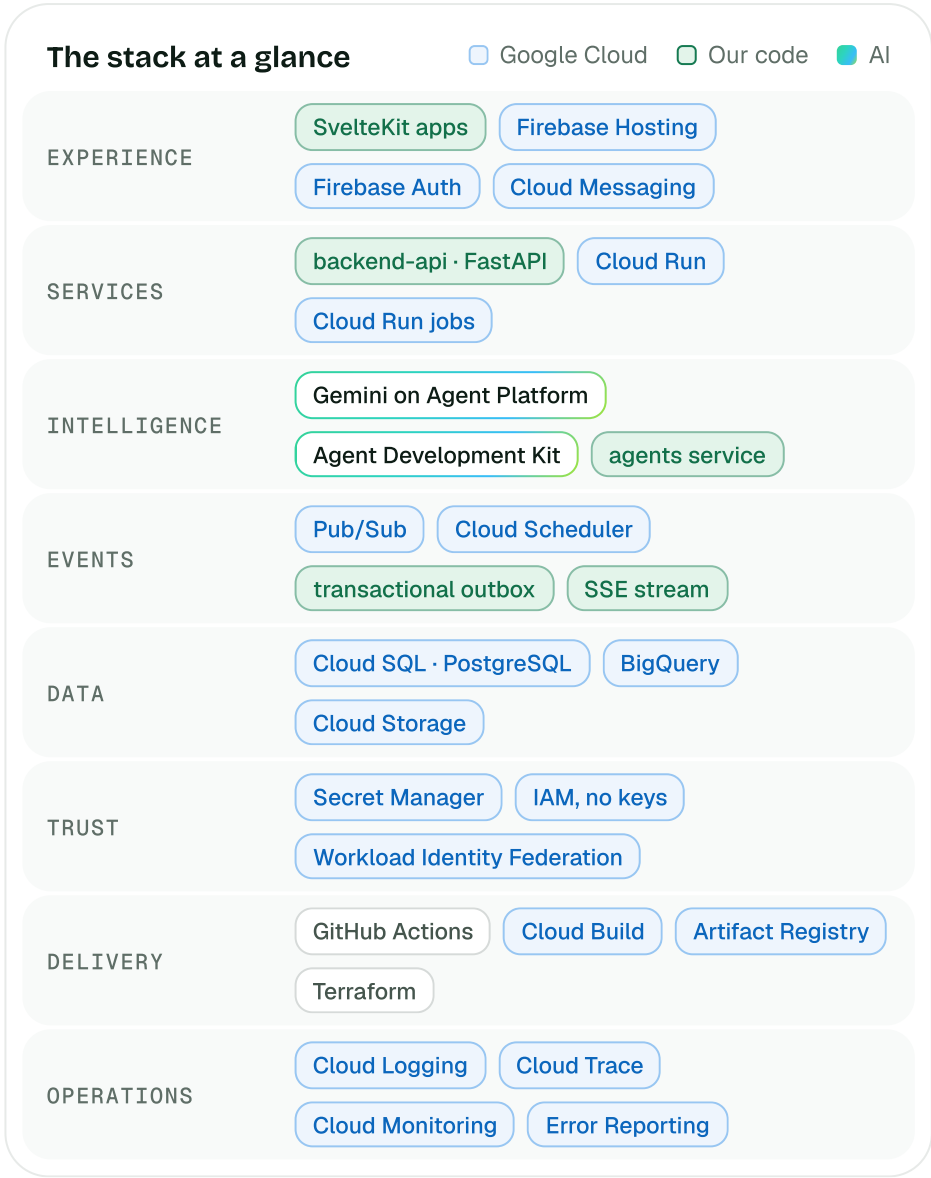
GOOGLE AI HACKATHON 2026 · RETAIL & COMMERCE · TECHNICAL ARCHITECTURE

Every near-expiry carton gets a second chance, routed by agents on Google Cloud.

Smart-Clearance finds an FMCG brand's stock that will not sell before its date, reads the carton's label, prices every way out, and waits for one person's yes. Then agents list the lot, offer it to kiranas, answer buyers, draft the GST papers and post the impact ledger. This page shows how it is built: Firebase at the edge, Cloud Run for the API and the agents, Gemini on Agent Platform for the words, Pub/Sub for the events, Cloud SQL as the system of record and BigQuery for the analytics.

- Landing page
- Workspace app
- Staff console
- Guided demo
- Source on GitHub

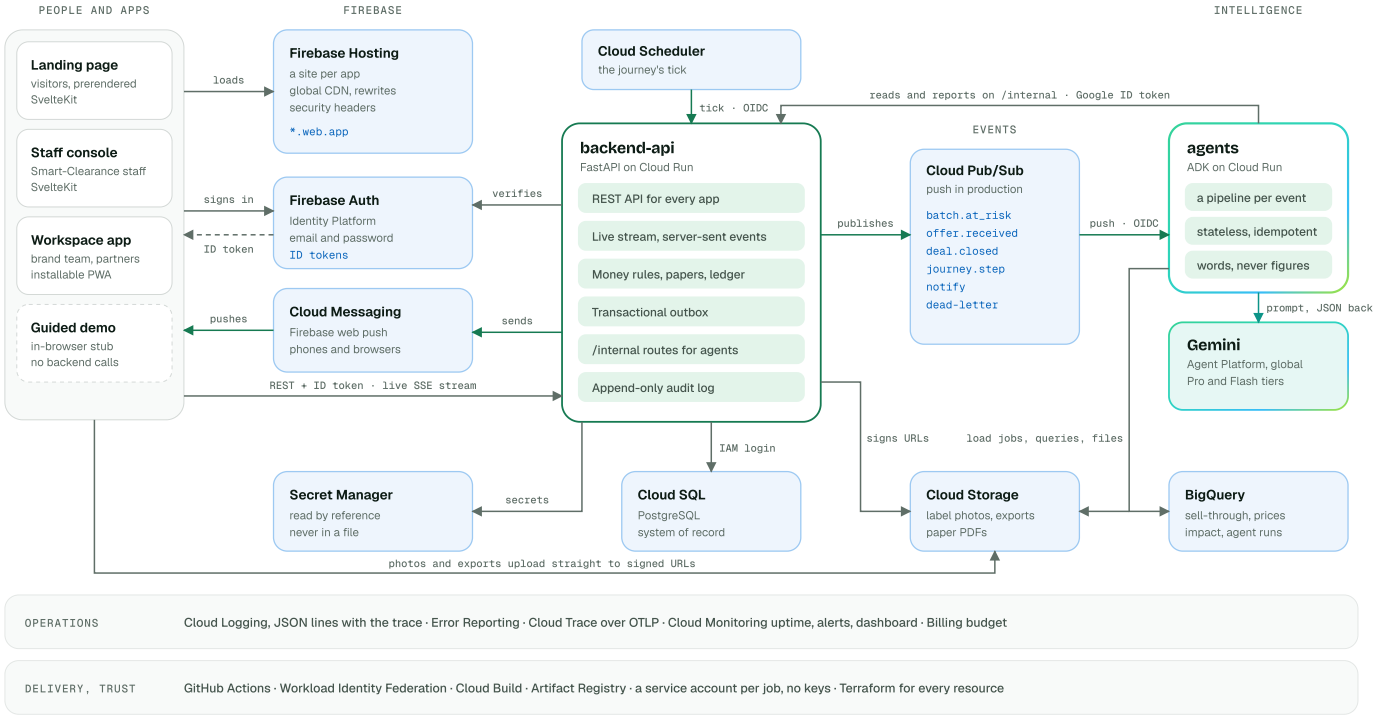
Source: github.com/aibuilder2026/smart-clearance · Live: smartclearance.web.app, munchly-smartclearance.web.app, smartclearance-console.web.app, smartclearance-demo.web.app. Every company, person and figure in Smart-Clearance is fictional; Munchly Foods is the demo client.



2 THE SYSTEM

One picture of everything that runs

Static apps on Firebase Hosting call one API on Cloud Run. The API is the only writer: every change commits in one transaction with its audit line and its outgoing events, then Pub/Sub carries the events to a stateless agents service that asks Gemini for words and reports back. Pushes reach phones through Firebase Cloud Messaging; screens move on over a live stream.



Blue boxes are managed Google Cloud and Firebase services; green outlines are Smart-Clearance's own code; the gradient outline is AI. Green arrows are events and pushes, grey arrows are requests, dashed arrows are replies.

Google Cloud Our code AI

3 WHAT THE BRIEF ASKS ABOUT

Firebase, Cloud Run and Gemini, and why Firestore is not here

Each Google product was picked for a job the product has to do. Where one would have been a poor fit, we said so and used the one that fits.

Firebase

Hosting, Authentication, Cloud Messaging, web apps

- **Hosting** serves each SvelteKit app as a static build on its own site: the prerendered landing page, the staff console, Munchly's workspace PWA and the guided demo. Deep links rewrite to the app, hashed assets cache as immutable, and the private apps send `noindex`, frame-denial and referrer headers.
- **Authentication** (Identity Platform) signs everyone in with email and password from one user pool. Sign-up is off, so only backend-api creates accounts; a password policy and email-enumeration protection are on. The API verifies each ID token with `firebase-admin`.
- **Cloud Messaging** carries the Notifier's pushes to members' phones and browsers through the PWA's service worker, so a distributor hears about a label photo request and an approver about a plan waiting.
- **Web apps and restricted keys**: each app has its Firebase web app config, written by Terraform, and a browser key limited to the identity APIs from that app's own origins.

WHY The apps are static over an API, so the whole frontend lives on Google's CDN with no server to run, and identity and push come from

Cloud Run

Two services, two jobs

- **backend-api**, FastAPI on Python, is the public API and the system's only writer. It serves the REST routes, the live stream, the agents' `/internal` routes and the tick.
- **agents**, Google ADK on Python, takes internal traffic only. Cloud Run's IAM lets in one identity, `sc-invoker`, which Pub/Sub's push subscriptions sign as.
- **Jobs** on the same image: `backend-api-migrate` runs Alembic as the schema owner, and `backend-api-hydrate` builds a synthetic world through the services. Cloud Build runs the migration before every deploy.
- Each runs as its own service account with an IAM database login, writes JSON logs that carry the trace, scales to zero, and rolls out one revision at a time behind a readiness check.

WHY Containers without servers: the same image runs on a laptop and in the cloud, service-to-service calls are signed with OIDC instead of shared secrets, and a quiet prototype costs next to nothing.

the same project with no password store or push service of our own.

Gemini on Agent Platform

Two tiers, through the Agent Development Kit

- **Gemini Pro** (`gemini-3.1-pro-preview`) writes where judgement shows: the Valuer's note on each exit, the Router's explanation of the split, the Negotiator's reply to a bid, and the evals' judge.
- **Gemini Flash** (`gemini-3.8-flash`) reads and drafts: Vision reads the printed label from a distributor's photo, the Data agent maps an unfamiliar DMS export, and the Lister, Outreach (offers in Hindi, English and Marathi), Paperwork and Impact write their copy.
- Called through `google-genai` on Agent Platform's global endpoint with the service account's own credentials: no API key exists. Every call asks for structured output against a schema.
- Gemini supplies words only. backend-api computes every rupee, the words are checked against those figures, and a template stands in for any reply that fails.

WHY A multimodal model reads a crumpled carton label as printed, writes fluent shop offers in Indian languages, and runs under the project's own IAM, quotas and audit logs.

Firestore `Not used, by design`

Cloud SQL for PostgreSQL holds the record instead

- The workspace is a ledger: money worked out by one set of rules, papers numbered in series, credit notes per line, and an audit log nobody may edit. That needs transactions across tables, constraints, triggers and SQL.
- **Live updates**, the usual reason to reach for Firestore, come from PostgreSQL itself: a change NOTIFYs as it commits, and members read it as server-sent events, cut to their role on the server.
- **Events leave only after commit**: each change writes its Pub/Sub messages to an outbox in the same transaction, so an agent never acts on a change that rolled back.
- The same PostgreSQL runs on a laptop in Docker, so every test runs against a real database, never a mock.

WHY Firebase does what it is best at here (hosting, sign-in and push), and the record lives where transactions, joins and an append-only audit log come built in.

4 EVERY GOOGLE CLOUD COMPONENT

What each service does here, and why we chose it

All of it lives in one Google Cloud project, written in Terraform. Nothing is made by hand in a console, and no service account has a key.

Edge and identity

Where people meet the product

Firebase Hosting

A site per app, released by `firebase-tools`: the landing page prerendered with its data, and the console, workspace and demo as single-page apps with deep-link rewrites and immutable caching for hashed assets.

WHY Static builds on a global CDN, each app with its own address, release and rollback.

Firebase Authentication

Identity Platform with email and password only, sign-up disabled, a password policy and email-enumeration protection. backend-api verifies ID tokens and makes accounts under a custom role limited to user management.

WHY Managed sign-in with standard ID tokens, and no password store of our own.

Firebase Cloud Messaging

The Notifier sends each notification to the member's registered devices through the workspace PWA's service worker. backend-api holds a custom role that may only send messages.

WHY One push channel for Android, installed iOS web apps and desktop browsers.

Compute

Containers, no servers

Cloud Run services

`backend-api` takes public traffic from the apps. `agents` takes internal traffic only, and only from `sc-invoker`, the identity Pub/Sub signs its pushes with.

WHY Scale to zero, a revision per release, service-to-service auth with OIDC tokens.

Cloud Run jobs

`backend-api-migrate` runs Alembic as the schema owner's IAM login; `backend-api-hydrate` builds the synthetic world through the services. Both use the API's image.

WHY One-off work with the same image and identity model, and no bastion host.

Intelligence

Gemini writes; the rules decide

Agent Platform, Gemini

Pro for the Valuer, Router, Negotiator's bids and the evals' judge; Flash for Vision's multimodal label read, the Data agent's column mapping, the Lister, Outreach, Paperwork and Impact. Global endpoint, structured output, retries with backoff on quota and server errors.

WHY Multimodal reading, Indian-language copy, and enterprise IAM and quotas in the same project.

Agent Development Kit

Each event runs one pipeline: `SequentialAgent` where a step needs the last one's result, `ParallelAgent` where it does not, `LlmAgent` with an `output_schema` for each model call, and custom steps for the deterministic work.

WHY Google's framework for Gemini agents: composable, traced, and testable on recorded responses.

Events and time

Decoupling a tap from the work

Pub/Sub

Topics per environment for the agents' events and the Notifier, ordered by batch, with dead-letter topics. Production pushes to Cloud Run with an OIDC token; a laptop pulls the local subscriptions. Each message carries its `traceparent`.

WHY At-least-once delivery with retries, ordering that keeps a batch's steps in sequence, and the same code locally with no emulator.

Cloud Scheduler

The journey's clock: `sc-tick` calls backend-api's tick every minute with an OIDC token. It runs the daily agents on the journey's own time, fires due timers, re-sends stalled events and flushes the outbox.

WHY A managed clock with the same keyless auth, so a compressed demo day and a real one run on one mechanism.

Data

Record, analytics, files

Cloud SQL for PostgreSQL 18

The system of record: clients, members, batches, cases, papers, the ledger and the audit log. IAM logins only, through the connector, over TLS; the app's login has data access only, the audit log refuses edits by trigger; backups and point-in-time recovery.

WHY A ledger needs transactions and constraints, and LISTEN/NOTIFY gives the live stream for free.

BigQuery

The agents' analytics: stock snapshots and secondary sales loaded from DMS exports, channel prices, the impact ledger, one row per agent run (model, latency, fallback, trace), and the evals' results.

WHY Sell-through across distributors and pincodes is an analytical query; load jobs never load a file twice.

Cloud Storage

Buckets for label photos, DMS exports and Paperwork's PDFs, per environment. Phones upload straight to short-lived signed URLs, and papers open through signed links. Build sources and Terraform state live here too.

WHY Files never pass through the API, and lifecycle rules tidy what is no longer needed.

Trust

No keys, no secrets in files

Secret Manager

The few secrets there are, such as the default password new accounts start on. Scripts generate them and pipe them in on stdin; services read them by reference at runtime.

WHY Nothing secret in a file, a command line, a log, Terraform state or git.

IAM and service accounts

An account per job: the API, the agents, the migrator, the invoker, the builder and two CI deployers. Developers impersonate a local account in code from their own credentials. Custom roles for user management and messaging.

WHY Least privilege, and every action in the audit logs under the identity that did it.

Workload Identity Federation

GitHub Actions trades its OIDC token for short-lived Google credentials, accepted only from this repository's ids and the protected `prod` environment, which only `main` may deploy to.

WHY CI deploys with no stored credential at all.

Delivery

From a merge to a revision

Cloud Build

Builds each image with the last one's layers as cache, pushes it, runs the migration job on it, deploys the service and checks it answers ready. CI starts it as `sc-builder`, an account that may do only that.

WHY Builds run inside Google's network, next to the registry and Cloud Run.

Artifact Registry

A private Docker repository for the API's and the agents' images, each tagged with its commit, with a cleanup policy. Cloud Build owns the images; Terraform ignores them.

WHY In-region, private, and every running revision traceable to a commit.

Operations

Seeing a request end to end

Cloud Logging and Error Reporting

Every service writes one JSON object a line with severity, source and the request's trace, so Logs Explorer groups a request's lines and Error Reporting groups stack traces.

`WHY` One place to read a request, from either service.

Cloud Trace, Telemetry API

OpenTelemetry spans over OTLP: each route, SQL statement and token check in the API, and each ADK agent, Gemini call and BigQuery job in the agents, all under the trace the browser started.

`WHY` Google's recommended route into Cloud Trace, joined to Cloud Run's own sampled requests.

Cloud Monitoring and budgets

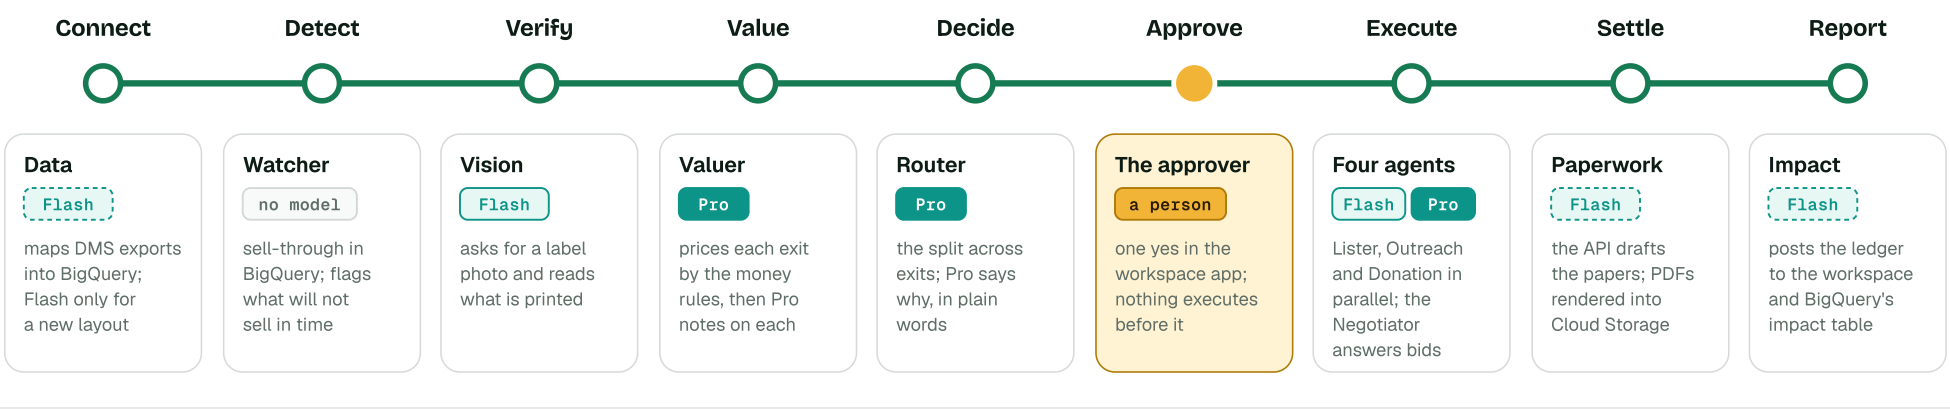
An uptime check on the API's readiness, alert policies for downtime, errors, latency, the database, dead letters, agent errors and fallbacks and push failures, log-based metrics, a dashboard and a billing budget.

`WHY` Alerts reach the operator by email before a person notices.

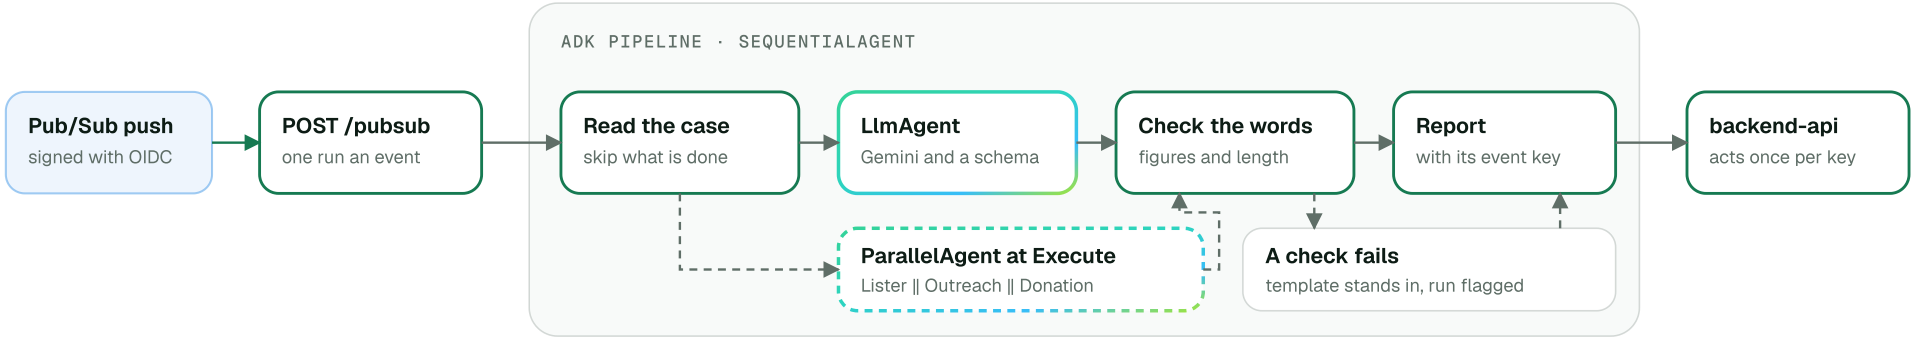
5 THE AGENTS

Ten agents along nine stops, and one person's yes

A batch at risk moves through the same stops every time. Agents do the reading, the writing and the chasing; backend-api does the arithmetic; and nothing leaves the godown until the approver says yes.



Each stop's agent and the Gemini tier it calls. A dashed tier is called only when needed (an export layout the Data agent has not seen, an optional cover note); the amber stop is a person.



Inside one run. A redelivered event reads the case first and resumes where the last run stopped; backend-api answers a repeated event key with a no-op, so nothing happens twice.

Figures are computed, never generated

backend-api works out every price, split, credit note and ledger line by one set of money rules. The agents read previews of those figures and write words around them.

The reserve never reaches a prompt

The lowest price a buyer may get is cut from every case and preview an agent reads, and tests try hostile buyer messages to prove it.

A buyer's words are data

Messages go into the Negotiator's prompt between delimiters they cannot forge, as quoted text, never as instructions.

Every word is checked

Each figure in the copy must be one of the computed ones, a reply to a bid states the decided price and nothing else, and offers keep their shop placeholder and push length.

Structured output Gemini can fill

Every schema field is required, with null meaning absent, and no schema holds a free-keyed map, after the first live eval run showed how those come back.

Bounded calls

Each model call has a deadline, retries quota and server errors with backoff, and counts against a per-run budget; past it, the template stands in.

6 APPLICATION FLOW

A batch's journey, request by request

Pick a stage to see which component calls which, in order. The core loop is the same at every stop: a person or a clock acts, backend-api commits, Pub/Sub carries the event, an agent works and reports, and every member's screen moves on.

The pages that follow take each stage in turn: the core loop and signing in first, then a batch's journey from Connect to Report.

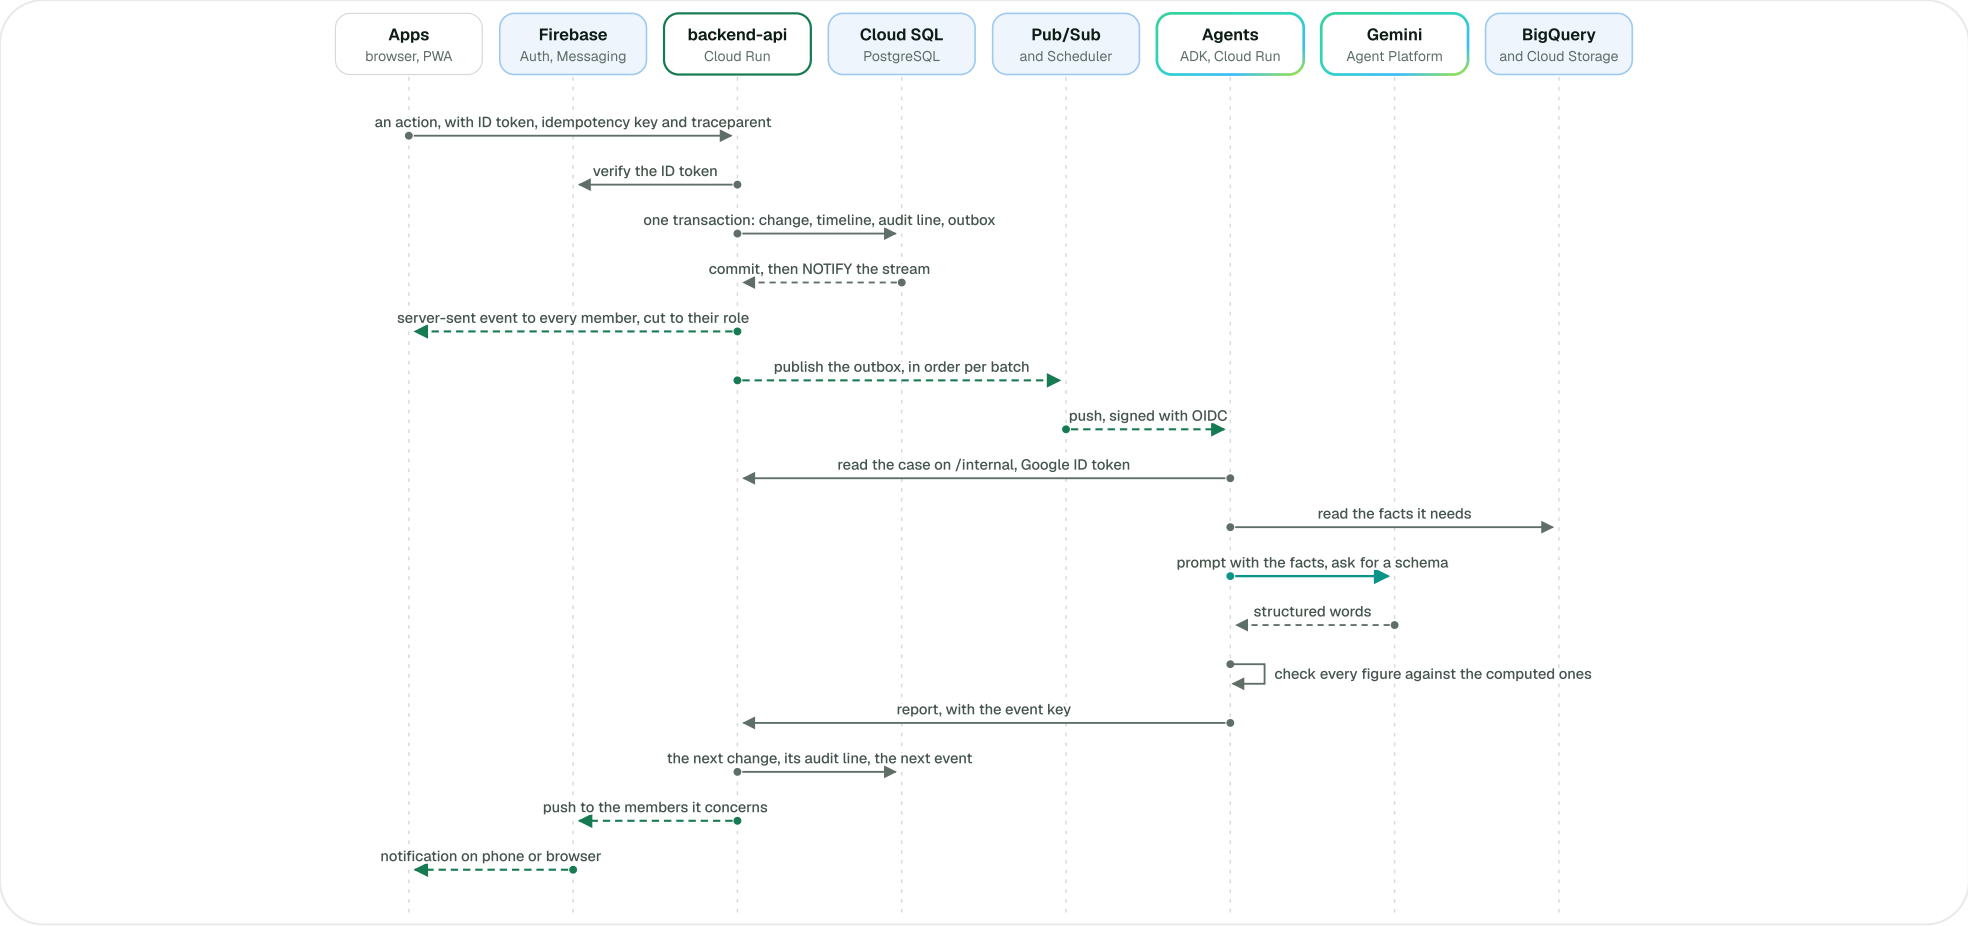
1. The core loop
2. Signing in
3. Connect · Data agent
4. Detect · Watcher
5. Verify · Vision
6. Value · Valuer
7. Decide · Router
8. Approve · a person
9. Execute · Lister, Outreach, Donation, Negotiator
10. Settle · Paperwork
11. Report · Impact

Solid grey: a request **Dashed grey**: a reply **Dashed green**: an event, push or stream **Teal**: a Gemini call **Amber**: a person's yes

The core loop

Every stop of the journey runs on this loop. A person acts in an app or the clock ticks; backend-api checks who it is, commits the change with its audit line and its outgoing events in one transaction, streams the change to every member, and publishes the events once the commit has landed. An agent picks the event up, gathers facts, asks Gemini for words, and reports back; the report is itself a change, so the loop turns again.

Firebase Authentication · Cloud Run · Cloud SQL · Pub/Sub · Agent Platform · BigQuery · Cloud Messaging

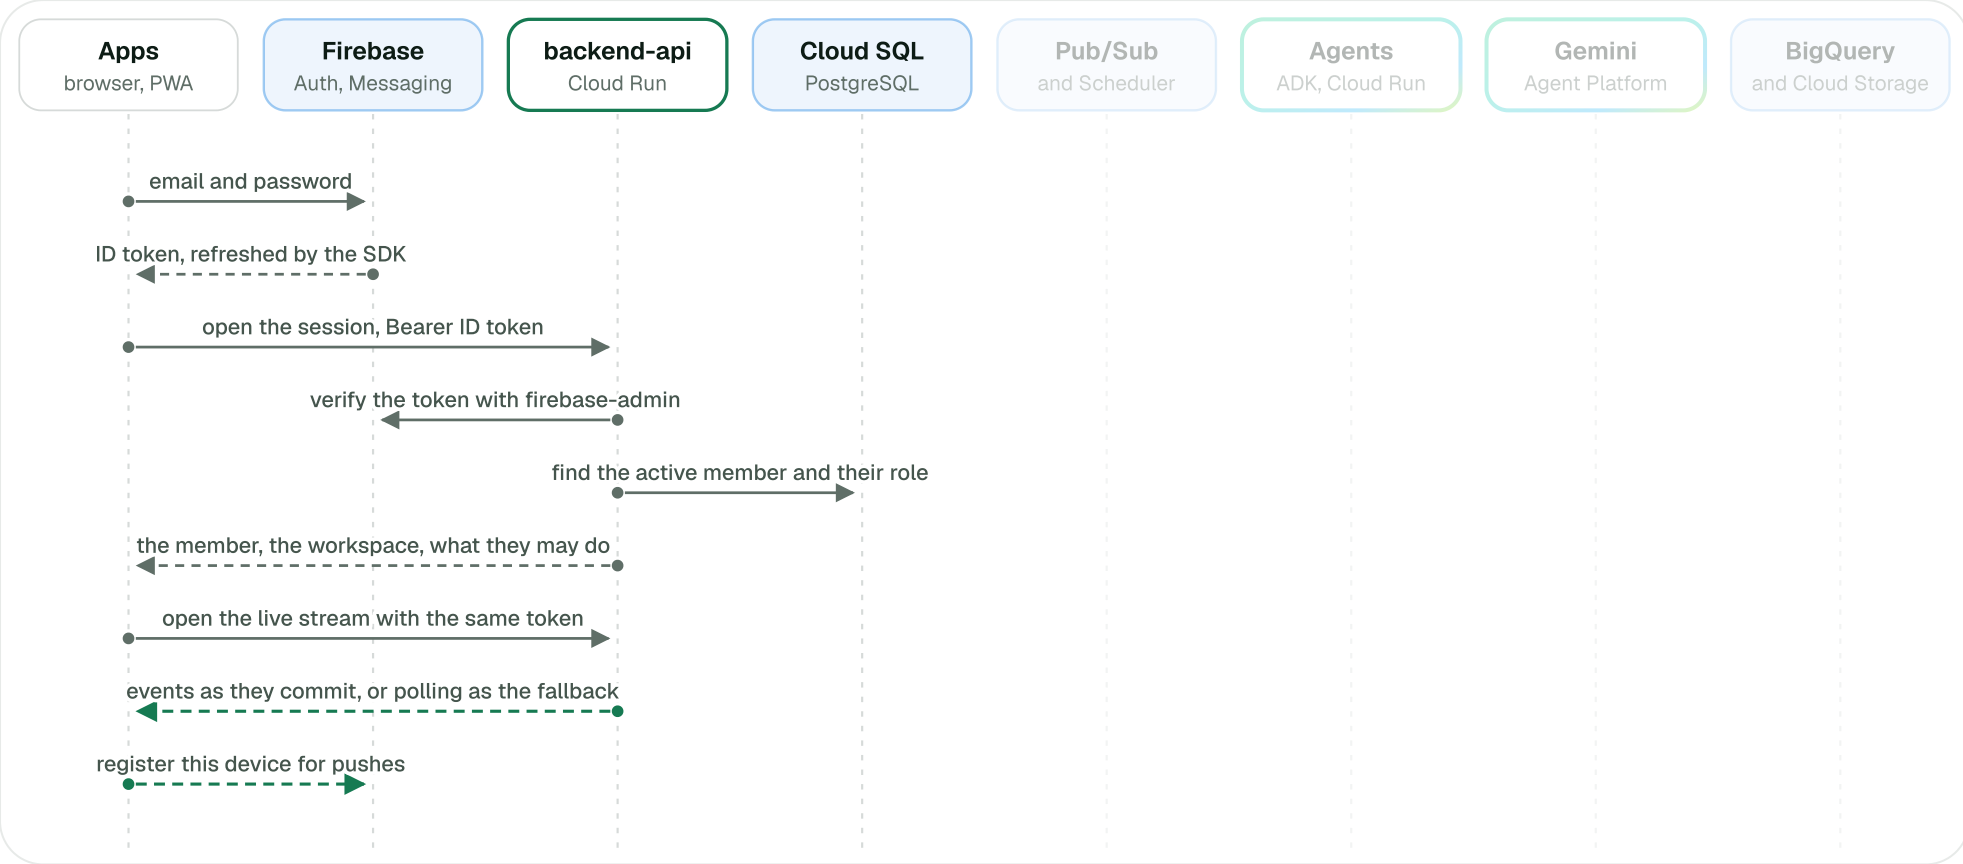


Each report names the event it answers. A redelivered event finds its work done, and backend-api answers the repeat with a no-op.

Signing in

Everyone signs in with email and password through Firebase Authentication. The browser holds the ID token, the Firebase SDK refreshes it, and every call to backend-api carries it. The API verifies the token, finds the active member and their role, and from then on every answer is cut to what that role may see.

Firebase Authentication · Identity Platform · Cloud Run · Cloud SQL

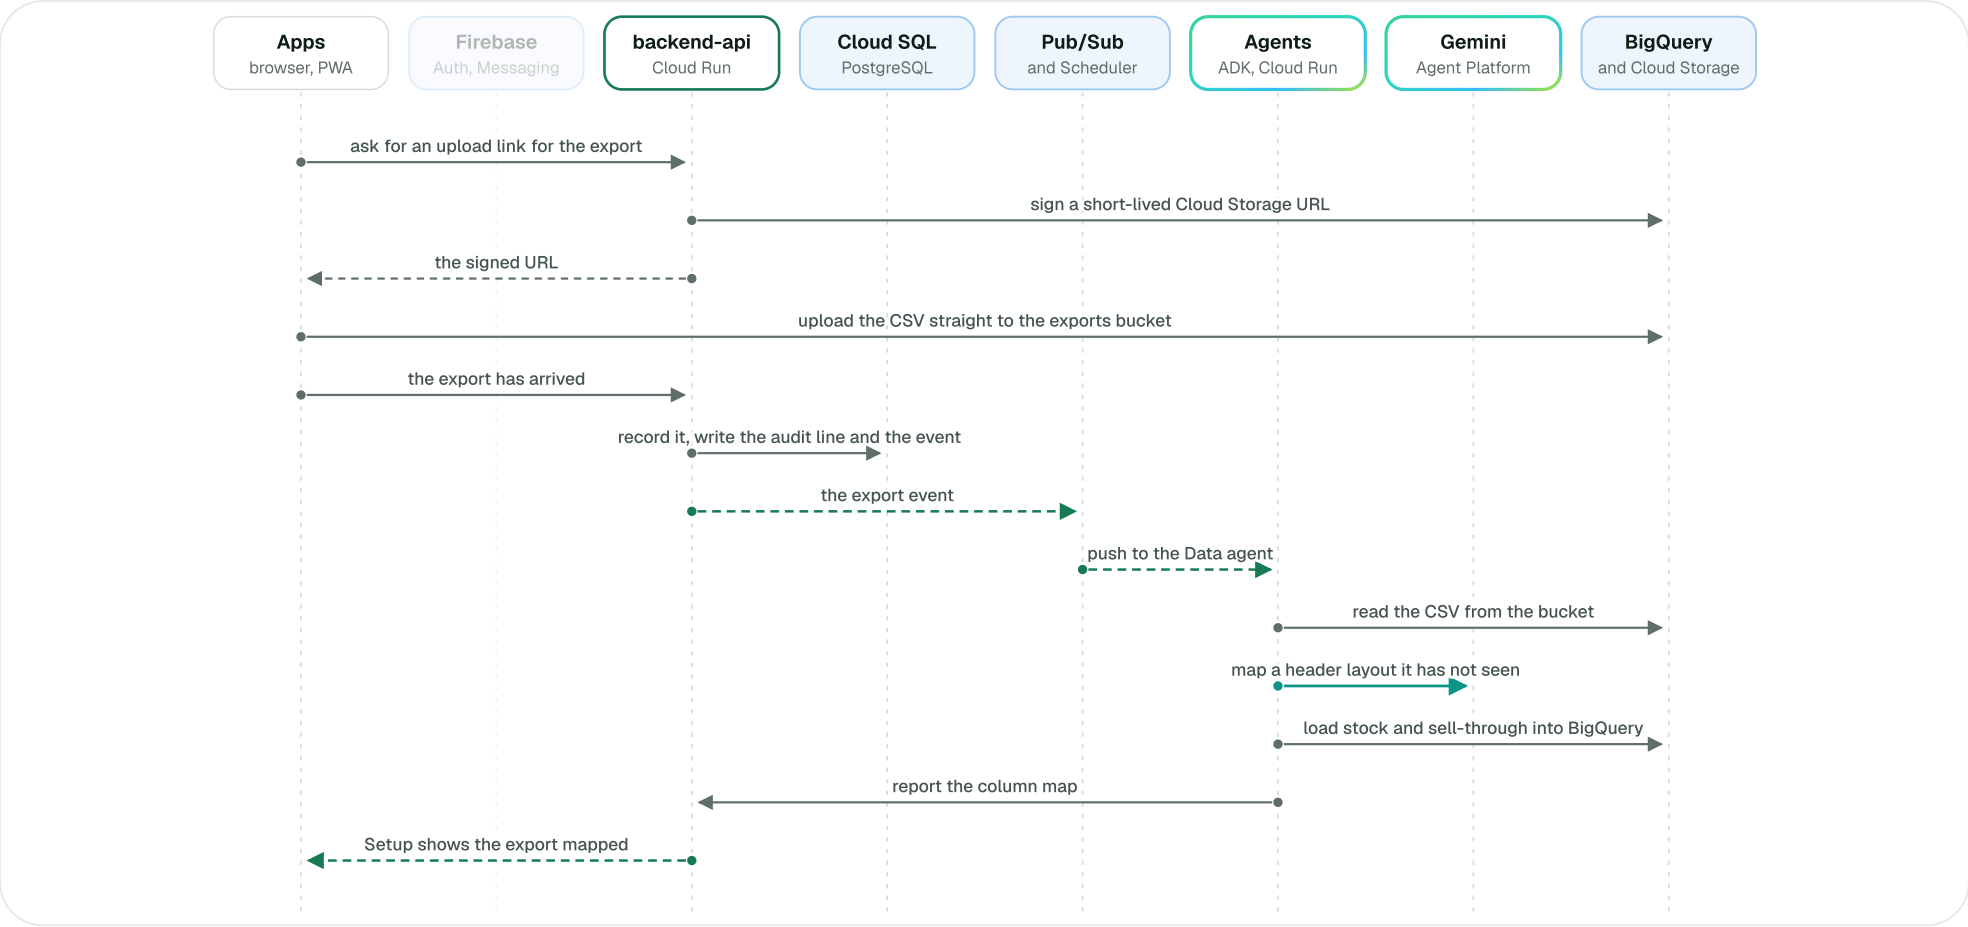


Sign-up is disabled, so only backend-api creates accounts. Email enumeration protection gives every wrong sign-in the same answer. A Firebase account that is not a member of this workspace is signed out again.

Connect · Data agent

The brand’s distributor management system exports stock and secondary sales as CSV. The file goes straight from the browser to Cloud Storage on a signed URL, never through the API. The Data agent reads it, maps its columns (with a saved layout, or Gemini Flash for headers it has not seen), and loads it into BigQuery with a load job.

Cloud Storage | Pub/Sub | Cloud Run | Agent Platform | BigQuery

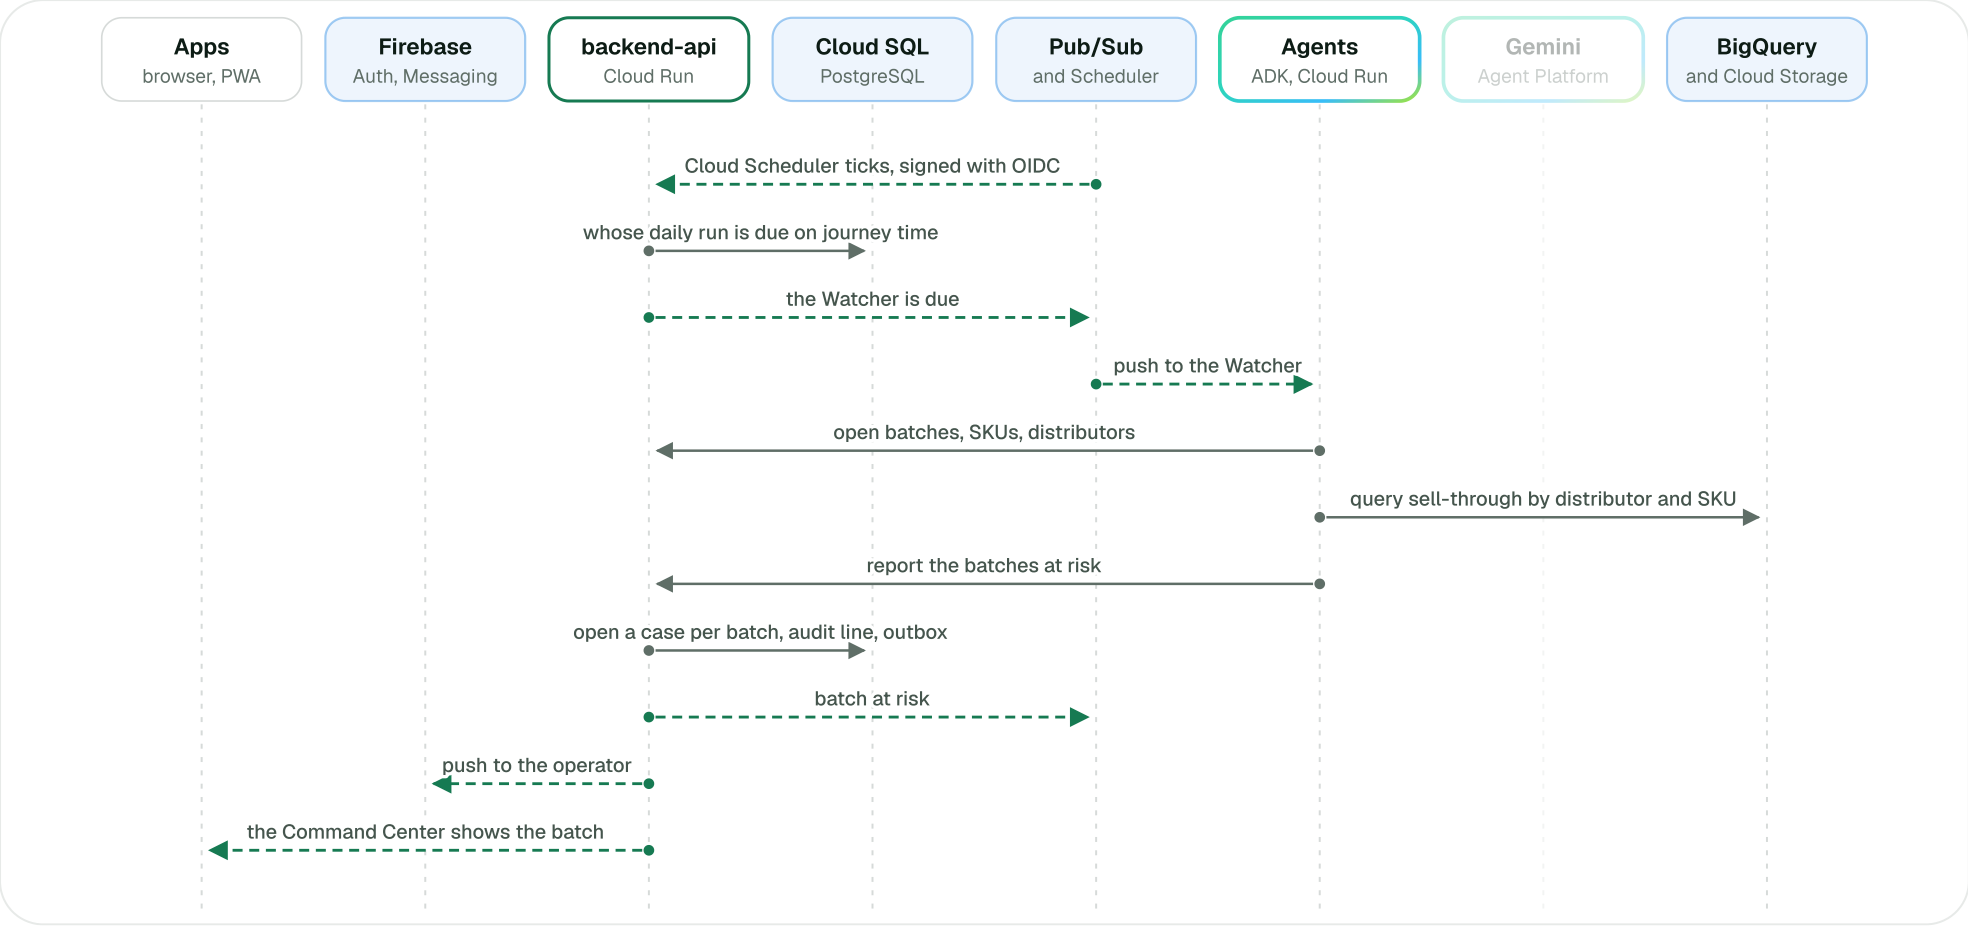


A load job never loads the same file twice. Staff can map a client’s first export from the console, so the workspace opens already connected.

Detect · Watcher

Cloud Scheduler ticks backend-api, which runs each client's daily jobs on the journey's own clock. The Watcher reads each open batch's sell-through from BigQuery and reports which batches will not sell before retailers stop taking them. backend-api opens a case for each and tells the operator.

Cloud Scheduler | Pub/Sub | Cloud Run | BigQuery | Cloud SQL | Cloud Messaging

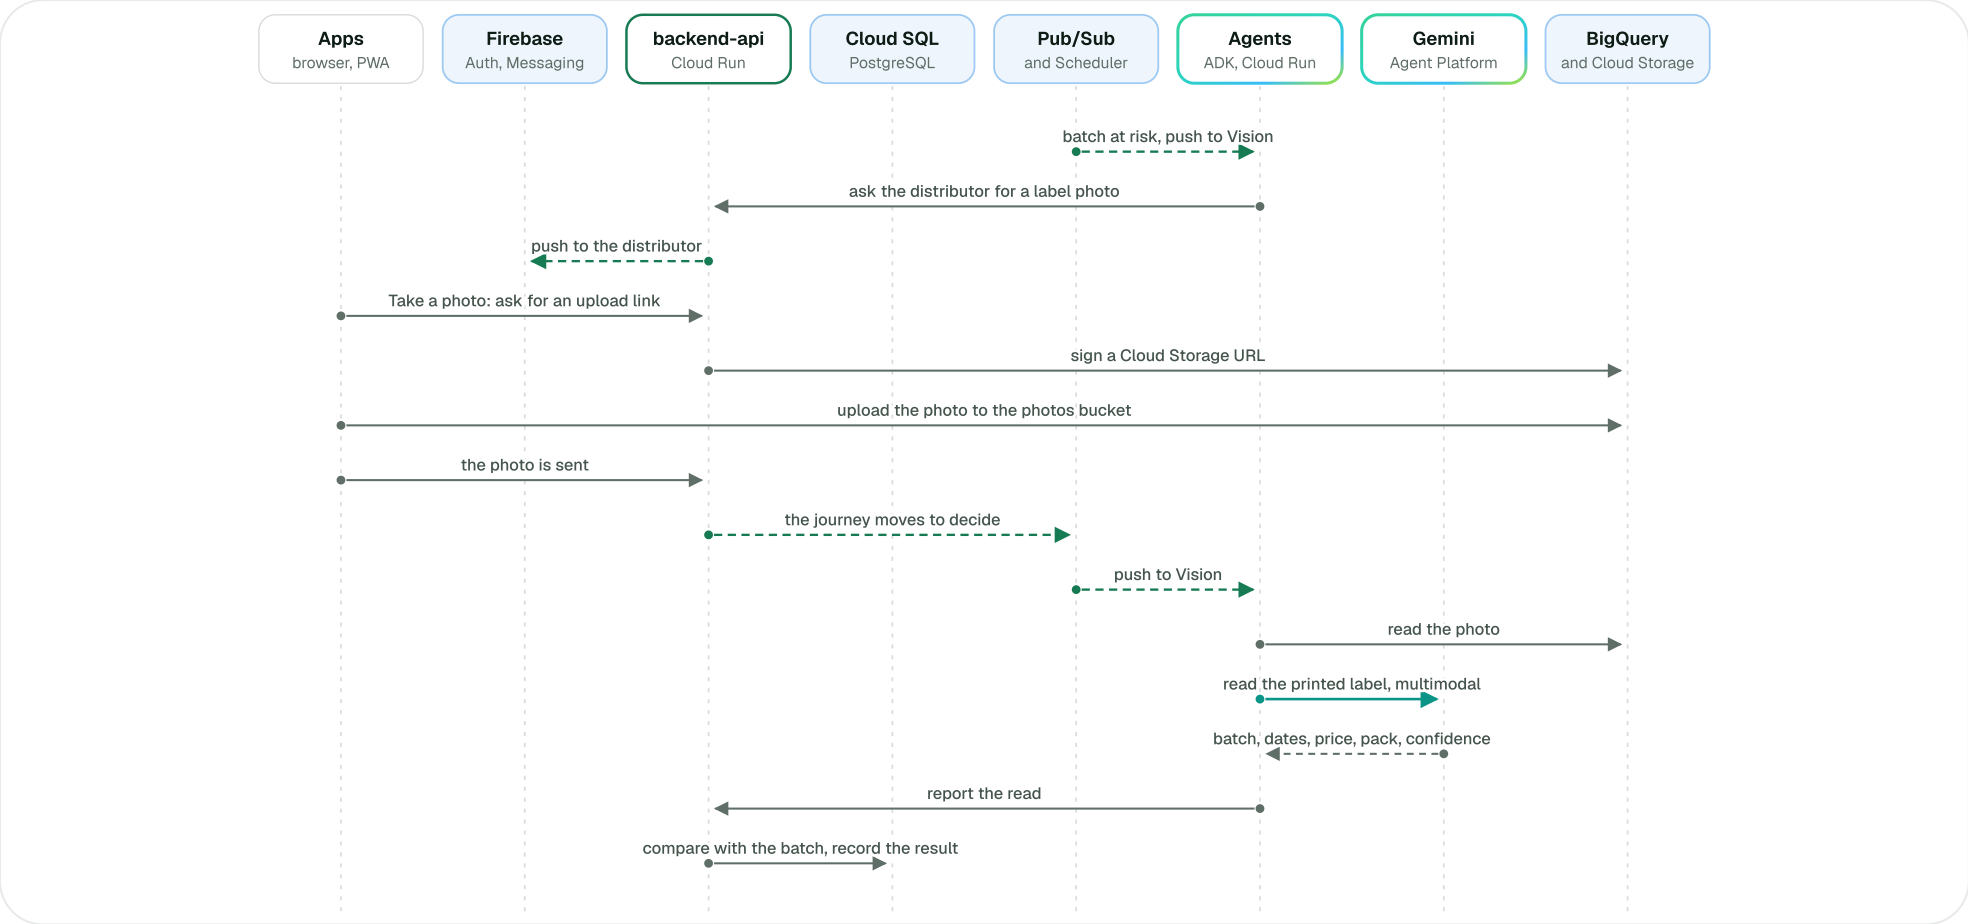


The Watcher needs no model: sell-through is a query. It flags a batch only once Setup is confirmed and the distributor has given its permission.

Verify · Vision

Before anything is priced, the batch is checked on the shelf. Vision asks the distributor for one photo of the carton's label; the phone uploads it straight to Cloud Storage; Gemini Flash reads what is printed, without being shown the record it is checked against. backend-api compares the read with the batch.

Cloud Messaging | Cloud Storage | Pub/Sub | Agent Platform | Cloud Run | Cloud SQL

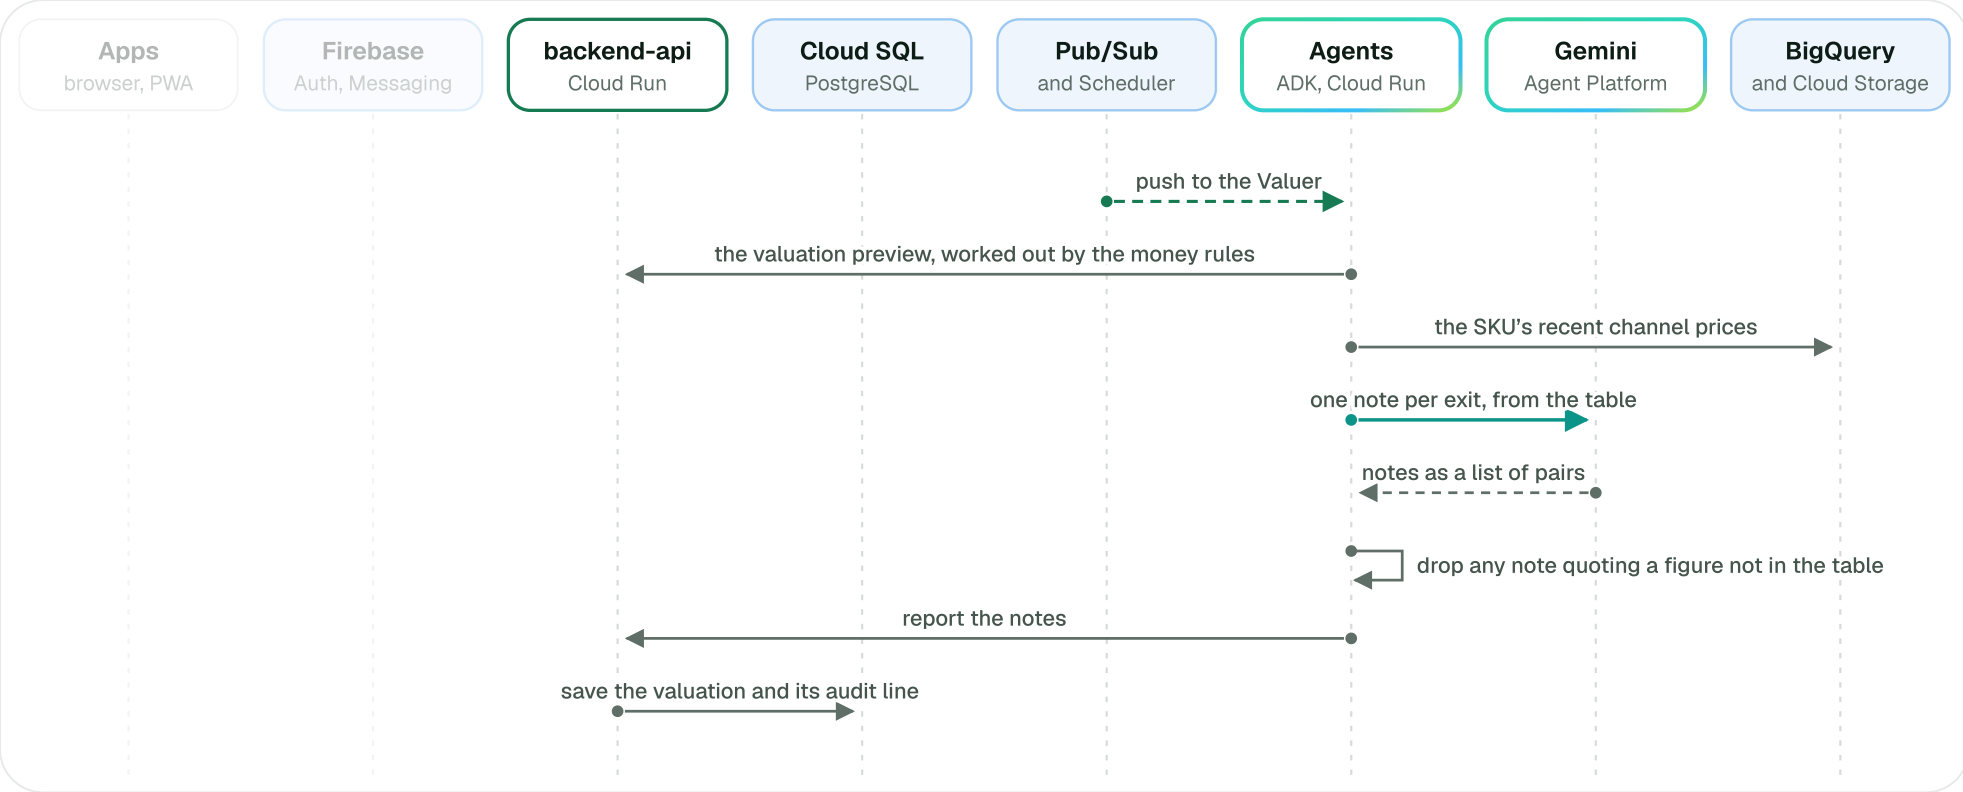


A label that cannot be read, or reads differently from the record, asks for another photo with the reason. A read Vision could not make is retried by Pub/Sub before it gives up.

Value · Valuer

backend-api prices every way out for the batch with the money rules: kiranas, the quick-commerce marketplace, a staff sale, a food bank, or the bin. The Valuer reads that table and the SKU's recent prices in BigQuery, and Gemini Pro writes one note per exit saying why it suits this batch now.

Pub/Sub | Cloud Run | BigQuery | Agent Platform | Cloud SQL

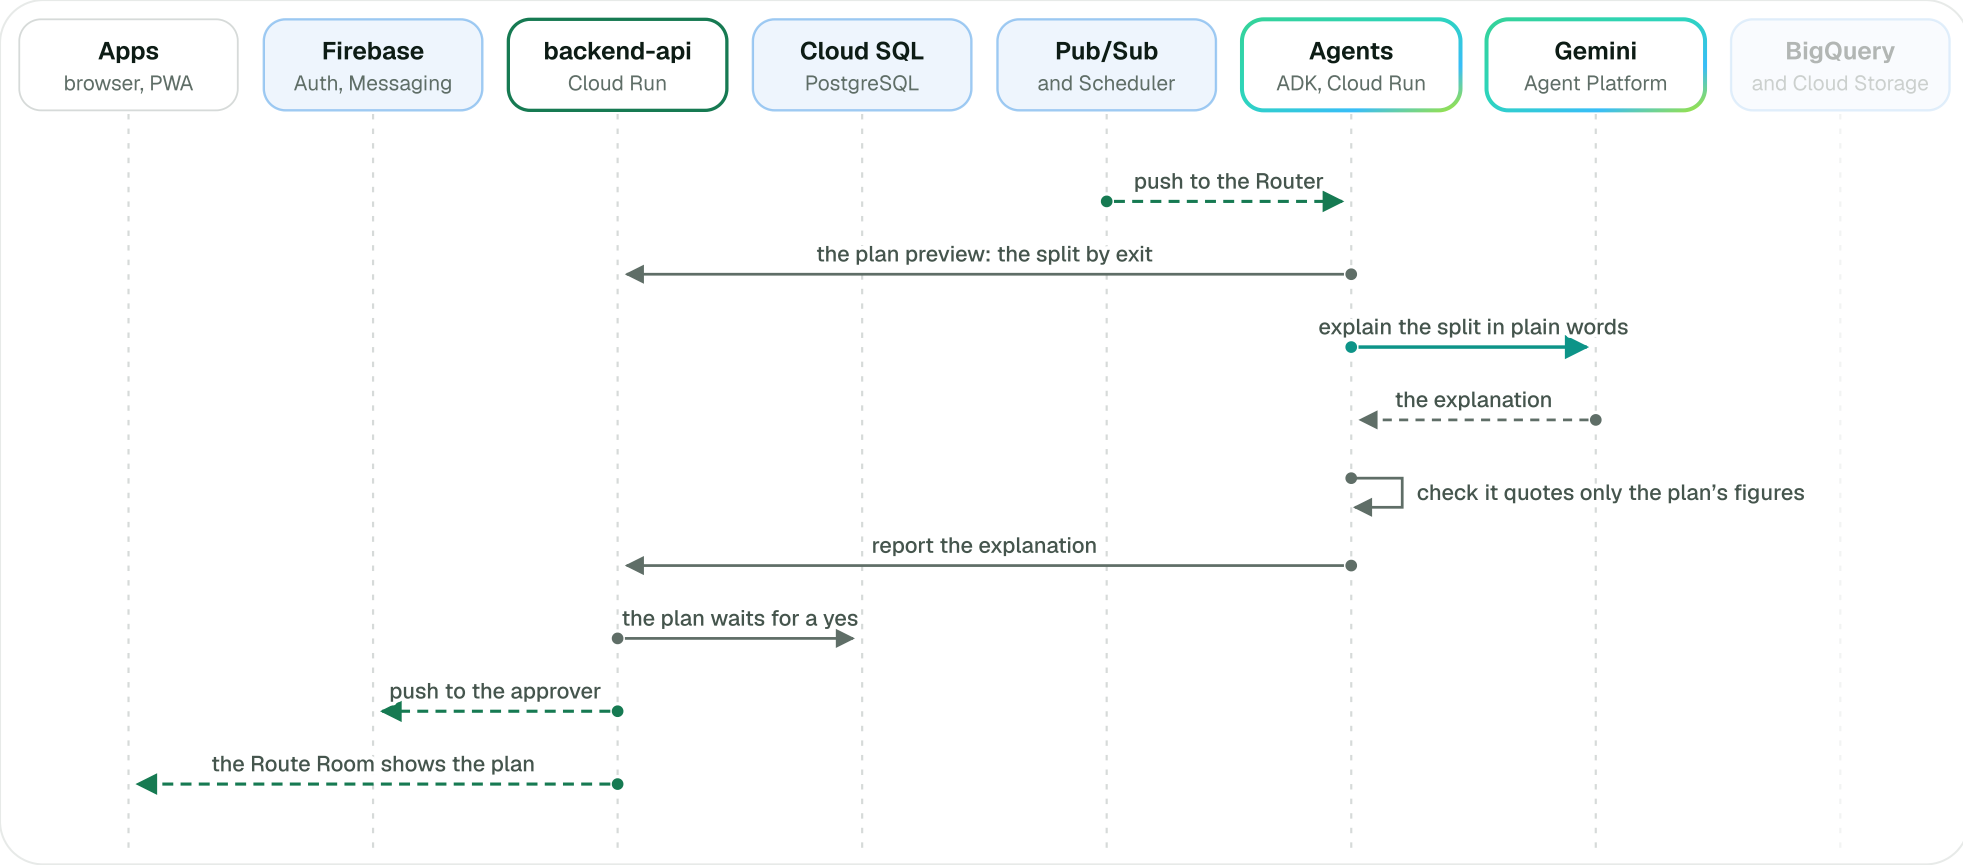


The table comes from backend-api, never from the model. A note that quotes a figure not in the table is dropped, and the template stands in.

Decide · Router

backend-api splits the batch across the exits by the client's gates and guardrails. The Router reads the plan and Gemini Pro explains it in plain words for the approver, quoting only the plan's own figures. The plan then waits for a person.

Pub/Sub | Cloud Run | Agent Platform | Cloud SQL | Cloud Messaging

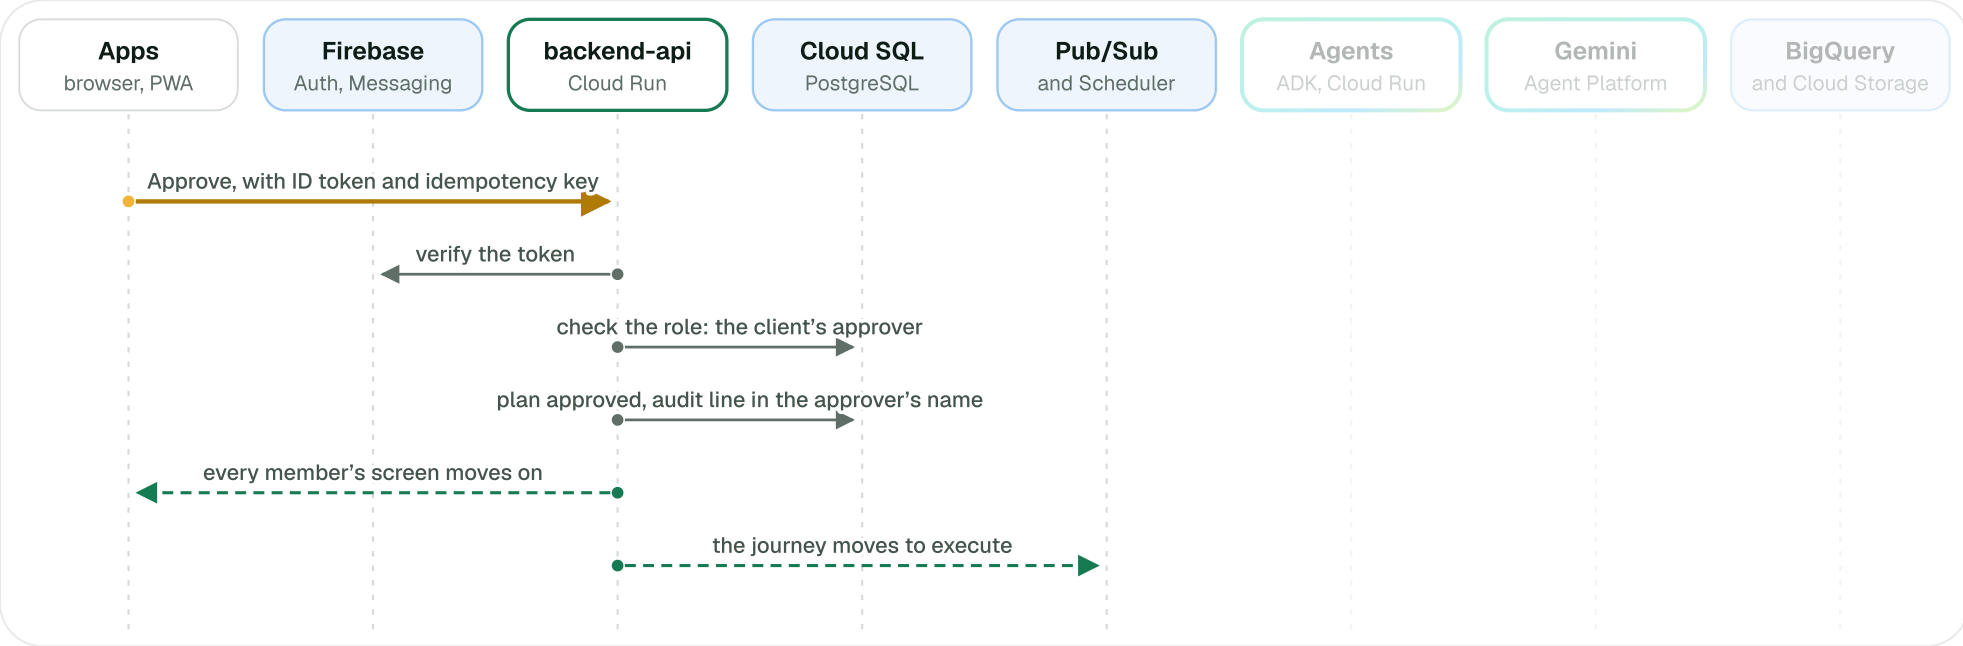


If the explanation fails a check, the approver still sees the plan, explained by the template.

Approve · a person

The approval is a person, never an agent. The approver reads the plan, the notes and the explanation, and says yes in the workspace app. backend-api checks the role, records the yes in the approver's name, and only then sends the event that starts execution.

Firebase Authentication | Cloud Run | Cloud SQL | Pub/Sub

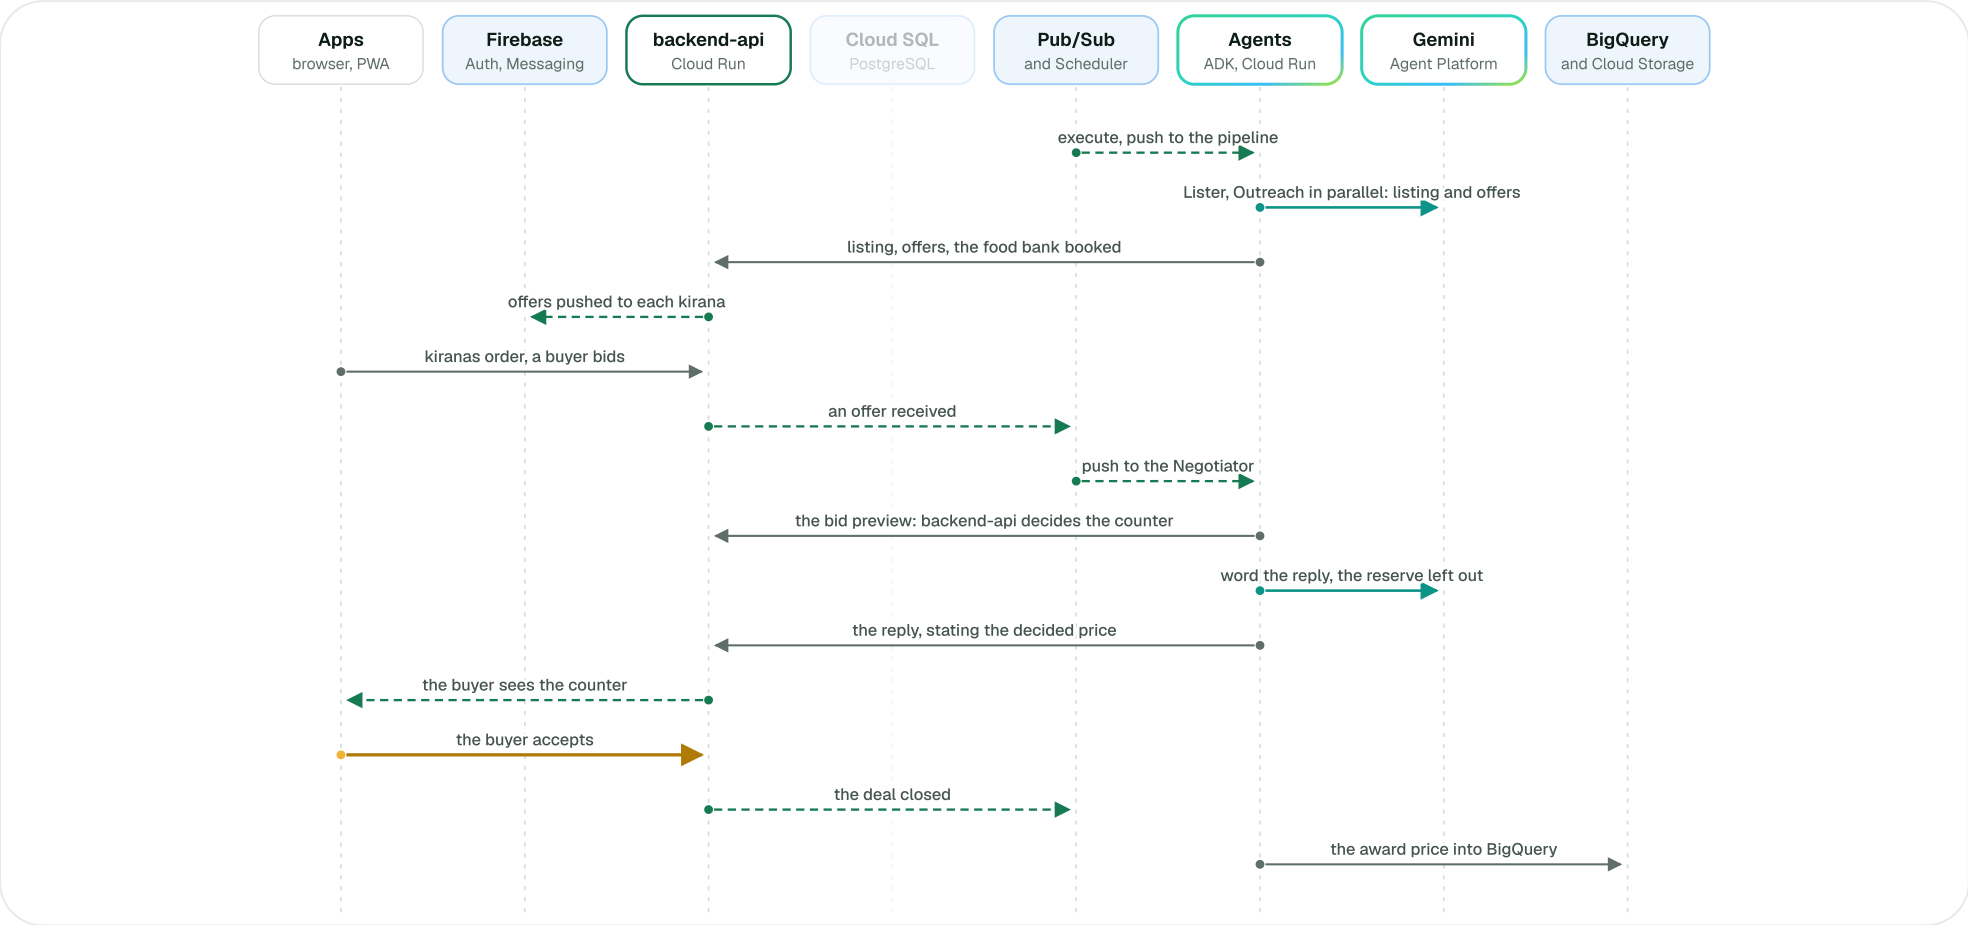


The yes carries an idempotency key, so a second tap on a slow network approves once.

Execute · Lister, Outreach, Donation, Negotiator

One event starts three agents at once in an ADK ParallelAgent: the Lister writes the marketplace listing, Outreach writes the kiranas' offer in Hindi, English and Marathi, and Donation books the food bank. Kiranas order in the app; a buyer bids, backend-api decides the counter, and the Negotiator words the reply.

Pub/Sub | Cloud Run | Agent Platform | Cloud Messaging | BigQuery | Cloud SQL

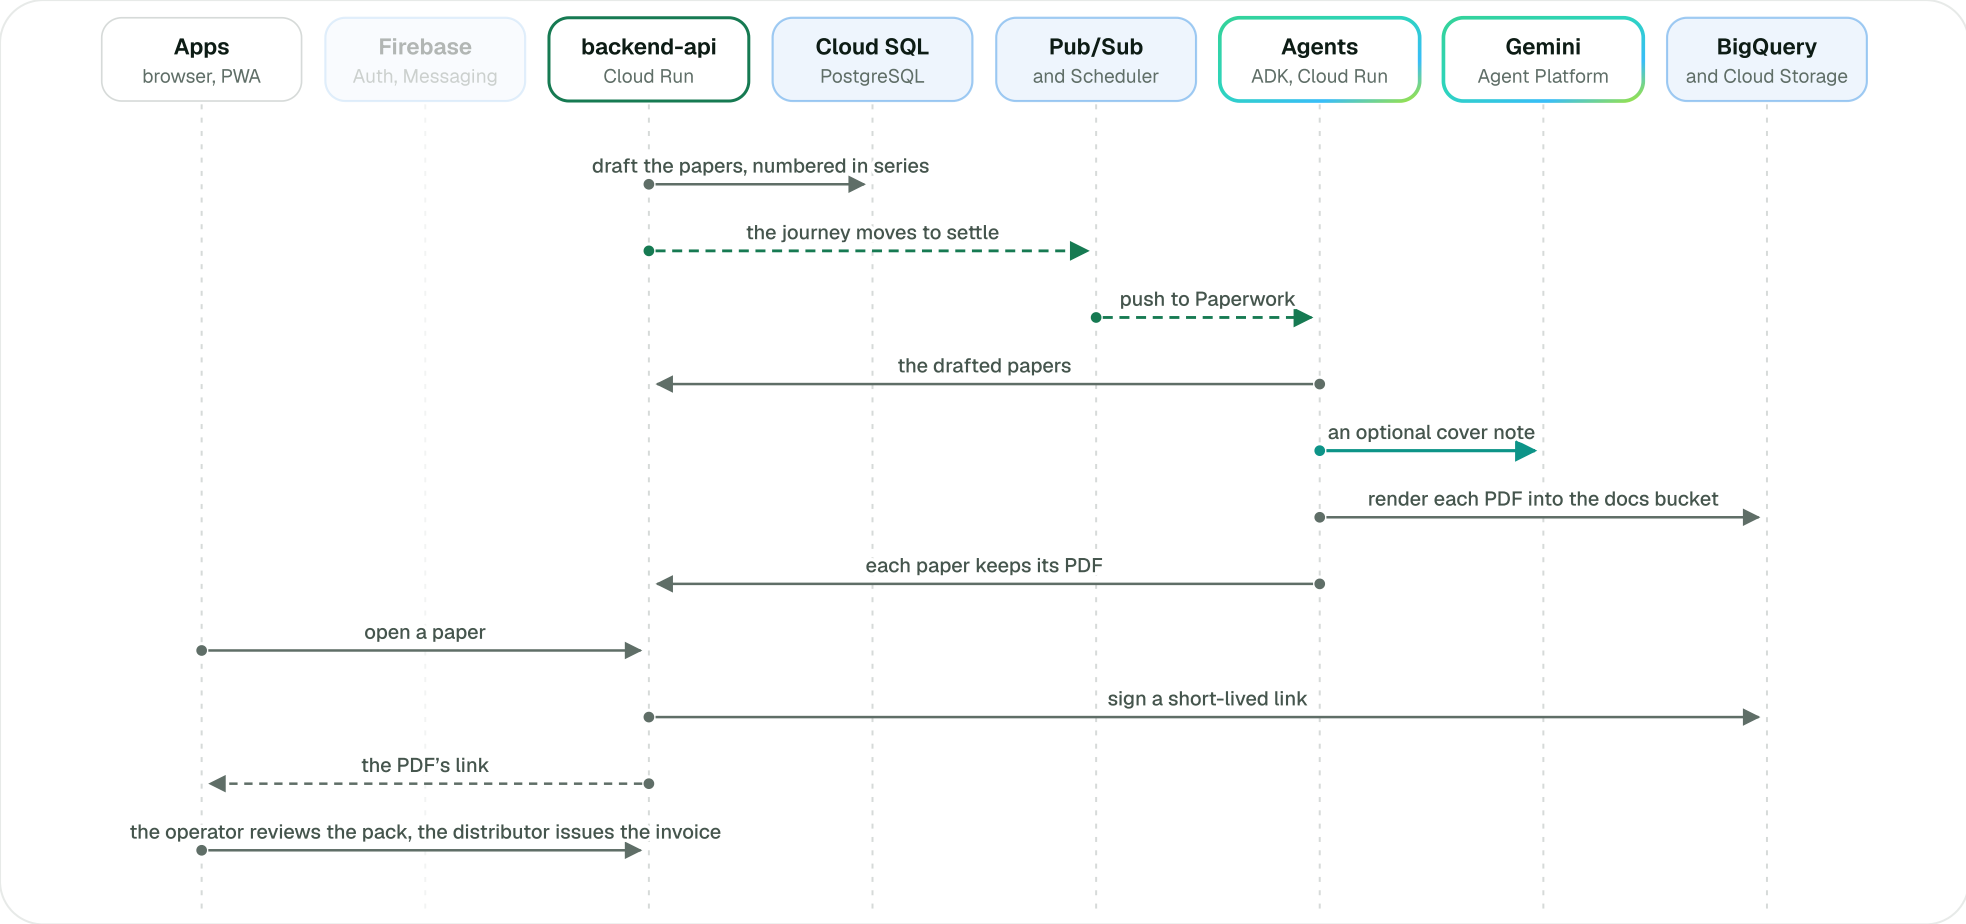


The reserve price never reaches a prompt, and a buyer's message is quoted between delimiters as data. A reply to a bid must state the decided price and nothing else.

Settle · Paperwork

Once every line of the plan is done, backend-api drafts the papers, numbered in their series: the tax invoice, the credit notes, the GST input-credit memo, the food bank's receipt. Paperwork renders each as a PDF with WeasyPrint into Cloud Storage, and members open them through short-lived signed links.

Cloud SQL | Pub/Sub | Cloud Run | Agent Platform | Cloud Storage

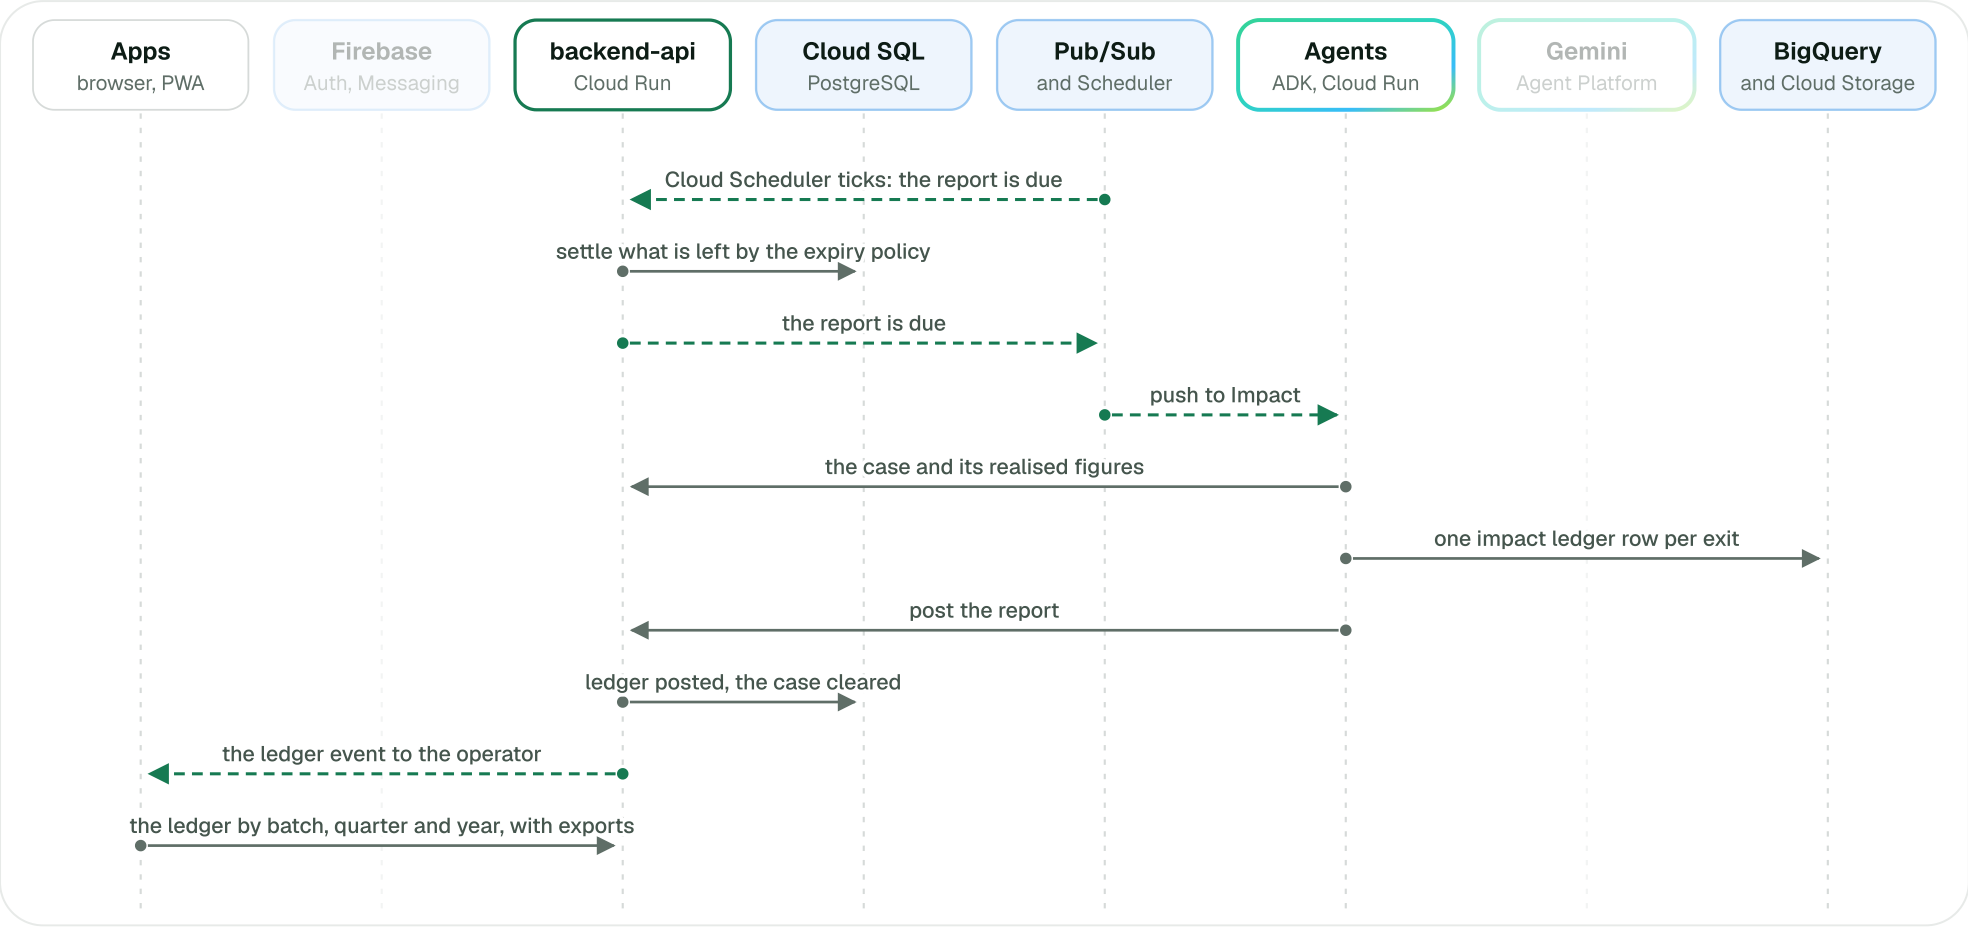


The figures on every paper are backend-api's. Gemini Flash may add a one-line cover note; if it fails, the paper ships without it.

Report · Impact

When the batch's report falls due, or the console calls expiry day, backend-api closes what is left by the client's expiry policy. Impact posts the ledger: what each exit recovered, the input credit kept and reversed, the kilos kept out of landfill, the meals. The rows go to BigQuery's impact table and to the workspace's ledger.

Cloud Scheduler | Pub/Sub | Cloud Run | Cloud SQL | BigQuery



The ledger reads by batch, by quarter of the Indian financial year and by year, and exports the BRSR table, the GST summary and the ledger itself.

7 IDENTITY AND SECURITY

Every caller signs as someone, and nobody holds a key

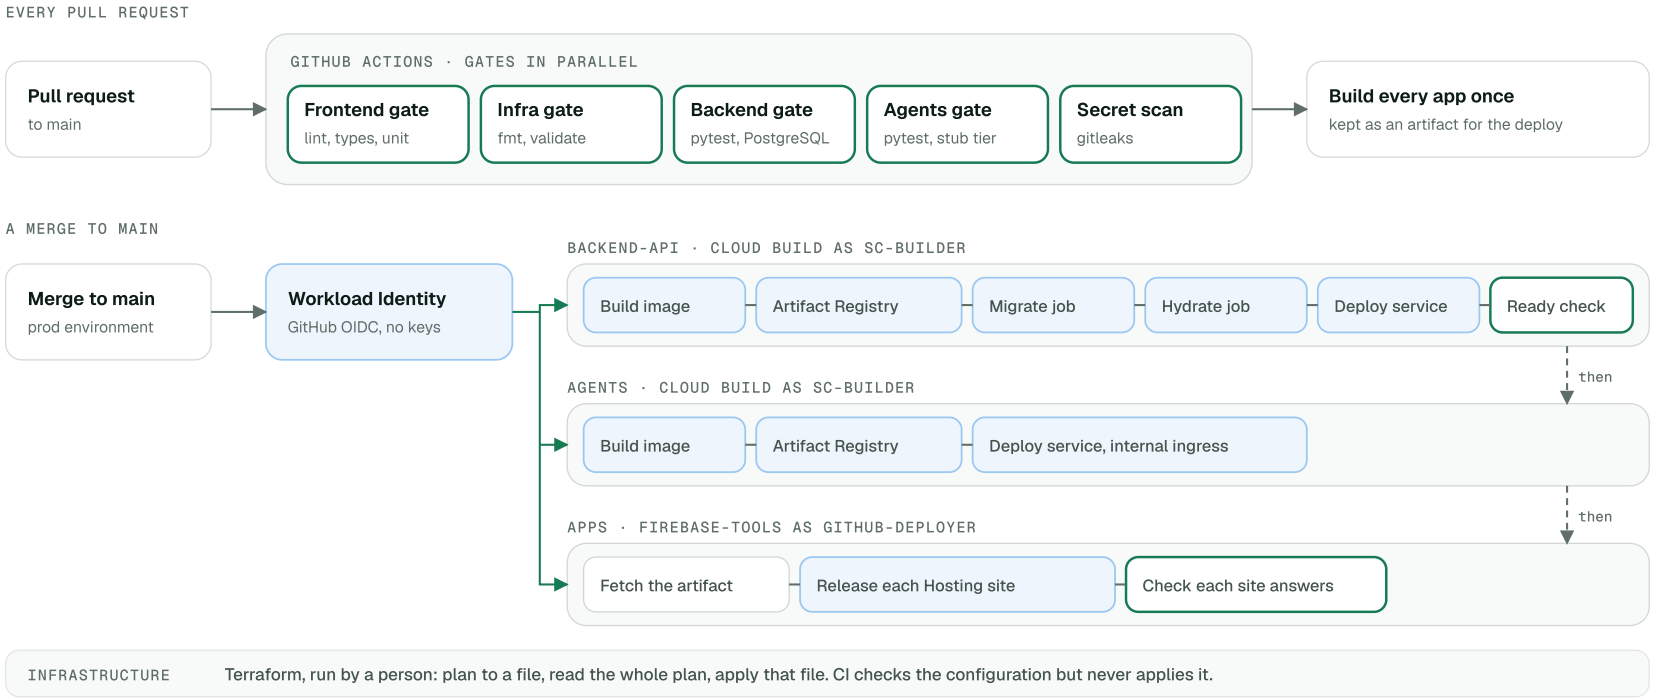
People sign in with Firebase; machines sign with OIDC tokens Google mints for a service account. Each account may do one job, and the database accepts IAM logins only.

WHO ACTS	SIGNS AS	MAY
A member in an app Firebase email and password	**`Firebase ID token`** verified by backend-api	**See a view cut to their role** a buyer never sees the reserve, a kirana sees its own offer, partners never see the brand's margins
Cloud Scheduler, Pub/Sub the clock and the event bus	**`sc-invoker`** OIDC token for a fixed audience	**Call the tick and the agents** Cloud Run's IAM lets nobody else into the agents service
The agents service on Cloud Run	**`sc-agents`** Google ID token to backend-api	**Read facts, report words** Agent Platform, BigQuery, Cloud Storage and `/internal`; never the database
backend-api and its jobs Cloud Run service and jobs	**`sc-api, sc-migrator`** IAM database logins	**Write the record** data access as the app, schema changes only as the migrator; the audit log refuses edits even to its owner
GitHub Actions on `main`, environment `prod`	**`github-backend, github-deployer`** Workload Identity Federation	**Start Cloud Build as sc-builder, release Hosting** nothing else
A developer's laptop their own gcloud login	**`sc-api-local, sc-agents-local`** impersonated in code	**Run the real stack locally** the same Firebase pool, secrets, topics and buckets, on a local PostgreSQL

8 CI/CD

GitHub Actions gates every change; Cloud Build ships it to Cloud Run

A pull request runs five gates in parallel and builds the apps once. A merge to `main` signs in to Google with Workload Identity Federation, has Cloud Build build, migrate and deploy backend-api and then the agents, and releases the apps' build to Firebase Hosting.



Jobs wait for the gates they need, and deploys queue instead of overlapping. Every action is pinned to a commit SHA, the workflow token is read-only, and only the deploying jobs may request an identity token.

Path-aware

A merge rebuilds backend-api only when `backend-api/` changed and the agents only when `agents/` changed; the apps release on every merge that touches the frontend, the designs or the infrastructure.

Migrate before deploy

Cloud Build runs the migration job on the new image before the service moves to it, so the schema is always ahead of the code reading it.

Ordered releases

backend-api deploys first, then the agents that call it, then the apps that read it, because the landing page prerenders against the live API.

Images owned by the build

Cloud Build tags each image with its commit; Terraform owns the service's configuration and ignores the image, so neither undoes the other.

9 TESTING

Proven at every layer, from a pixel to a live Gemini reply

Fast gates run on every pull request. The browser suites (accessibility, end to end, parity, live and journey) and the live-model evals run when we ask for them, on a real machine, against real builds and real services.

Gates

EVERY PULL REQUEST

`Vitest · pytest · ruff · ESLint · svelte-check · gitleaks`

- Frontend: unit tests, drift tests that hold ported CSS to the design, golden texts of every role's screens, seed and icon checks.
- Backend: pytest on a real PostgreSQL service container, the live journey walked over HTTP with in-memory Pub/Sub, Storage and Cloud Messaging.
- Agents: every pipeline on recorded Gemini responses, its exact requests, its fallbacks, and the reserve never reaching a model.
- Contracts: the money rules exist in JavaScript, TypeScript and Python, and shared fixtures hold all three to the same answers.

Accessibility

ON REQUEST

`Playwright · @axe-core/playwright · WCAG 2.2 AA`

Each app's production build, scanned in desktop, tablet and phone viewports, light and dark, with zero violations. Keyboard specs cover focus in sheets and alerts, the menu-button pattern and sign-in by keyboard; a motion spec fails any endless animation. A coverage check fails the run if an app uses a design-system component that no scan had on screen.

End to end

ON REQUEST

`Playwright · Chromium · Firefox · WebKit`

The console's flows, the workspace journey on its stub, and the guided demo played beat by beat, on each app's build. Smoke runs in Firefox and in WebKit on iPad and iPhone profiles.

Parity

ON REQUEST

`Playwright · pixelmatch`

Each app's build is compared with its approved design3 prototype, pixel by pixel, at desktop, tablet and phone sizes in both themes. The port cannot drift from the design that was signed off.

Live

ON REQUEST

`Vitest + jsdom · Playwright · Firebase Auth`

The workspace's live tests draw its screens on what backend-api actually answered at moments of the journey, recorded from the running API. The live end-to-end suite drives the landing page and the console against the real API and real Firebase sign-in, and reads each change back from the audit log under its own trace.

Journey

ON REQUEST

`Playwright · Firebase custom tokens · live Gemini`

Munchly Chips, Munchly Mango and their leftover variants. From a console reset to the impact report, every person acts in the real UI in turn: the operator, the distributor, each ordering kirana, the buyer and the food bank, with the agents on live Gemini. Sign-in uses custom tokens minted by an impersonated service account, so no password is handled. Each run records a video, a still per step and a report of findings, against the local stack or production.

Evals

ON REQUEST

`live Gemini · deterministic scorers · Gemini Pro judge`

Every agent's eval set runs through the same writer the pipeline calls: label photos rendered with blur, glare and tape for Vision, odd DMS headers for the Data agent, hostile buyer messages for the Negotiator. Scorers check the hard rules, a judge scores the words on a rubric, and results land in BigQuery.

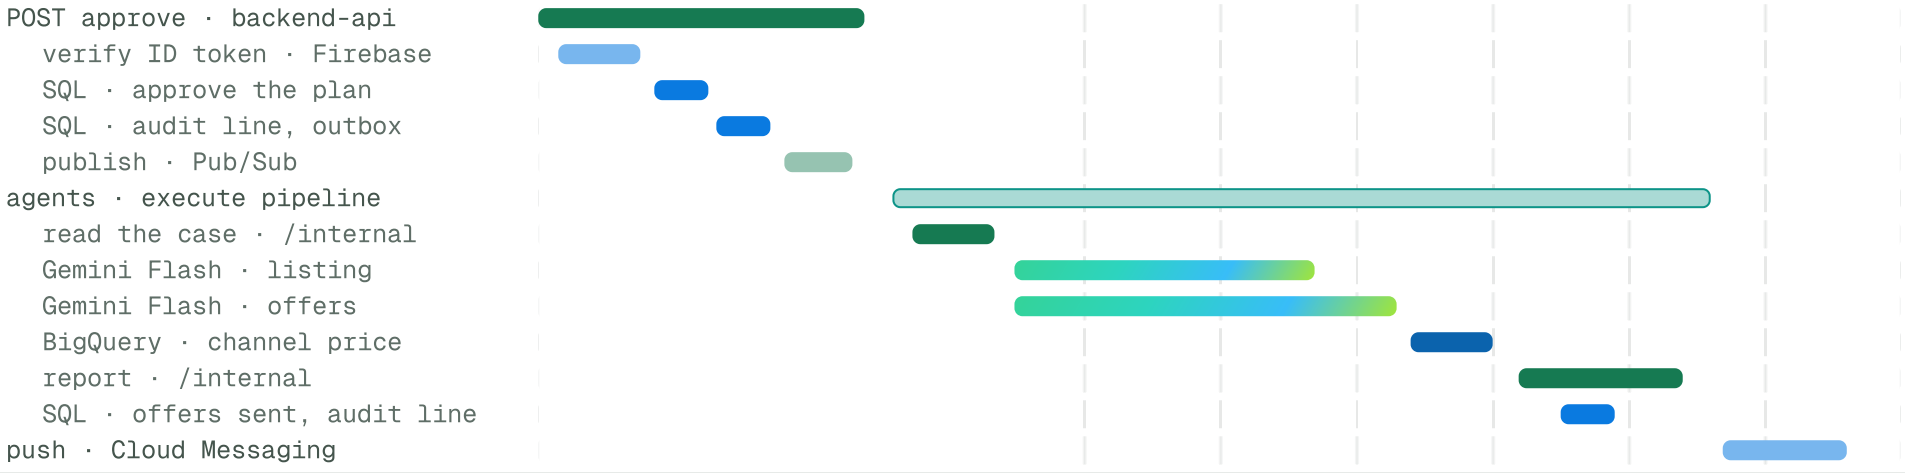
SUITE	DESIGN PROTOTYPE	APP CODE	REAL BROWSER	BACKEND-API	POSTGRESQL	FIREBASE AUTH	PUB/SUB, STORAGE	AGENTS	GEMINI
Frontend gate	●	●							
Backend gate				●	●	◐	◐		
Agents gate				◐			◐	●	◐
Accessibility		●	●						
End to end		●	●						
Parity	●	●	●						
Live		●	●	●	●	●			
Journey		●	●	●	●	●	●	●	●
Evals								●	●

● the real thing ◐ recorded, faked or in memory

10 OBSERVABILITY

One trace, from a tap to an agent's report

The apps send a `traceparent` with every call. Cloud Run keeps it, backend-api continues it into its spans, log lines and audit rows, and the outbox hands it to Pub/Sub, so the agents' run and its Gemini calls land in the same trace.



An illustrative trace as Cloud Trace draws it. Bars show order and nesting; their lengths are illustrative.

Logs carry the trace

Every JSON line written during a request names its trace and span, so Logs Explorer shows it under Cloud Run's own request line.

Audit rows carry the trace

Each change's audit row keeps the trace of the request that made it, so a line in the console leads back to its request.

Agent runs carry the trace

Each run's report and its BigQuery row hold its trace id, with model, latency and whether it fell back.

Sampled where it counts

A request Cloud Run sampled is always kept, a share of the rest is sampled, and every request keeps its trace id in its logs regardless.

11 PERFORMANCE, COST AND SCALE

Measured on production; it runs for about ₹1,350 to ₹3,600 a month

Measured on production between 8 and 10 October 2026: Cloud Run, Cloud SQL and BigQuery in `asia-south1`, and Gemini on Agent Platform's global endpoint. Prices are Google's INR list prices for `asia-south1` on 10 October 2026.

Measure	Result
An API call on Cloud Run	102 ms at the median, 412 ms at p95, over 7,251 calls
An agent's decision on Gemini Pro	8.7 s at the median; a Flash step takes 2 to 4 s
Gemini, a call	₹0.06 to ₹0.50
Gemini, a batch's whole journey	₹1.70 for the Masala Chips, about 0.01% of the ₹21,152 it recovered
Cloud SQL on `db-f1-micro`	27% CPU at its peak, 11% on average
The whole journey, recorded on production	77 of 77 steps from a reset to the impact report, in 9 min 46 s
First paint on a 3G line (the builds)	the console 2.6 s (from 8.1 s); the landing page 2.4 s (from 6.4 s), its loader lifted at 7.8 s (from 14.6 s)

Each agent on production: 784 agent runs and 257 Gemini calls, read from BigQuery's `agent_runs`. The 90th percentile includes the retries on the Pro preview's shared quota.

Agent	Model	Median	90th percentile	₹ a call
Paperwork	Flash	2.1 s	5 s	0.06
Lister	Flash	2.5 s	17 s	0.15
Outreach	Flash	3 s	9.2 s	0.22
Vision	Flash	4 s	48.6 s	0.34
Negotiator	Pro	5.6 s	14.9 s	0.15
Valuer	Pro	9.2 s	26.6 s	0.50

Agent	Model	Median	90th percentile	₹ a call
Router	Pro	10.8 s	38.4 s	0.26

Flash does the volume; Pro only the decisions a person approves.

Running it, a month:

Service	What runs	₹ a month
Cloud SQL for PostgreSQL 18	`db-f1-micro` (a shared vCPU, 0.6 GB) at ₹1.21 an hour, zonal; 10 GB SSD, backups, point-in-time recovery	1,115
Cloud Run	backend-api (1 vCPU, 512 MiB) and the agents (1 vCPU, 1 GiB), 0 to 2 instances, billed per request; most of it is the live stream while people use the app	100 to 600
Gemini on Agent Platform	Flash at ₹144 and ₹720 a million tokens in and out, Pro at ₹192 and ₹1,152; about ₹2 a batch's journey	130 to 1,900
Pub/Sub, Cloud Scheduler, BigQuery	the events, the minute's tick, the analytics tables	0
Firebase Hosting, Authentication, Cloud Messaging	four sites, one user pool, the pushes	0
Secret Manager, Cloud Build, Cloud Trace, Logging, Monitoring	the secrets, the releases, spans, JSON logs, alerts	0
In all	at the prototype's traffic, from light use to heavy	**1,350 to 3,600**

The budget, GBP 30 (about ₹3,840), emails at 50%, 90% and 100%. A full eval run on live Gemini costs about ₹260 and runs on request only.

The free tiers it runs inside: Cloud Run's 2 million requests and 180,000 vCPU-seconds a month; BigQuery's 1 TiB of queries and 10 GiB stored; Pub/Sub's 10 GiB of messages; Cloud Scheduler's 3 jobs (it runs 2); Secret Manager's 6 secret versions; Cloud Trace's 2.5 million spans; Logging's 50 GiB; Cloud Build's 2,500 build-minutes; Firebase Hosting, email sign-in and Cloud Messaging; GitHub Actions on a public repository. Every library in the stack is free and open source.

How it scales. Each limit is a setting in Terraform. The Cloud SQL tier is the `db_tier` variable: `db-f1-micro` at about ₹1,115 a month, `db-g1-small` at about ₹3,200, or 2 vCPU and 8 GB at about ₹11,900. Each Cloud Run service's instance ceiling (2 today) is set in `run.tf` and `agents.tf`, and costs nothing until it is used. The agents are stateless and a batch's events stay in order whatever the number of instances, so more instances need no other change. For Gemini, Provisioned Throughput on Agent Platform would end the shared quota's 429s.

12 DESIGN DECISIONS

What was chosen, and why

Decision	Why
Design first, in design3, then port	one source of truth for designs, reviewed on a board before anything is built; tests hold the port to it
Static SvelteKit builds on Firebase Hosting	the apps are a prerendered page and client-side apps over an API; every byte is cacheable at the edge, and there is no rendering server to run
Firebase for hosting, sign-in and pushes; Cloud SQL for the record, not Firestore	relational truth with transactions across a change, its audit line and its outbox; SQL views the agents read; `LISTEN/NOTIFY` and SSE for the live screens
Only backend-api writes; the agents report to it	one place for every rule, an audit log complete by construction, and no database credentials outside the API
Gemini writes the words, backend-api computes the figures	money is never a model's guess; every figure is held to `money.js`; a model's sentence that quotes another figure is dropped
Two Gemini tiers	Pro where a person decides on the words (prices, the split, a counter); Flash where reading and drafting must be fast and cheap
Stateless ADK pipelines, one run an event	a run survives a cold start and a redelivery; nothing to resume; Pub/Sub's retries and dead letters do the rest
A state machine orchestrates; a table dispatches	the next step comes from code in backend-api and each event's pipeline from a fixed table, so the same facts always lead to the same steps; no model routes work
ADK workflow agents, no model-driven delegation	`SequentialAgent` and `ParallelAgent` give order and parallelism without a model picking the next agent; a writer cannot transfer
The agents on Cloud Run, not Agent Runtime	Pub/Sub pushes straight to them; they scale to zero beside backend-api, with one way to deploy, sign in, log and trace; backend-api is the record, so no managed sessions are needed
Gemini on Agent Platform, as a service account	no API key anywhere; the model ids are Terraform variables; Pro and Flash on the global endpoint, where the Gemini 3 models are served
An outbox, published after the commit	an event never leaves for a change that rolled back, and never goes missing for one that committed; ordering per batch
Cloud Scheduler for the clock	the journey moves on with every instance scaled to zero; the console can fire any timer at once

Decision	Why
A journey day from 1 to 1,440 minutes	a demo can walk the whole story in an hour on the real stack
The stub and the live source behind one interface	the guided demo and the prototype keep working with no backend, and a live build carries none of the stub
Email and password for everyone, no email sent	one shared user pool; accounts made by the API on a default password an operator hands over
Secrets by reference, no keys, IAM database logins	nothing secret in a file, a command line, state or git
Cloud Run scaling to zero, `db-f1-micro`	about GBP 9 a month for the platform; a cold start is covered by the console's and the workspace's splash
The photos and papers kept for good, the exports for 30 days	a batch's Record and its PDFs are its evidence; the exports are loaded into BigQuery
Cloud Build releases the services; firebase-tools the apps; Terraform owns neither's content	the Hosting provider cannot upload files, and Cloud Build owns the image
Every Terraform plan saved, read in full, then applied	nothing in the cloud changes unread
OpenTelemetry over OTLP to the Telemetry API	Google's route into Cloud Trace; logs, audit rows and agent runs carry every request's trace, sampled or not
WCAG 2.2 AA with zero axe violations	held by suites on the real apps, not by review alone

13 KNOWN GAPS AND RISKS

What is not done yet, and where it shows

- **One user pool** serves local development and production, so the same accounts sign in to both.
- **Cold starts.** Cloud Run scales to zero, so the first request after a quiet spell waits a few seconds. Cloud SQL is a shared core without an SLA, and the rate limiter counts per instance.
- **No custom domain:** the apps live on Firebase Hosting's own addresses, and the staff email domain is `smartclearance.example`.
- **Terraform runs from a workstation;** CI checks it but never plans. The browser suites, the journey suites and the evals are not in CI.
- **Firefox** cannot start in the agent's sandboxed shell, so the e2e suite's Firefox run needs a normal machine.
- **The prototypes' hosted copies** run on their mocks and are pinned to older commits.
- **Traces:** the apps start a trace per API call, not per click; the migrate and hydrate jobs are not traced.
- **The agents:** `infra/prod/monitoring.tf` describes the model timeout as 20 s where the code gives 30 s an attempt; the `sc-agents-harvest` entry point names a module that does not exist (`eval-harvest.sh` is the working route); the Data eval set's summary predates the shelf cases' removal.
- **BigQuery's** `shelf_counts` and the `cases.shelf` column are unused since the shelf check was retired.
- **The WCAG 2.2 criteria axe cannot check** (2.4.11, 2.5.7, 3.2.6, 3.3.7, 3.3.8) await a manual pass.
- **The landing page** ships about 160 kB of JavaScript, gzipped; the console's first load on a slow line is bound by its code in waves.
- **The agents' autonomy levels** (suggest, ask, act) are set per client and audited, but the live journey does not read them yet: an agent is on or off, and the journey's two people's gates are fixed.
- **Gemini 3.1 Pro is a preview** on a shared quota, which answers 429 under load; each call retries, then falls back to its template. Provisioned Throughput, or Pro on general availability, would end it.
- **Reject** sits beside Approve on the plan, drawn but not built: a plan is approved or left waiting.

14 WHAT COMES NEXT

The roadmap, area by area

Area	Next
Integrations	WhatsApp Business for the kiranas' offers, beside the in-app push; partner APIs for near-expiry marketplaces; DMS connectors (Bizom, Tally) instead of nightly CSV exports; live e-invoicing (IRN) and e-way bills through a GST Suvidha Provider
The agents	reject and re-plan: a rejection's reason becomes a rule, and the Valuer and the Router come back with two or three new plans, each with its money; prices and splits learned from every closed deal; demand forecasting per pincode with BigQuery ML and Agent Platform; Gemini Pro on general availability; offers in Tamil, Telugu, Bengali and Kannada
The product	full-credit take-back beside destruction at the godown; native iOS and Android apps on the same API, with offline capture at the godown and background sync; pharma and OTC profiles under the Drugs and Cosmetics Rules
The platform	self-serve onboarding for manufacturers; an own domain and single sign-on; data residency for enterprise clients; Provisioned Throughput for Gemini; Cloud Run and Cloud SQL sized up through their Terraform settings

A APPENDIX A

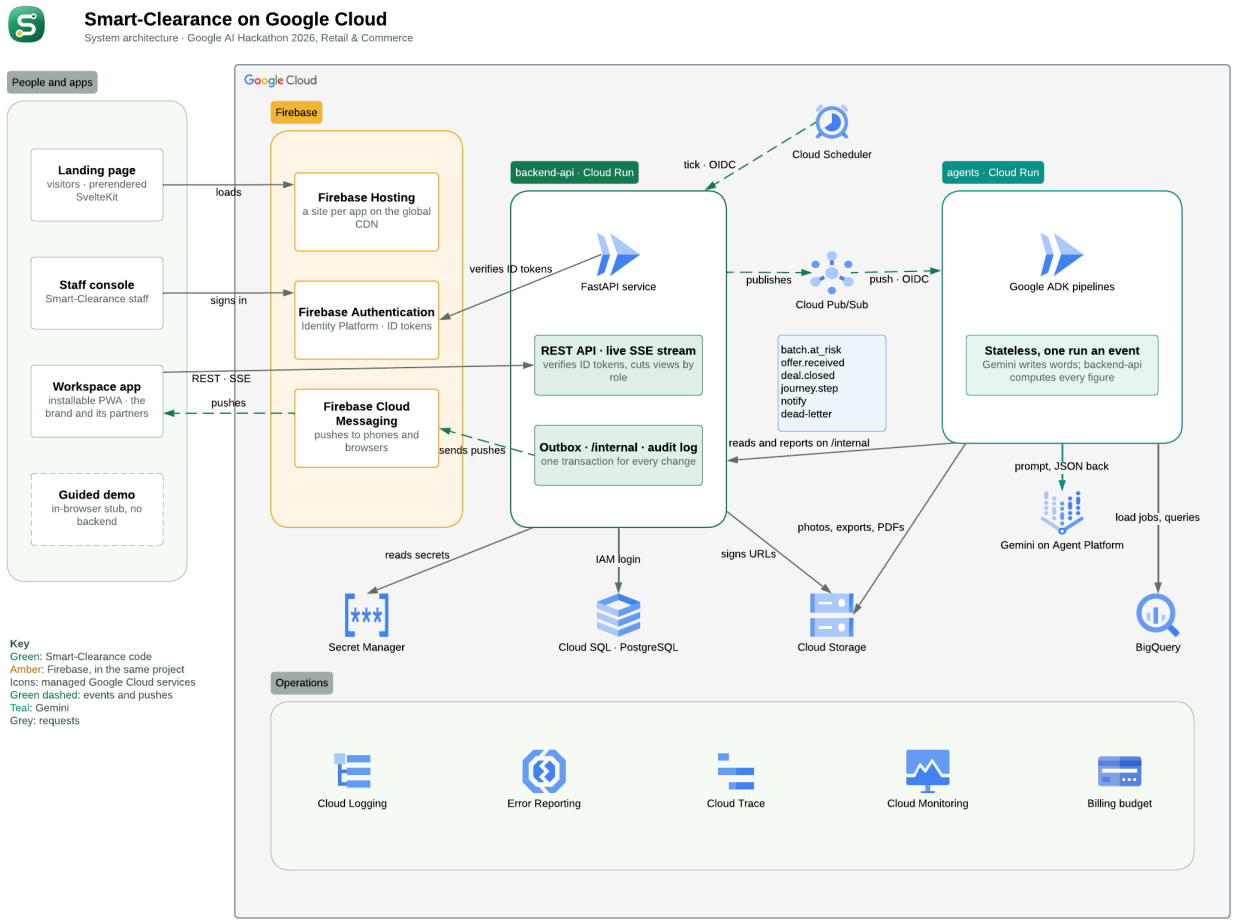
The diagrams

Every diagram in ARCHITECTURE.md, each a Lucidchart document drawn with Google Cloud's own icons, a page each. The link under each opens its source.

A.1	**Smart-Clearance on Google Cloud**	Smart-Clearance Architecture on Google Cloud, page 1
A.2	**Google Cloud, layer by layer**	Smart-Clearance Google Cloud and AI, page 1
A.3	**The core loop**	Smart-Clearance Application Flow · the core loop
A.4	**A batch's journey**	Smart-Clearance Application Flow · a batch's journey
A.5	**A case's phases**	Smart-Clearance · A case's phases
A.6	**Destroyed at the godown, against evidence and a second yes**	Smart-Clearance · Destroyed at the godown, against evidence and a second yes
A.7	**Ten agents along nine stops**	Smart-Clearance Architecture on Google Cloud, page 2
A.8	**How the agents are orchestrated**	Smart-Clearance · How the agents are orchestrated
A.9	**The agents and their models**	Smart-Clearance Google Cloud and AI, page 6
A.10	**A model call, bounded**	Smart-Clearance Google Cloud and AI, page 7
A.11	**backend-api's layers**	Smart-Clearance · backend-api's layers
A.12	**Pub/Sub and Cloud Scheduler**	Smart-Clearance Google Cloud and AI, page 3
A.13	**Where the data lives**	Smart-Clearance Google Cloud and AI, page 4
A.14	**The data: the platform's tables**	Smart-Clearance · The data: the platform's tables
A.15	**The data: the journey's tables**	Smart-Clearance · The data: the journey's tables
A.16	**The frontend's packages**	Smart-Clearance · The frontend's packages
A.17	**Rendering and data**	Smart-Clearance · Rendering and data
A.18	**design3's structure**	Smart-Clearance · design3's structure
A.19	**Who signs as whom**	Smart-Clearance Google Cloud and AI, page 5
A.20	**Cloud Run: two services, two jobs**	Smart-Clearance Google Cloud and AI, page 2
A.21	**GitHub Actions and Cloud Build**	Smart-Clearance Architecture on Google Cloud, page 3
A.22	**Local development**	Smart-Clearance · Local development
A.23	**One request, one trace**	Smart-Clearance · One request, one trace

A.1 Smart-Clearance on Google Cloud

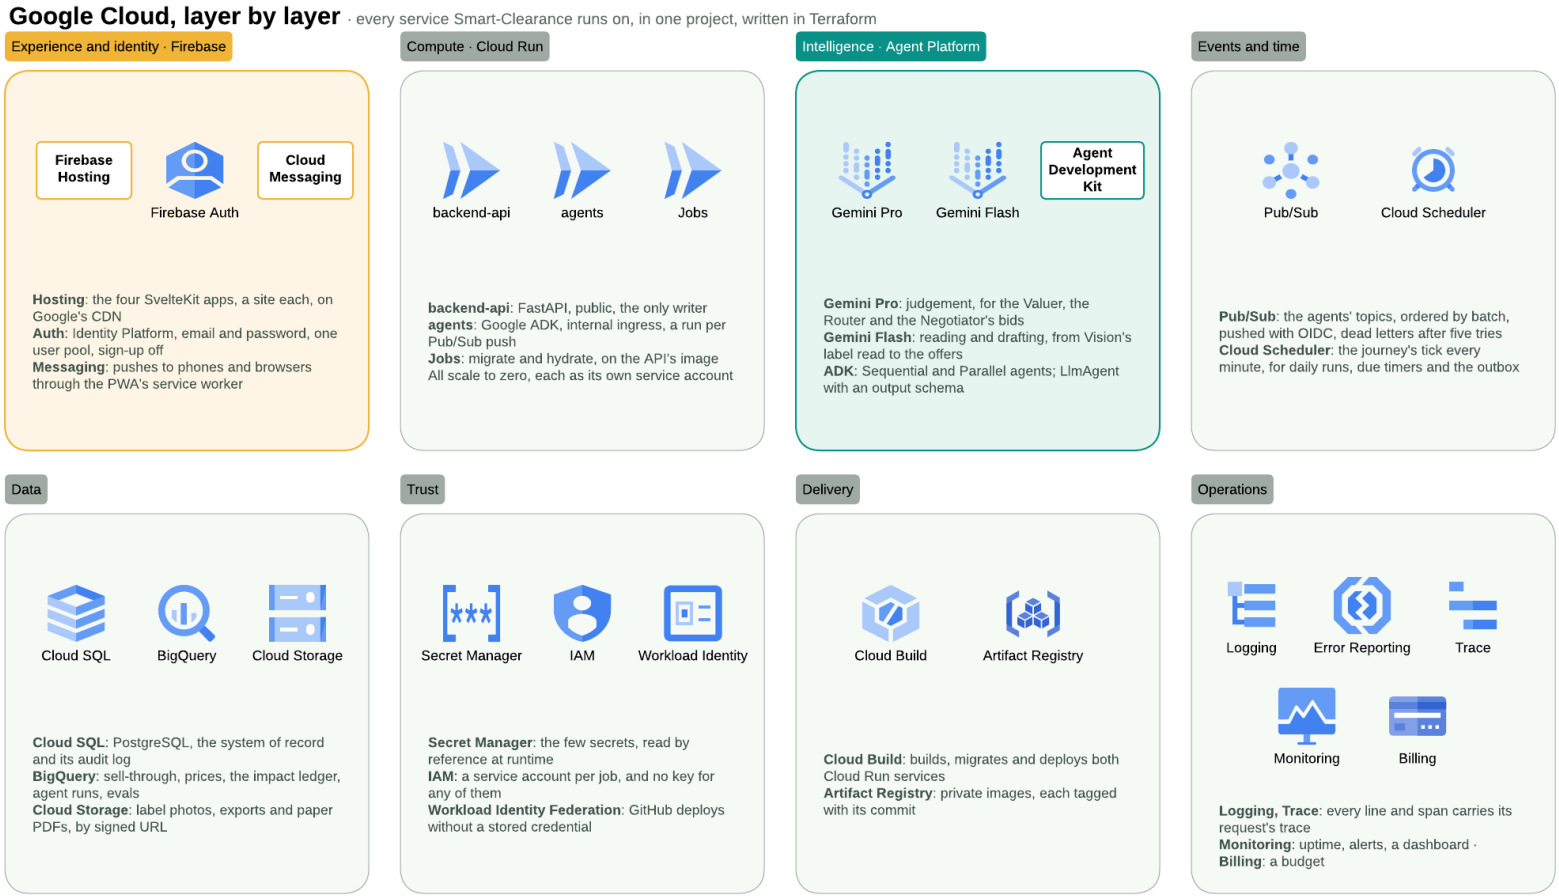
Smart-Clearance on Google Cloud: the four apps, Firebase, backend-api and the agents on Cloud Run, Pub/Sub and Cloud Scheduler, Gemini on Agent Platform, Cloud SQL, Cloud Storage, BigQuery, Secret Manager and operations



Lucid: Smart-Clearance Architecture on Google Cloud, page 1 · https://lucid.app/lucidchart/0e33214e-b46f-4cfb-9413-b2c40de7d1c6/view

A.2 Google Cloud, layer by layer

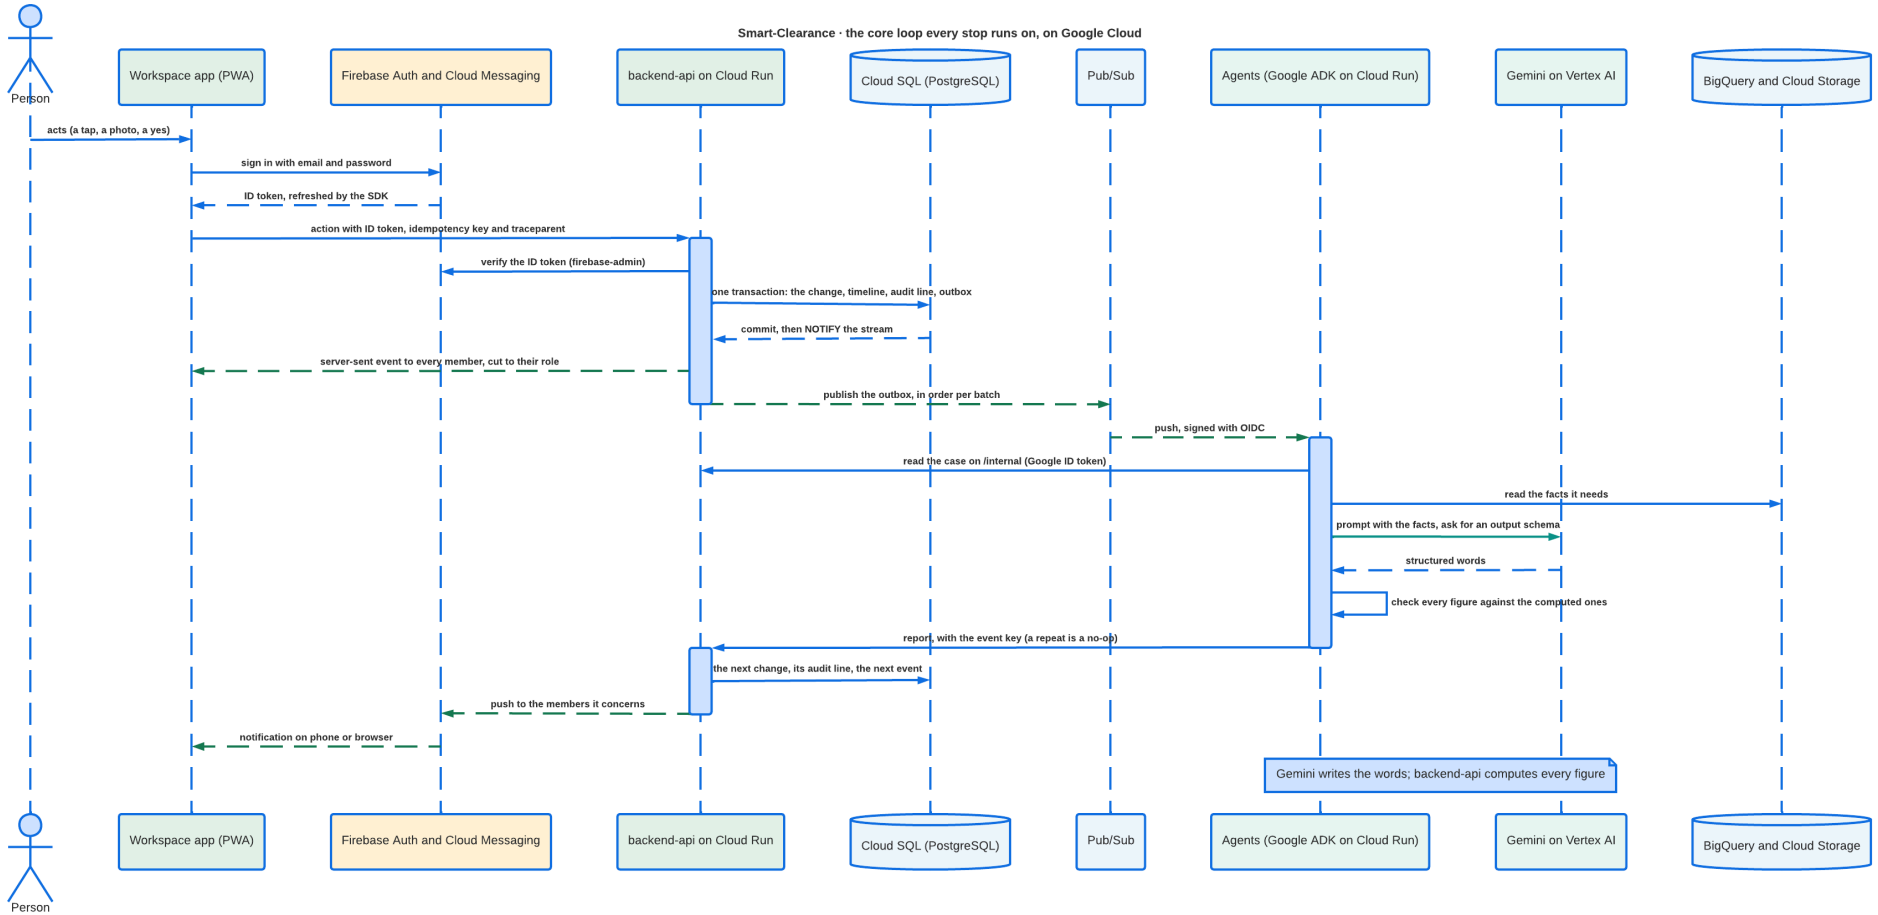
Google Cloud, layer by layer: Firebase, Cloud Run, Agent Platform, events and time, data, trust, delivery and operations



Lucid: Smart-Clearance Google Cloud and AI, page 1 · https://lucid.app/lucidchart/631bbb66-9b43-4e20-abf0-f44fa6a68174/view

A.3 The core loop

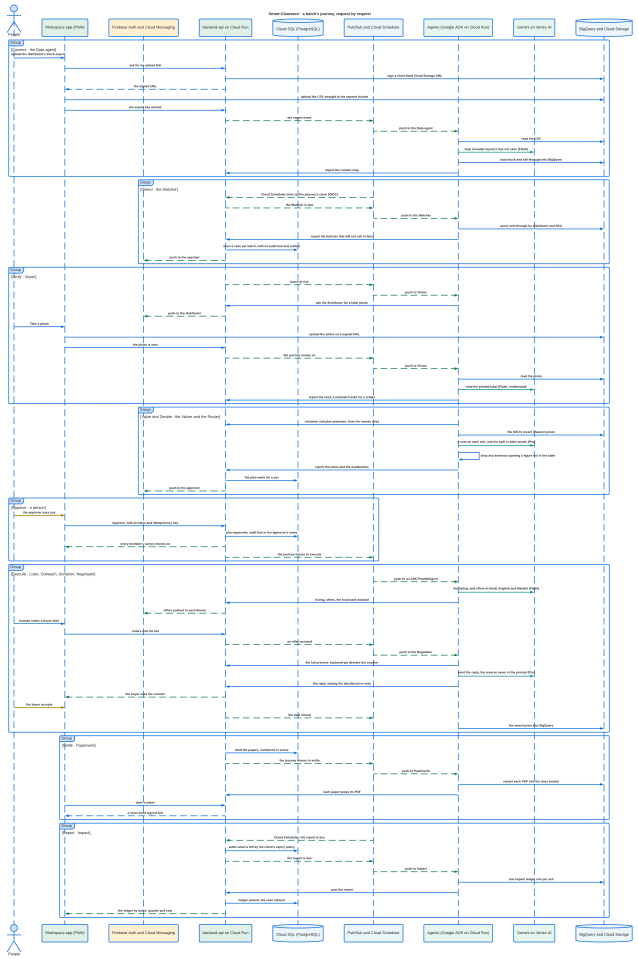
The core loop: a person acts, backend-api commits the change with its outbox, Pub/Sub pushes to the agents, the agents read facts, ask Gemini, check the words and report, and the members are told



Lucid: Smart-Clearance Application Flow · the core loop · https://lucid.app/lucidchart/2965aa52-3190-4d68-8791-9c95c8fd7bbc/view

A.4 A batch's journey

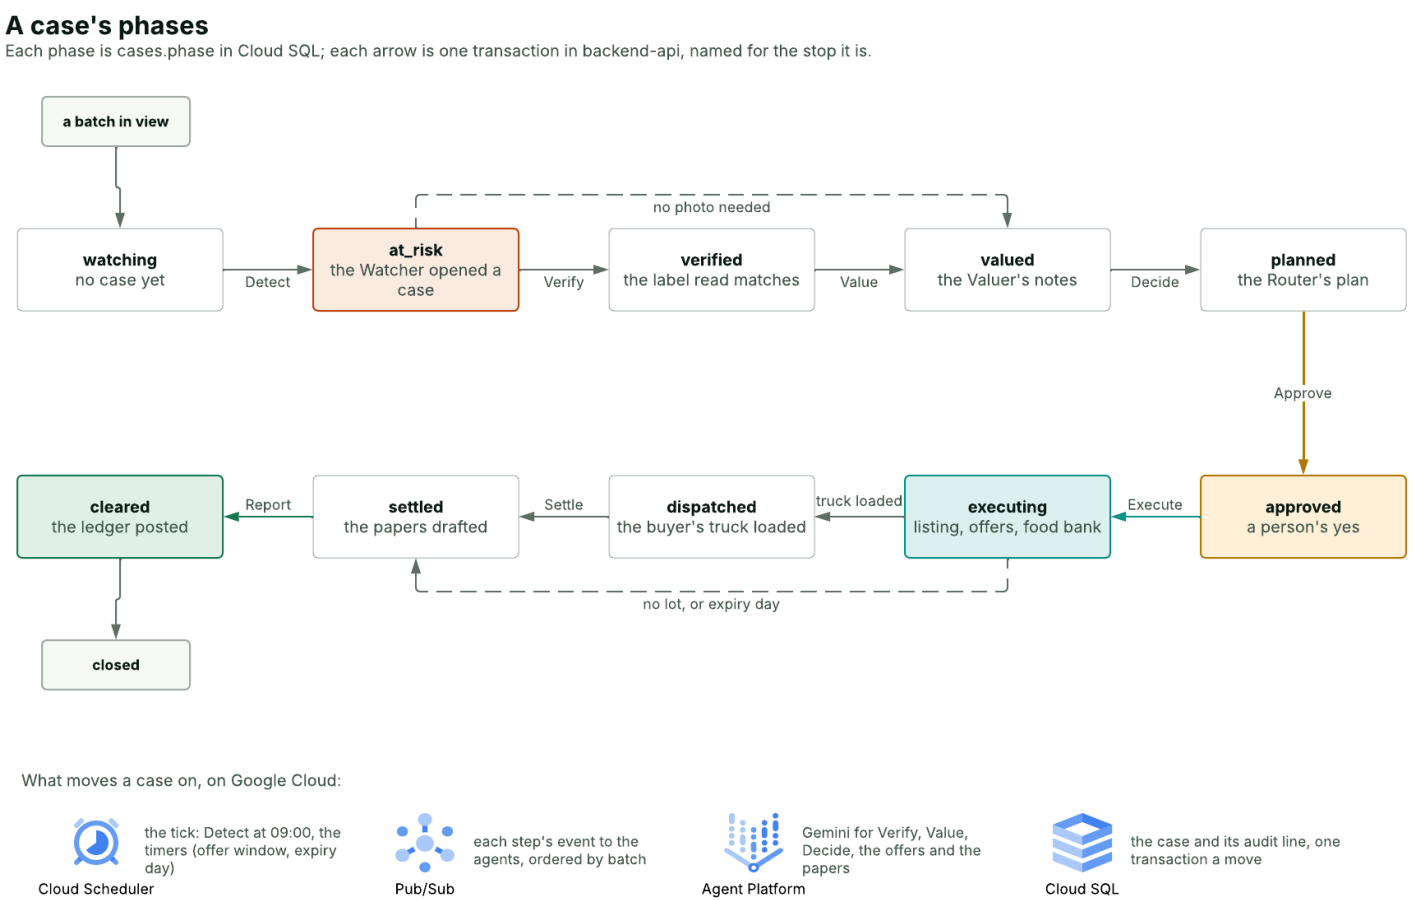
A batch's journey, request by request: Connect, Detect, Verify, Value and Decide, Approve, Execute, Settle and Report, across the workspace app, Firebase, backend-api, Cloud SQL, Pub/Sub and Cloud Scheduler, the agents, Gemini, BigQuery and Cloud Storage



Lucid: Smart-Clearance Application Flow · a batch's journey · https://lucid.app/lucidchart/c2efb94b-7aa2-4cc1-a6cc-af4980ff255f/view

A.5 A case's phases

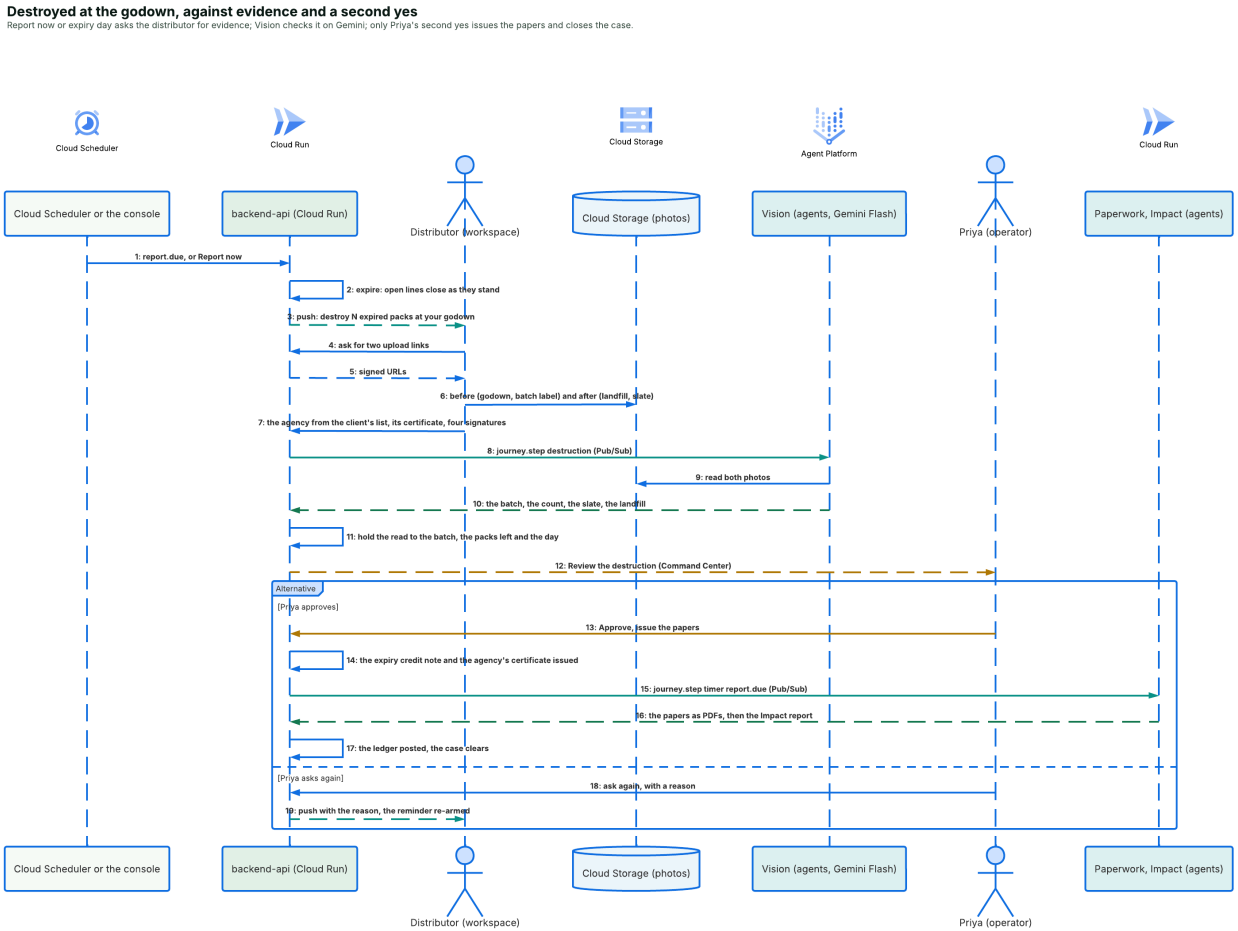
A case's ten phases, from watching to cleared: Detect, Verify (or no photo needed), Value, Decide, a person's Approve, Execute, the buyer's truck or straight to Settle, then Report; Cloud Scheduler, Pub/Sub, Agent Platform and Cloud SQL move it on



Lucid: Smart-Clearance · A case's phases · https://lucid.app/lucidchart/2605b0cc-b62e-4e50-8844-c788931a94db/view

A.6 Destroyed at the godown, against evidence and a second yes

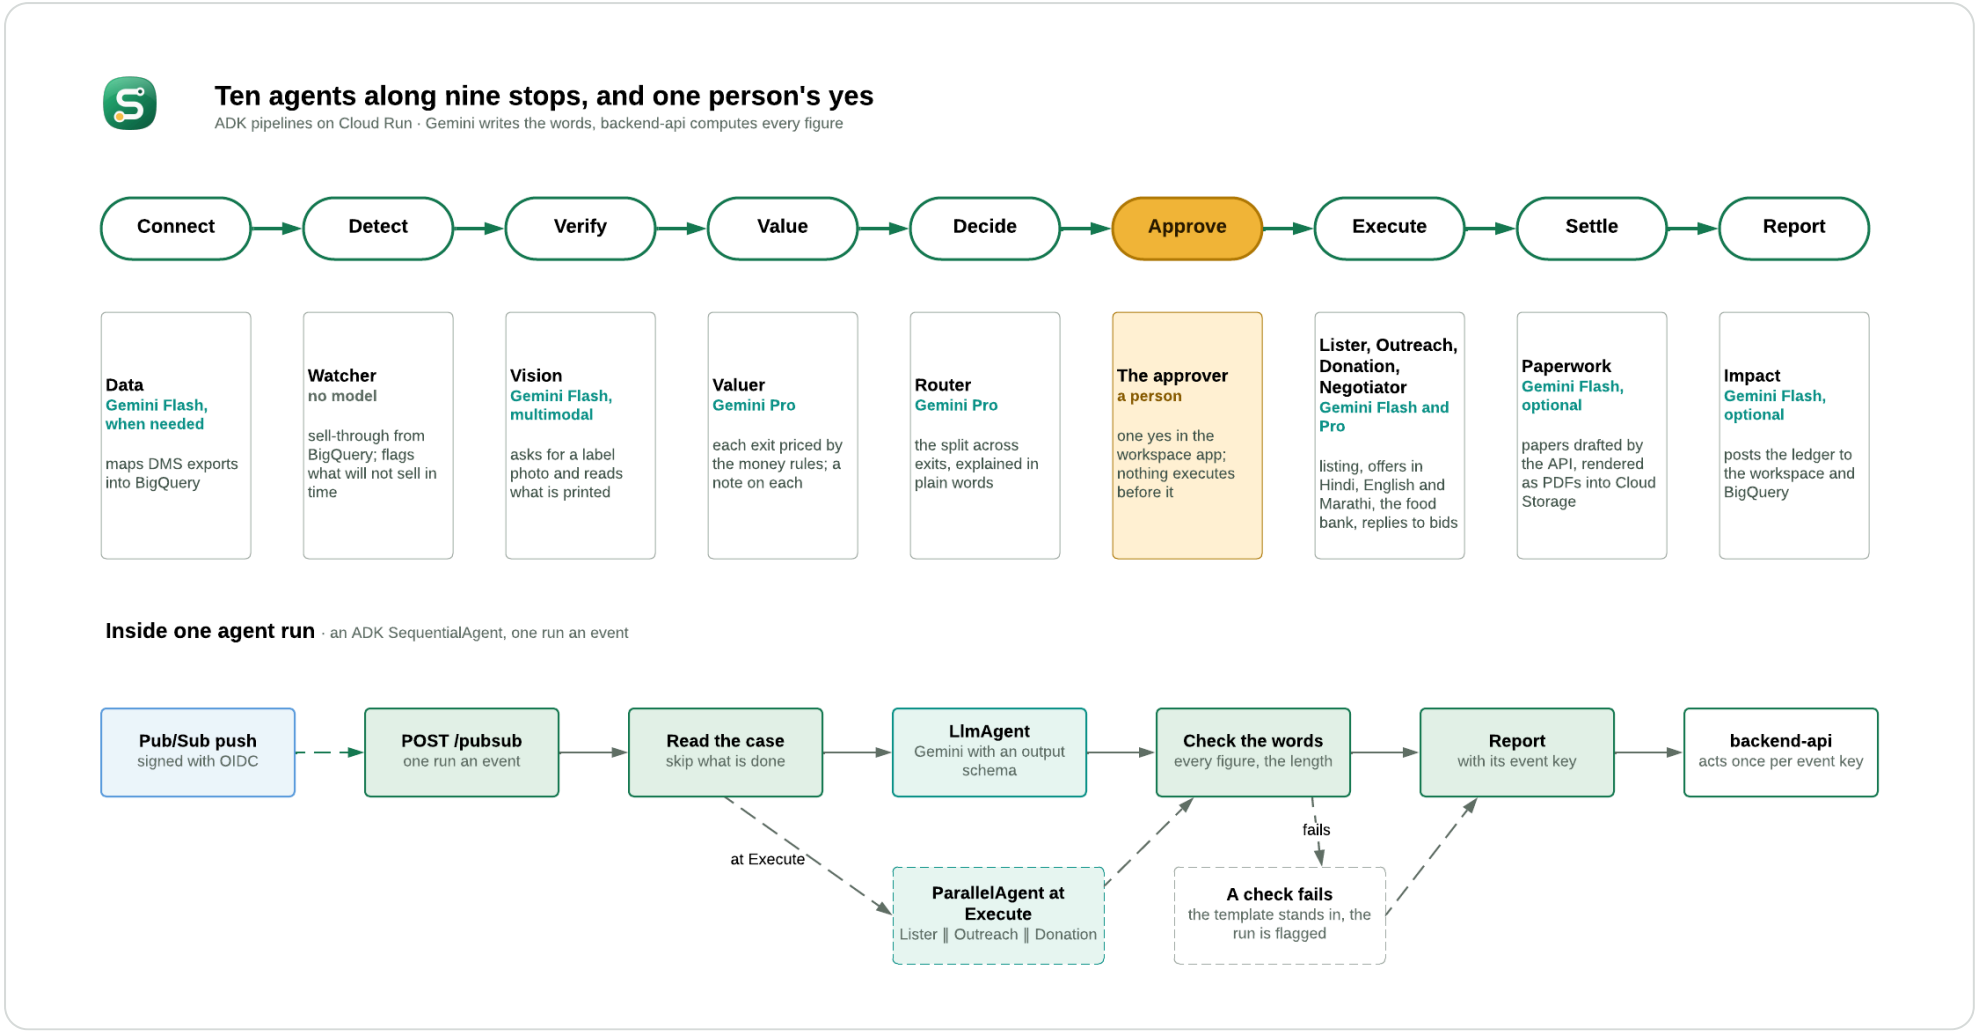
The destruction at the godown as a sequence: the report's timer or Report now, the distributor's two photos to Cloud Storage and four signatures, Vision's read on Gemini, Priya's review, and on her yes the papers and the ledger; else she asks again



Lucid: Smart-Clearance · Destroyed at the godown, against evidence and a second yes · https://lucid.app/lucidchart/08863960-0759-4e8b-8e46-69137532ad9a/view

A.7 Ten agents along nine stops

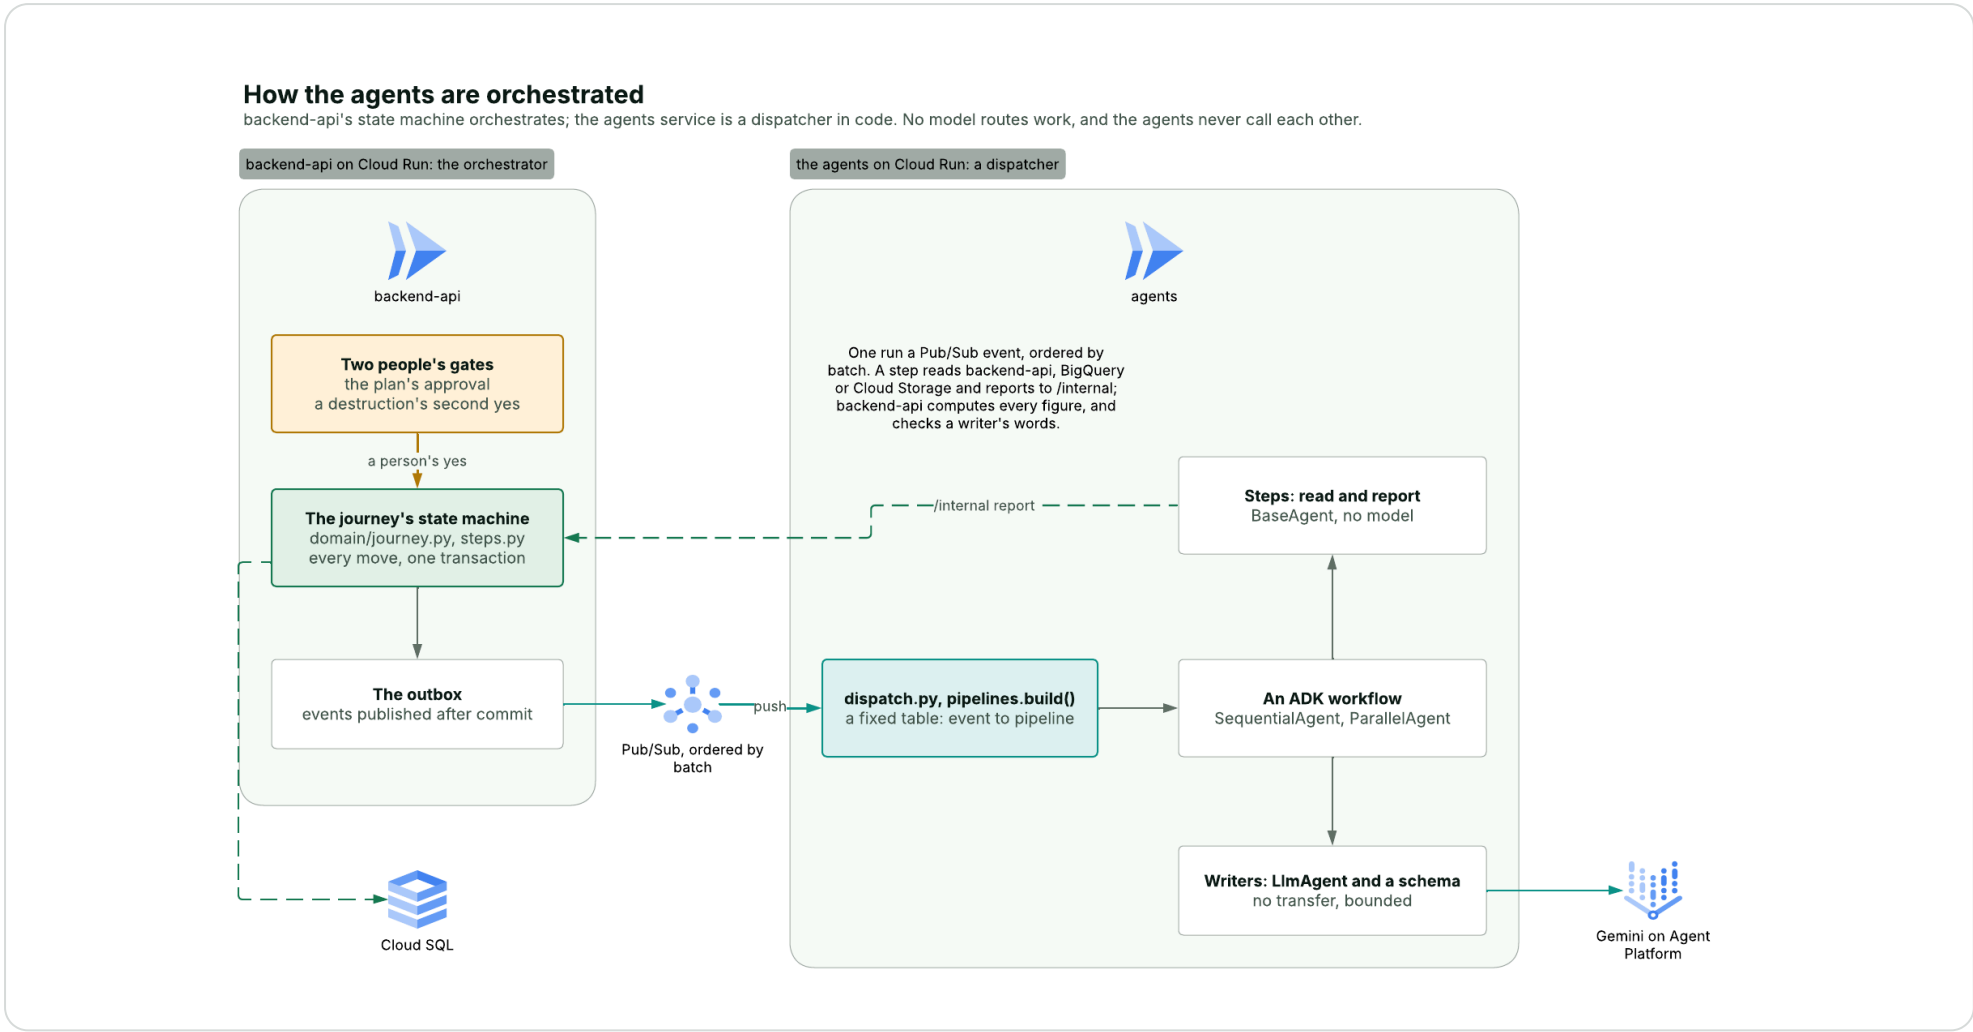
Ten agents along nine stops, and one person's yes; inside one agent run, an ADK SequentialAgent from the Pub/Sub push to the report



Lucid: Smart-Clearance Architecture on Google Cloud, page 2 · https://lucid.app/lucidchart/0e33214e-b46f-4cfb-9413-b2c40de7d1c6/view

A.8 How the agents are orchestrated

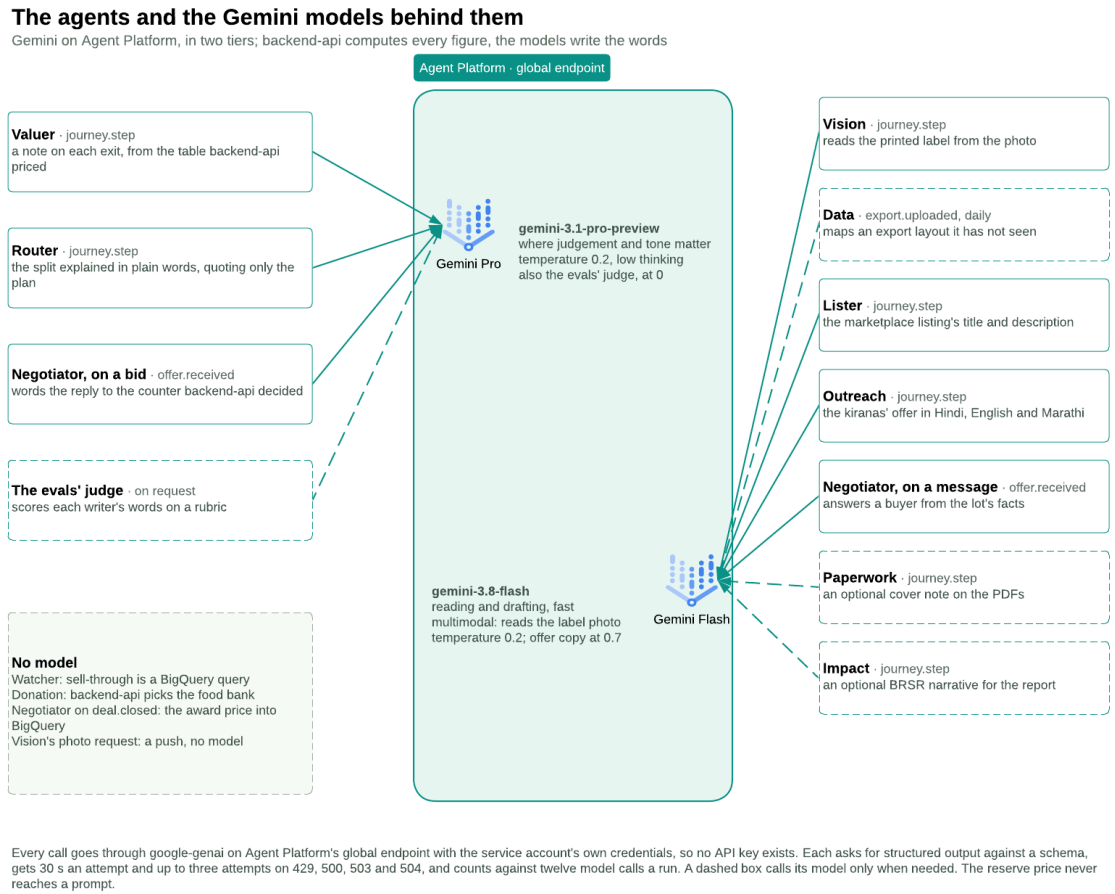
backend-api on Cloud Run orchestrates: the state machine, two people's gates and the outbox; Pub/Sub pushes each event to the agents on Cloud Run, where dispatch.py picks an ADK workflow of steps and bounded writers, the writers calling Gemini on Agent Platform; steps report to /internal, and the record lives in Cloud SQL



Lucid: Smart-Clearance · How the agents are orchestrated · https://lucid.app/lucidchart/78ca2917-3019-4930-b162-7c8f6671bf68/view

A.9 The agents and their models

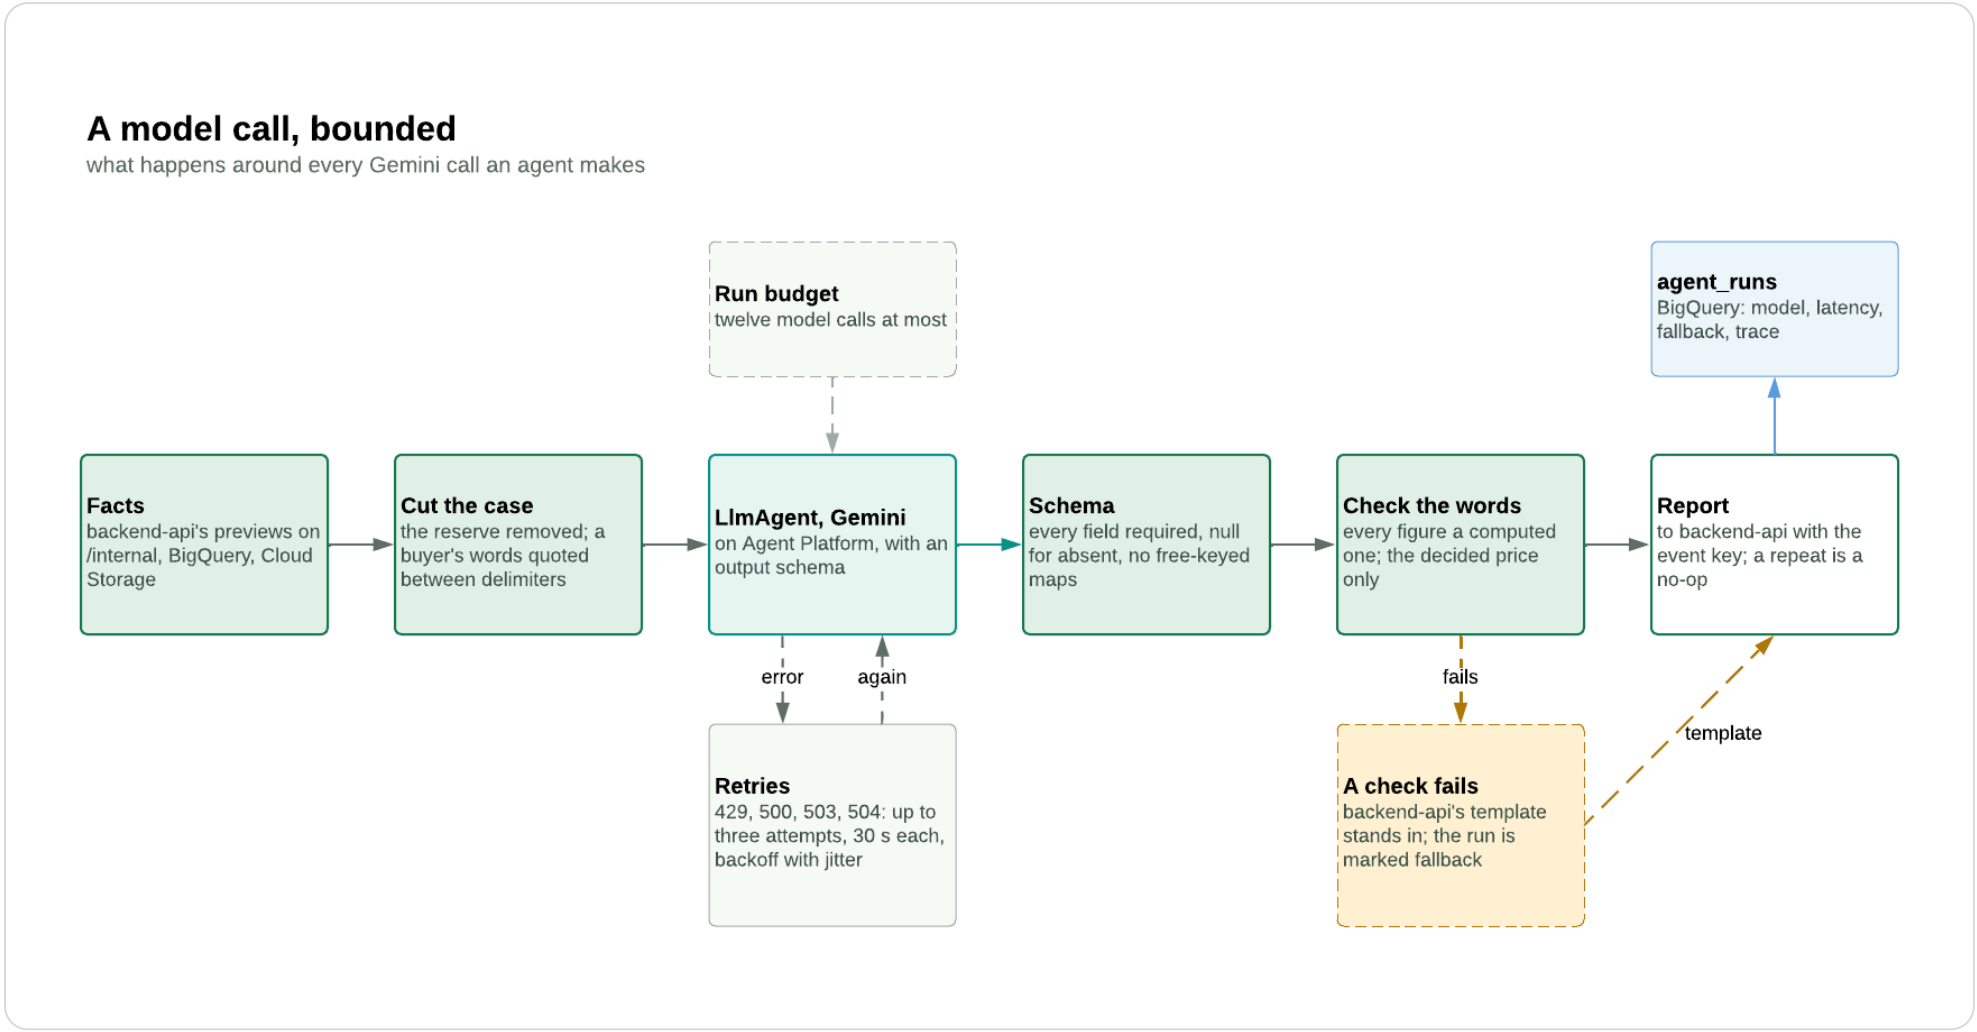
The agents and the Gemini models behind them: Pro for the Valuer, the Router, the Negotiator's bids and the evals' judge; Flash for Vision, Data, the Lister, Outreach, the Negotiator's messages, Paperwork and Impact; no model for the Watcher, Donation and the award



Lucid: Smart-Clearance Google Cloud and AI, page 6 · https://lucid.app/lucidchart/631bbb66-9b43-4e20-abf0-f44fa6a68174/view

A.10 A model call, bounded

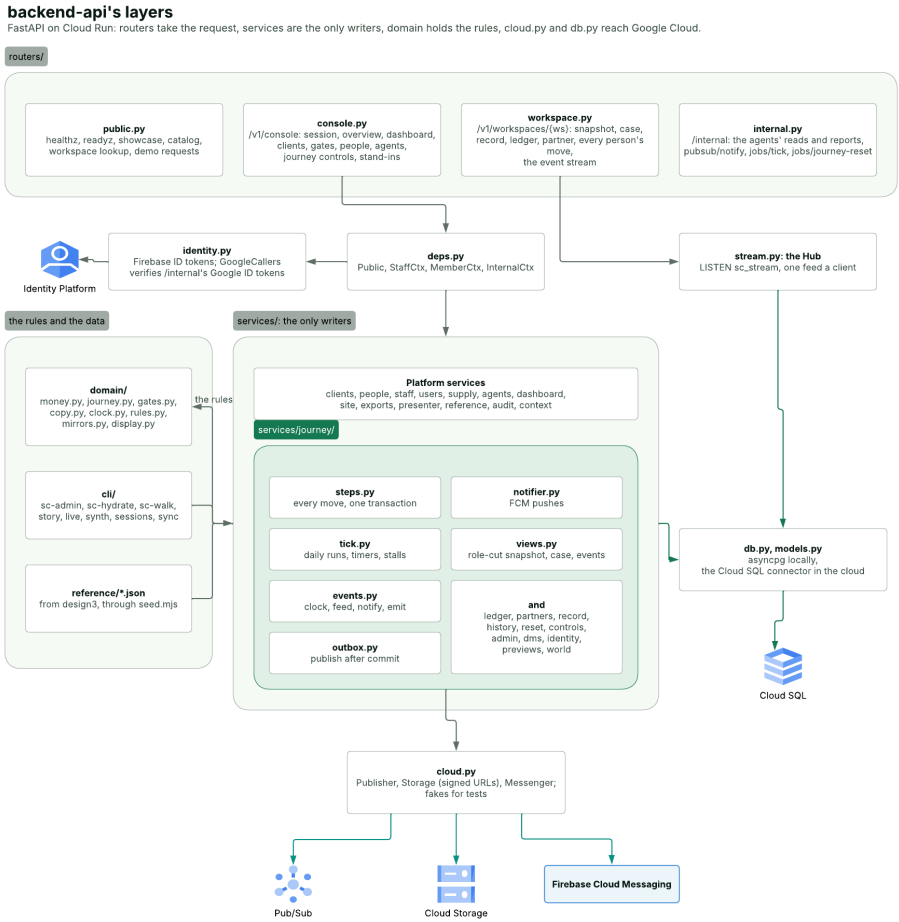
A model call, bounded: facts, the case cut, the LlmAgent with its schema, retries, the run budget, the checks, the template that stands in, and the report



Lucid: Smart-Clearance Google Cloud and AI, page 7 · https://lucid.app/lucidchart/631bbb66-9b43-4e20-abf0-f44fa6a68174/view

A.11 backend-api's layers

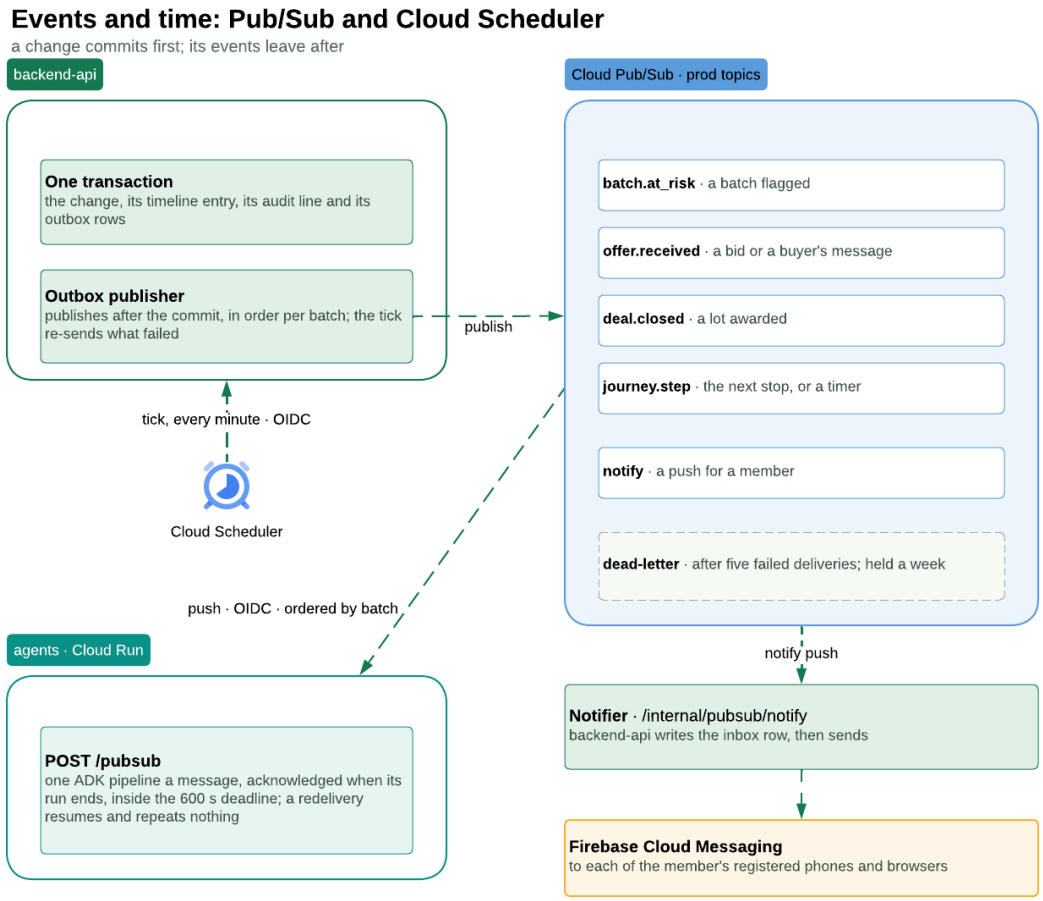
backend-api's layers: four routers, deps.py and identity.py (Identity Platform), stream.py, the services with services/journey/, domain, cli and the reference data, db.py to Cloud SQL, and cloud.py to Pub/Sub, Cloud Storage and Firebase Cloud Messaging



Lucid: Smart-Clearance · backend-api's layers · https://lucid.app/lucidchart/150c62da-d469-4561-b174-202f85d143fa/view

A.12 Pub/Sub and Cloud Scheduler

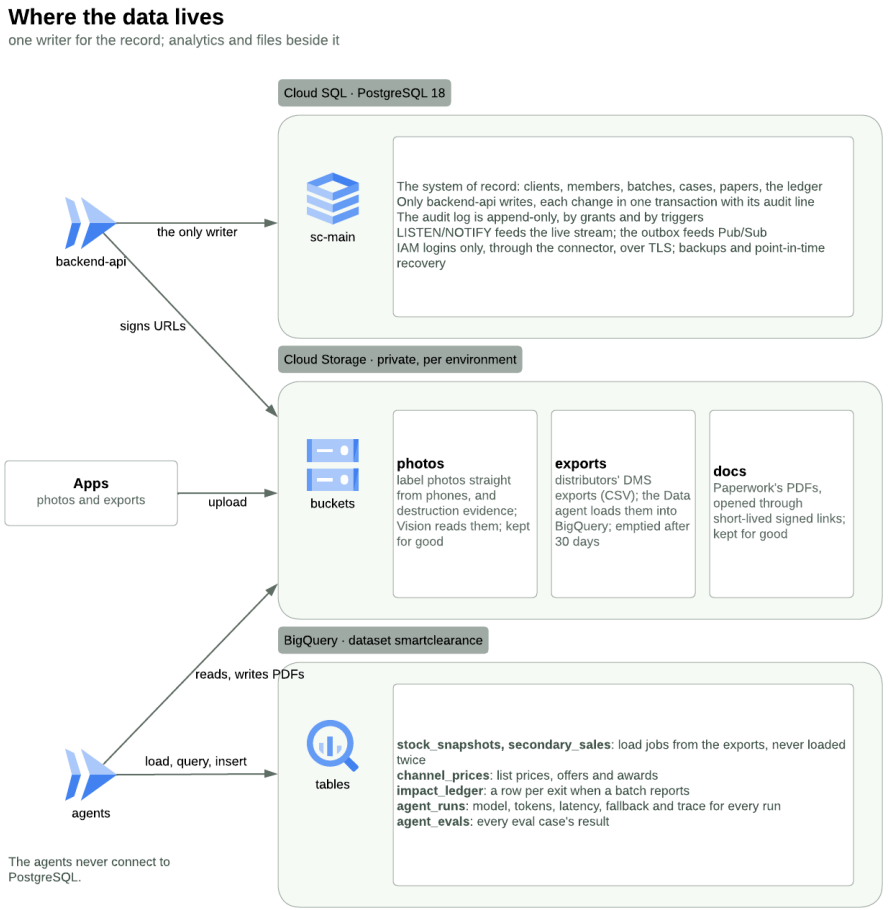
Events and time: one transaction writes the change and its outbox rows; the outbox publisher sends them after the commit; Cloud Scheduler ticks every minute; Pub/Sub pushes to the agents and to the Notifier; FCM reaches the member's devices



Lucid: Smart-Clearance Google Cloud and AI, page 3 · https://lucid.app/lucidchart/631bbb66-9b43-4e20-abf0-f44fa6a68174/view

A.13 Where the data lives

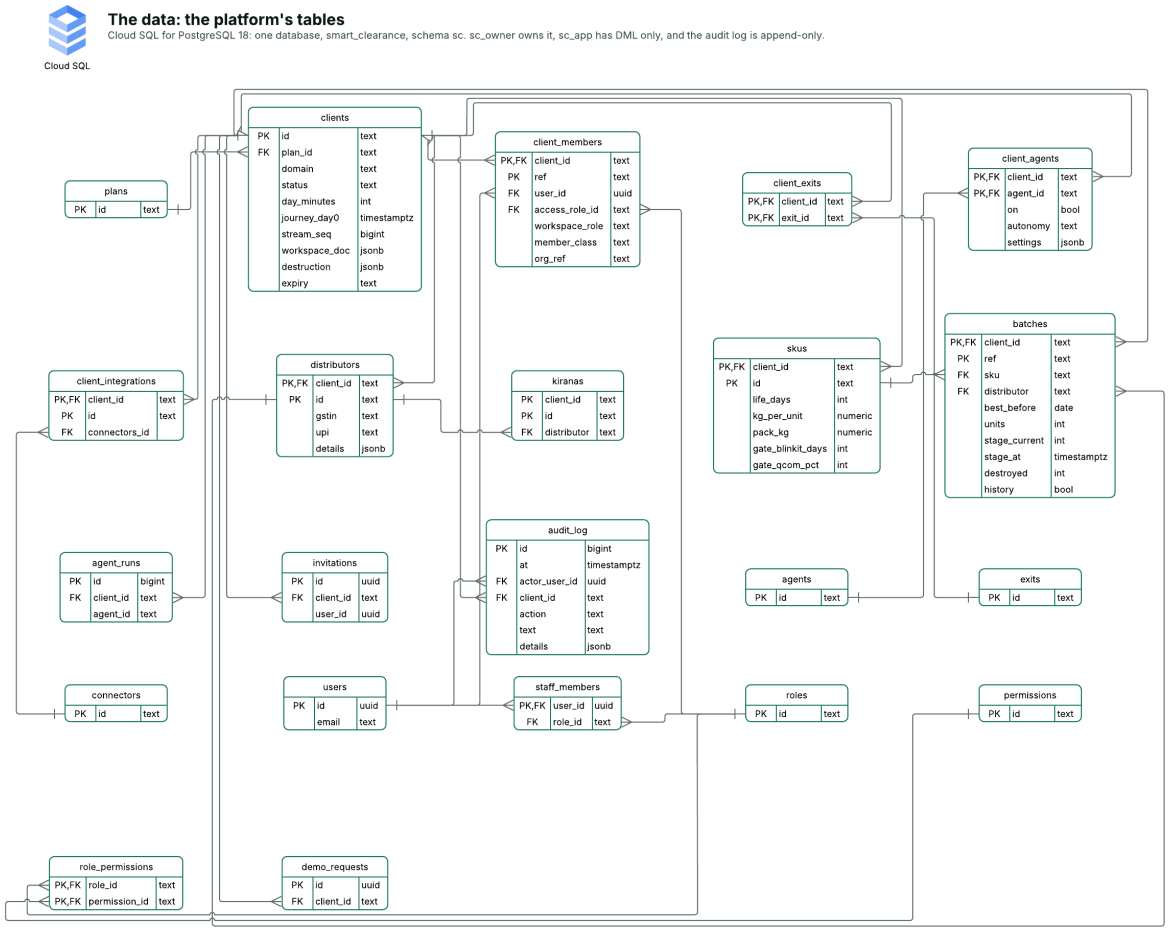
Where the data lives: Cloud SQL sc-main is the system of record, written only by backend-api; Cloud Storage holds photos, exports and docs; BigQuery holds the analytics the agents load, query and insert



Lucid: Smart-Clearance Google Cloud and AI, page 4 · https://lucid.app/lucidchart/631bbb66-9b43-4e20-abf0-f44fa6a68174/view

A.14 The data: the platform's tables

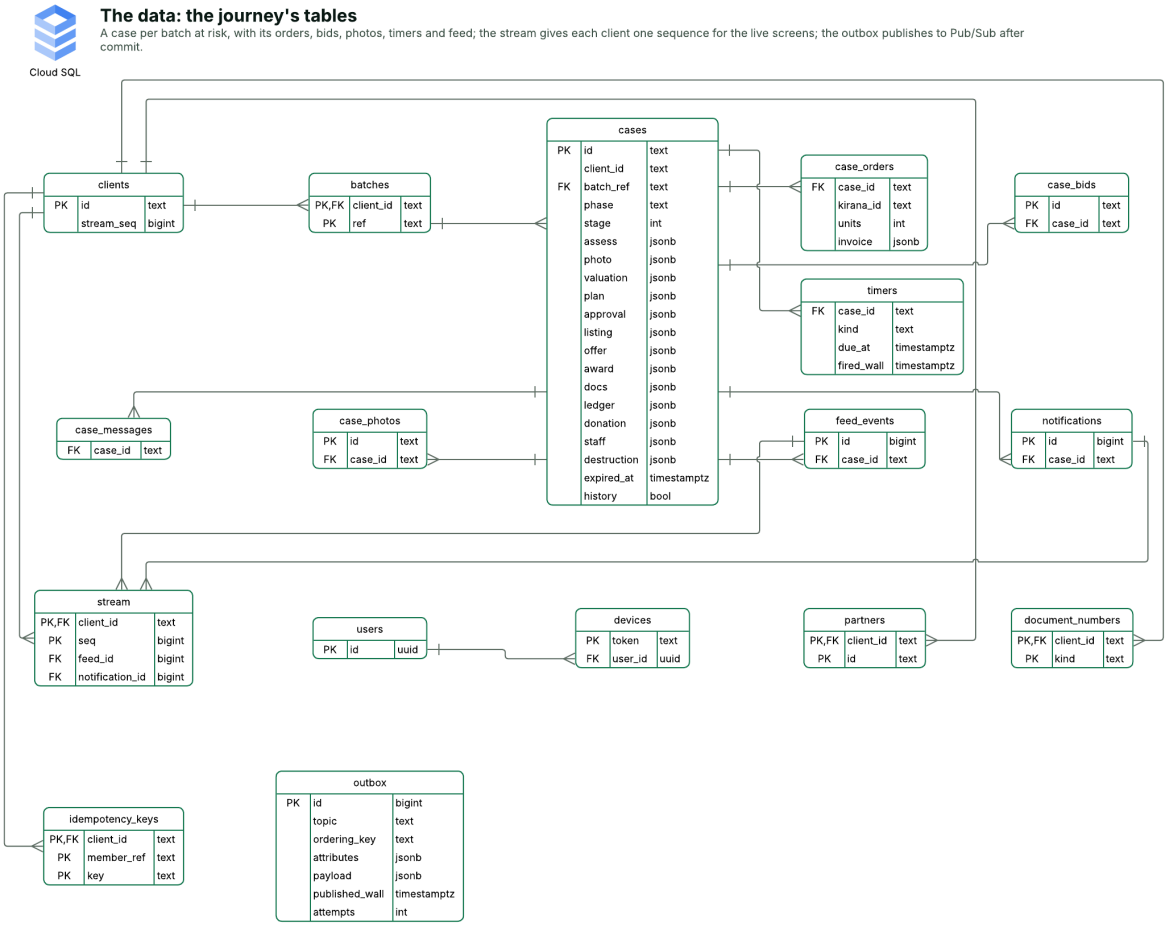
Entity relationship diagram of the platform's tables in Cloud SQL: plans, clients, members, exits, agents, integrations, distributors, kiranas, SKUs, batches, agent runs, invitations, the audit log, users, staff, roles and permissions, demo requests



Lucid: Smart-Clearance · The data: the platform's tables · https://lucid.app/lucidchart/13537dc0-591f-486a-a0d6-e326cf875e3a/view

A.15 The data: the journey's tables

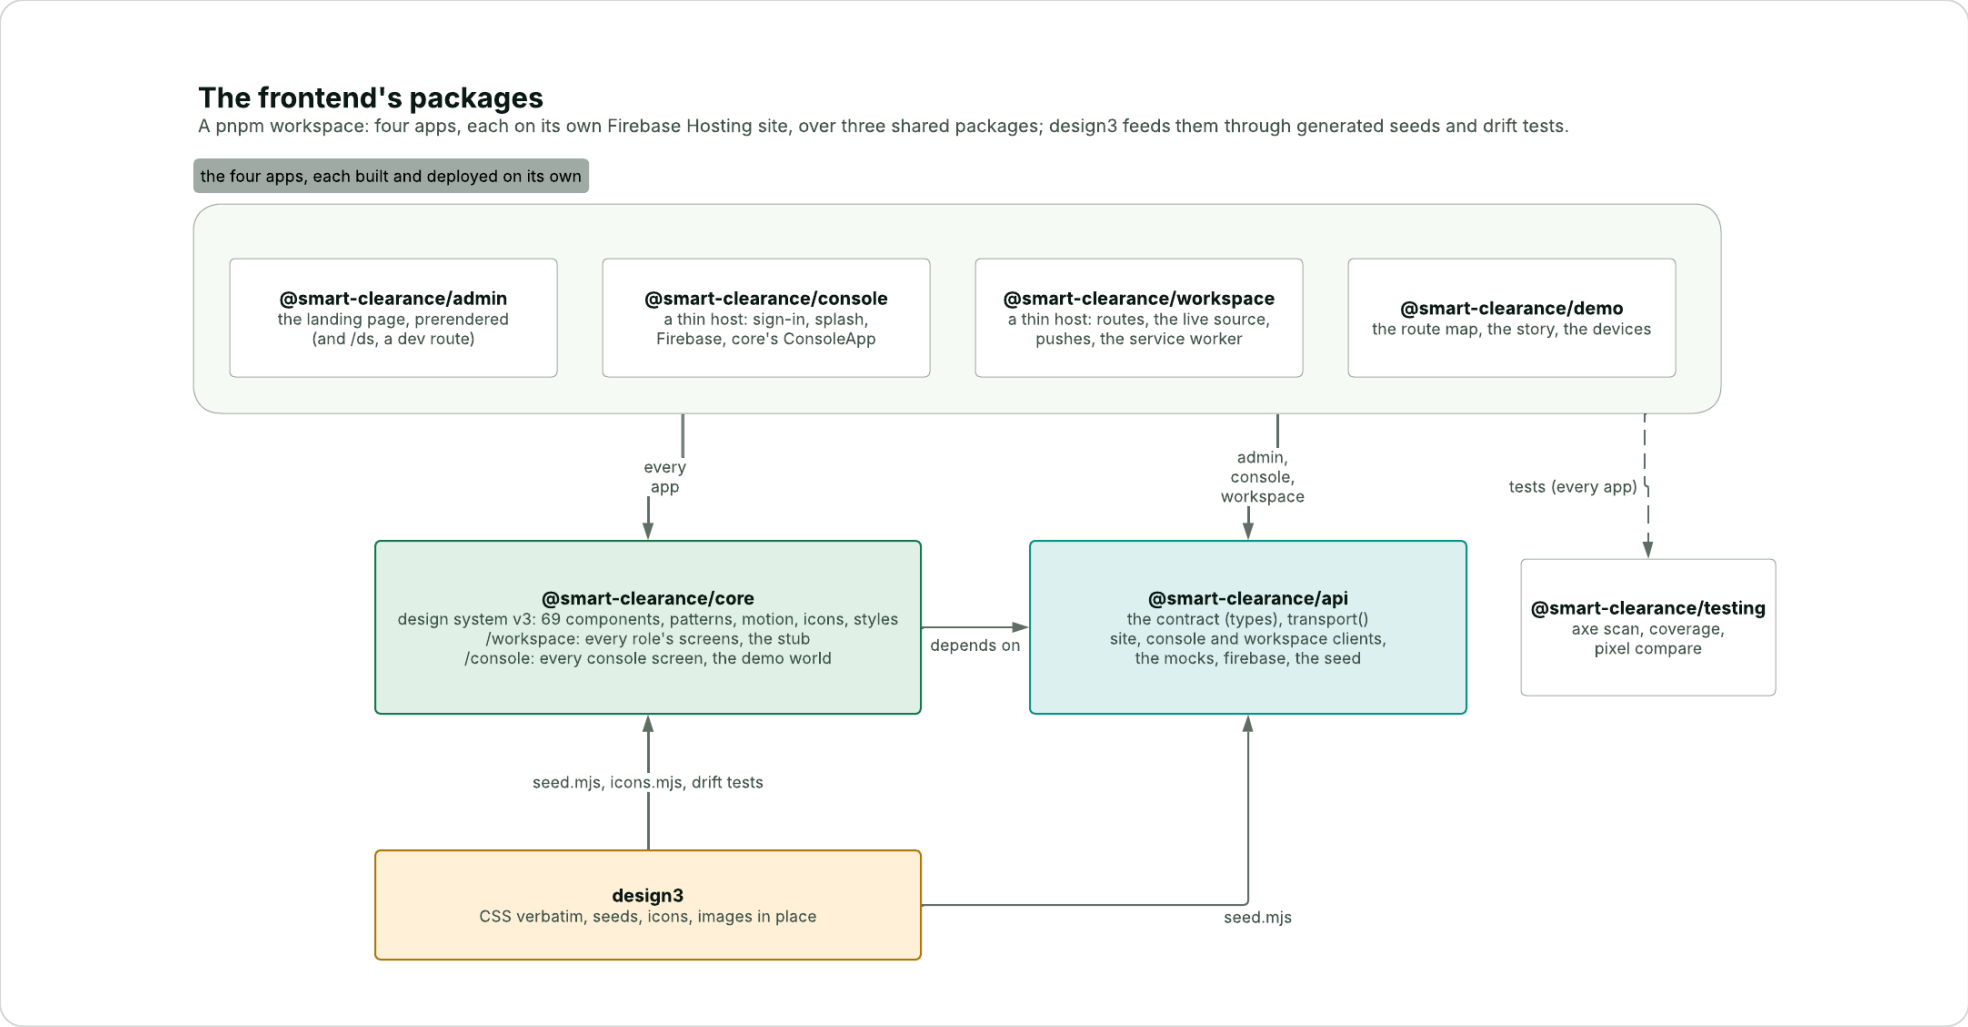
Entity relationship diagram of the journey's tables in Cloud SQL: cases with their orders, bids, messages, photos, timers, feed and notifications, the stream, devices, partners, document numbers, idempotency keys and the outbox



Lucid: Smart-Clearance · The data: the journey's tables · https://lucid.app/lucidchart/97c4c4bb-b174-4608-b5ce-25d4dd660ce4/view

A.16 The frontend's packages

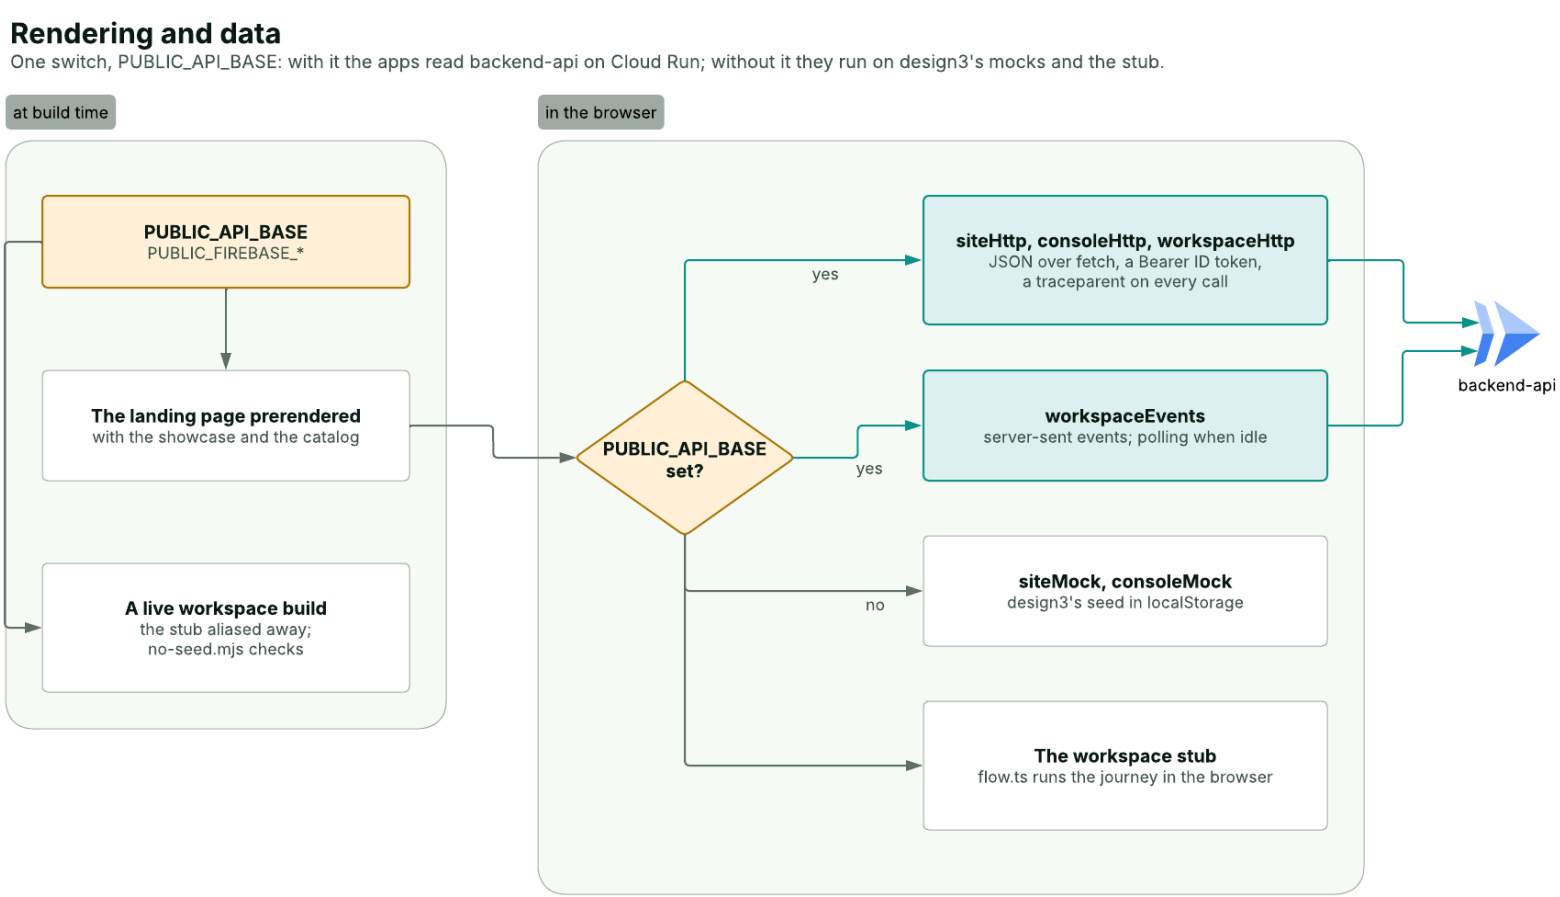
The frontend's packages: the four apps over core and api, core depending on api, testing for every app, and design3 feeding core and api through seed.mjs, icons.mjs and drift tests



Lucid: Smart-Clearance · The frontend's packages · https://lucid.app/lucidchart/dea5e30f-002c-430f-93b0-0c00b1f66895/view

A.17 Rendering and data

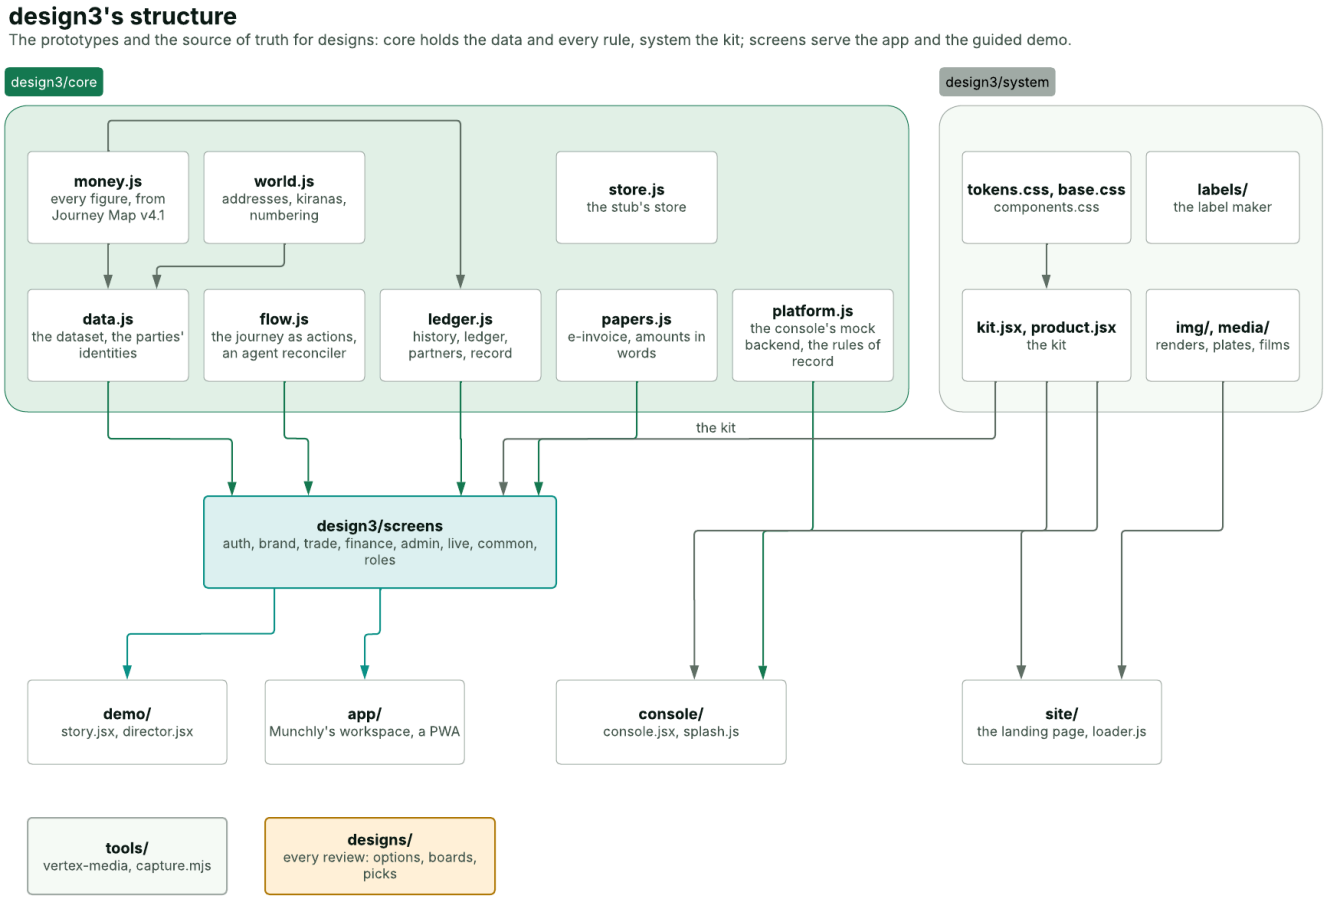
Rendering and data: PUBLIC_API_BASE at build time; in the browser, with it the HTTP clients and the event stream reach backend-api on Cloud Run, without it the mocks and the workspace stub run



Lucid: Smart-Clearance · Rendering and data · https://lucid.app/lucidchart/d7a7eca7-4e9c-4d7b-bf2f-7bbc4f9ea07c/view

A.18 design3's structure

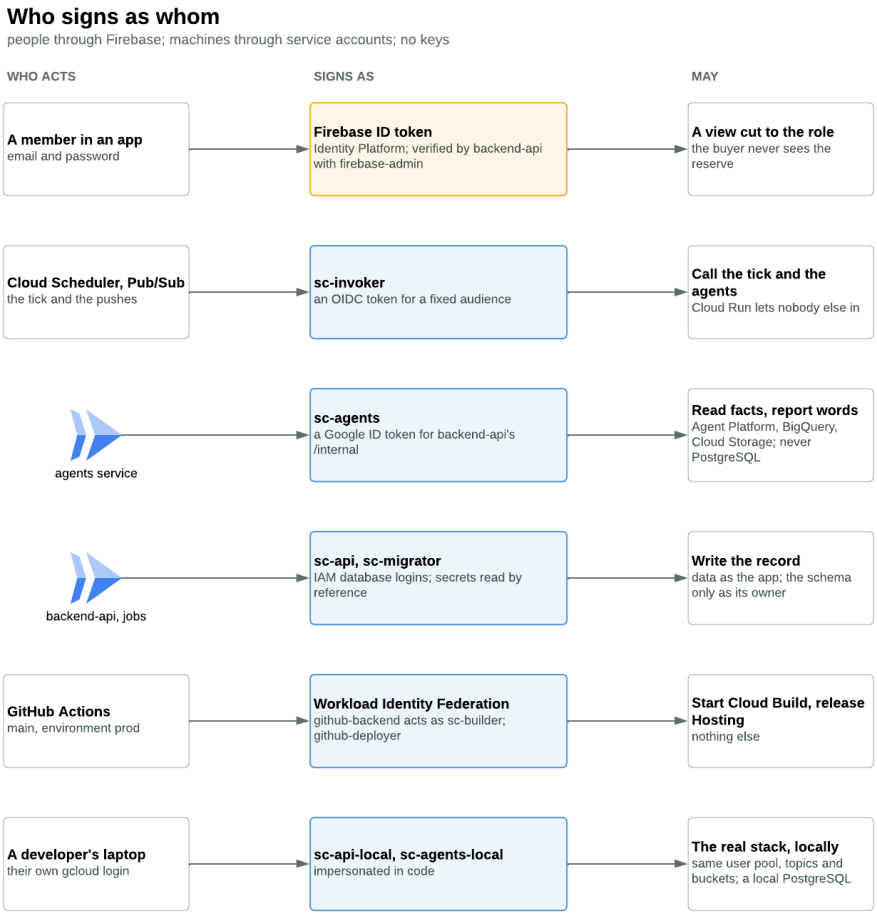
design3's structure: core (money, data, world, store, flow, ledger, papers, platform) and system (tokens, kit, labels, images) feed the screens, which serve the demo and the app; the kit also serves the console and the landing page



Lucid: Smart-Clearance · design3's structure · https://lucid.app/lucidchart/11c44a73-6506-48ab-adfb-98338cb9af35/view

A.19 Who signs as whom

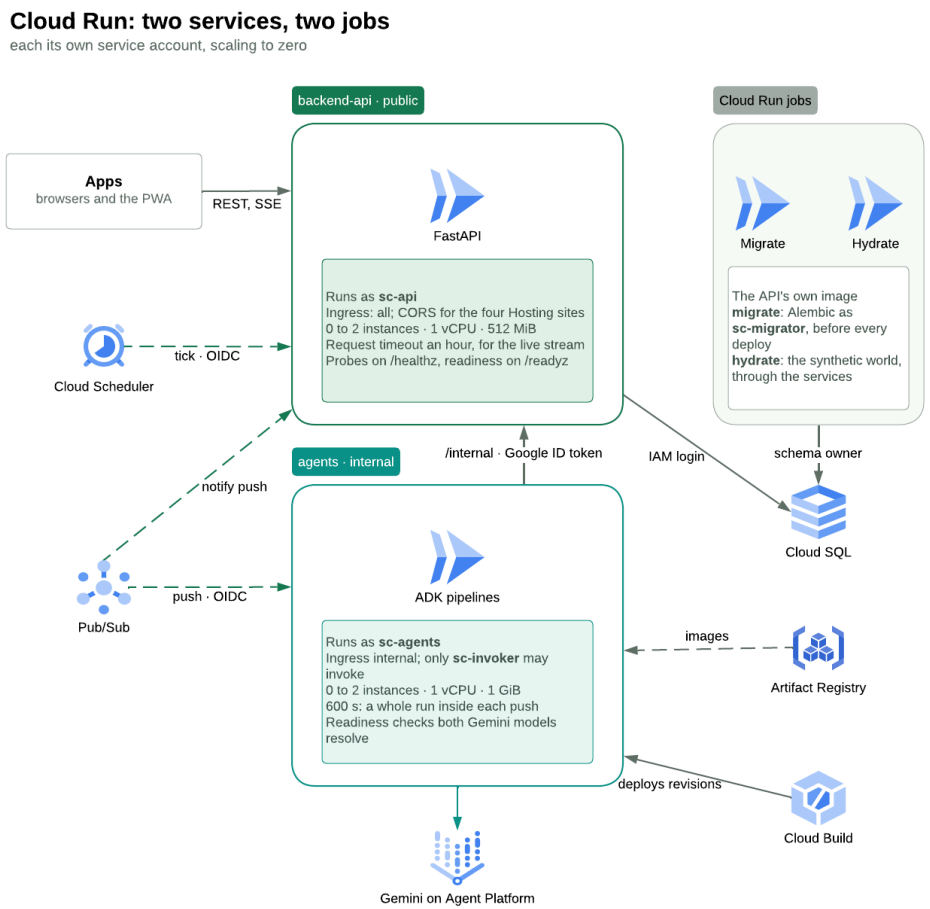
Who signs as whom: a member through a Firebase ID token; Cloud Scheduler and Pub/Sub as sc-invoker; the agents as sc-agents; backend-api and its jobs as sc-api and sc-migrator; GitHub Actions through Workload Identity Federation; a developer's laptop as sc-api-local and sc-agents-local



Lucid: Smart-Clearance Google Cloud and AI, page 5 · https://lucid.app/lucidchart/631bbb66-9b43-4e20-abf0-f44fa6a68174/view

A.20 Cloud Run: two services, two jobs

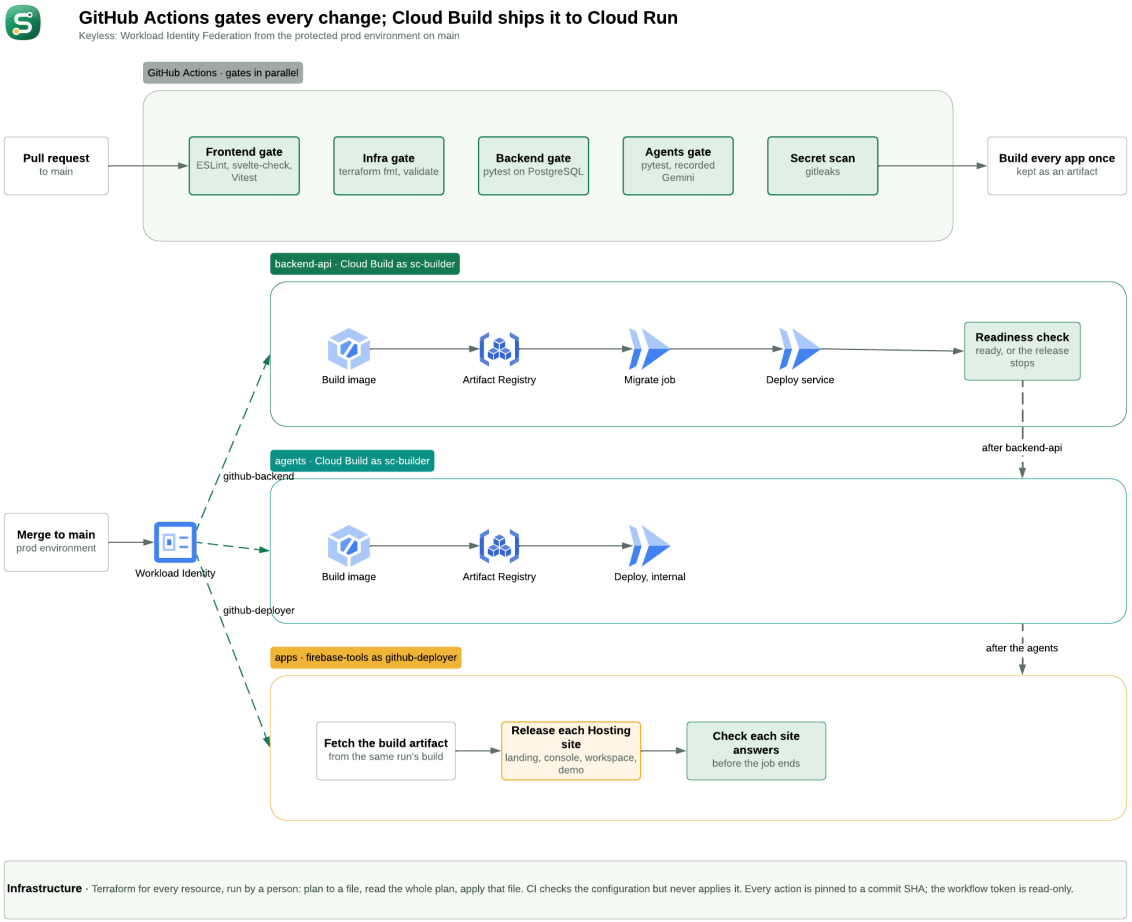
Cloud Run: two services, two jobs. backend-api is public as sc-api; the agents are internal as sc-agents; migrate and hydrate run the API's image; Cloud Scheduler and Pub/Sub push with OIDC



Lucid: Smart-Clearance Google Cloud and AI, page 2 · https://lucid.app/lucidchart/631bbb66-9b43-4e20-abf0-f44fa6a68174/view

A.21 GitHub Actions and Cloud Build

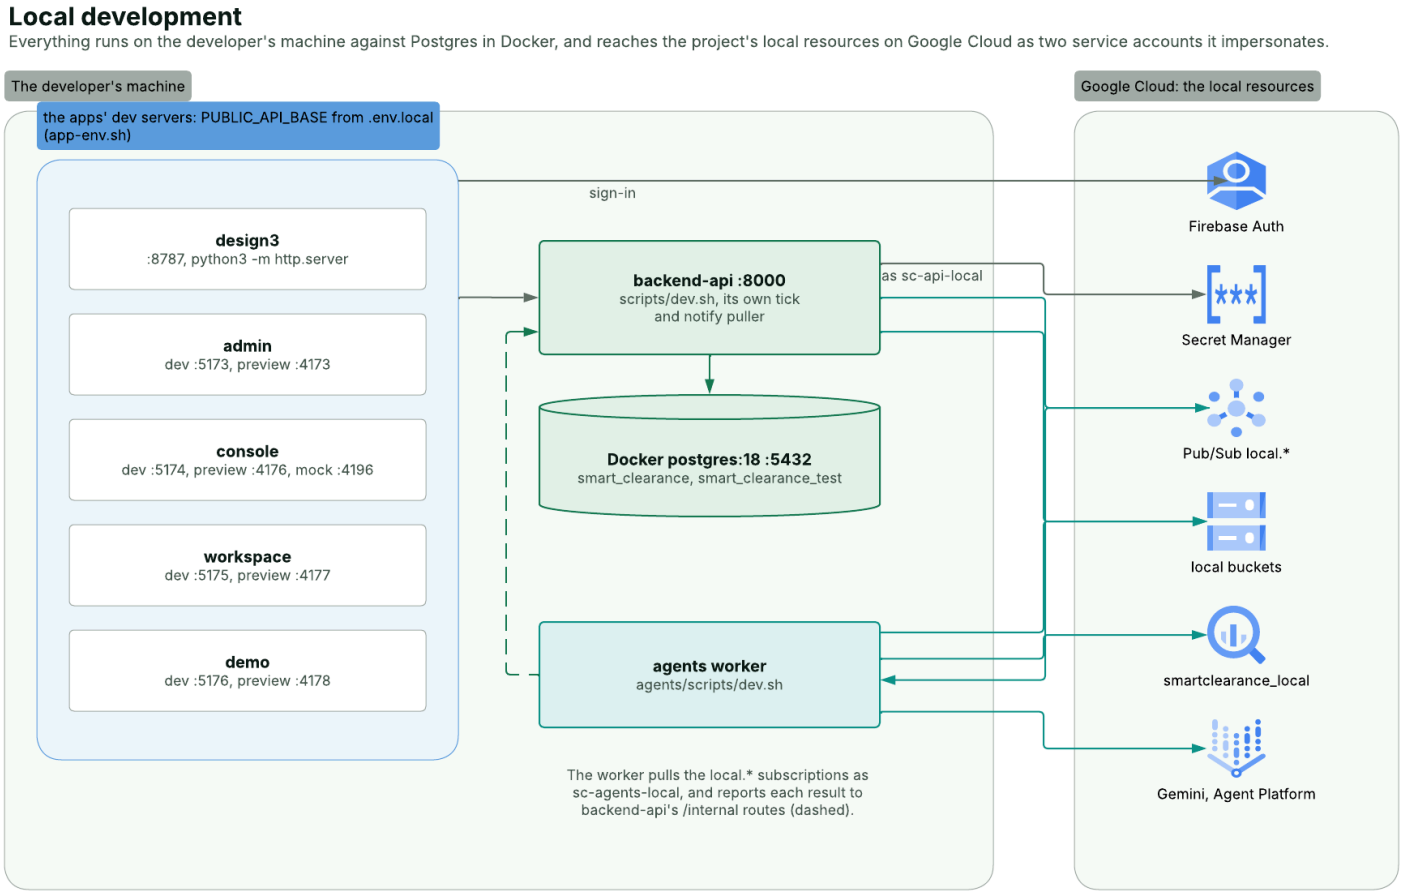
GitHub Actions gates every change; Cloud Build ships it to Cloud Run: five gates in parallel and one build on every pull request; on main, backend-api, then the agents, then the four Hosting sites, keyless



Lucid: Smart-Clearance Architecture on Google Cloud, page 3 · https://lucid.app/lucidchart/0e33214e-b46f-4cfb-9413-b2c40de7d1c6/view

A.22 Local development

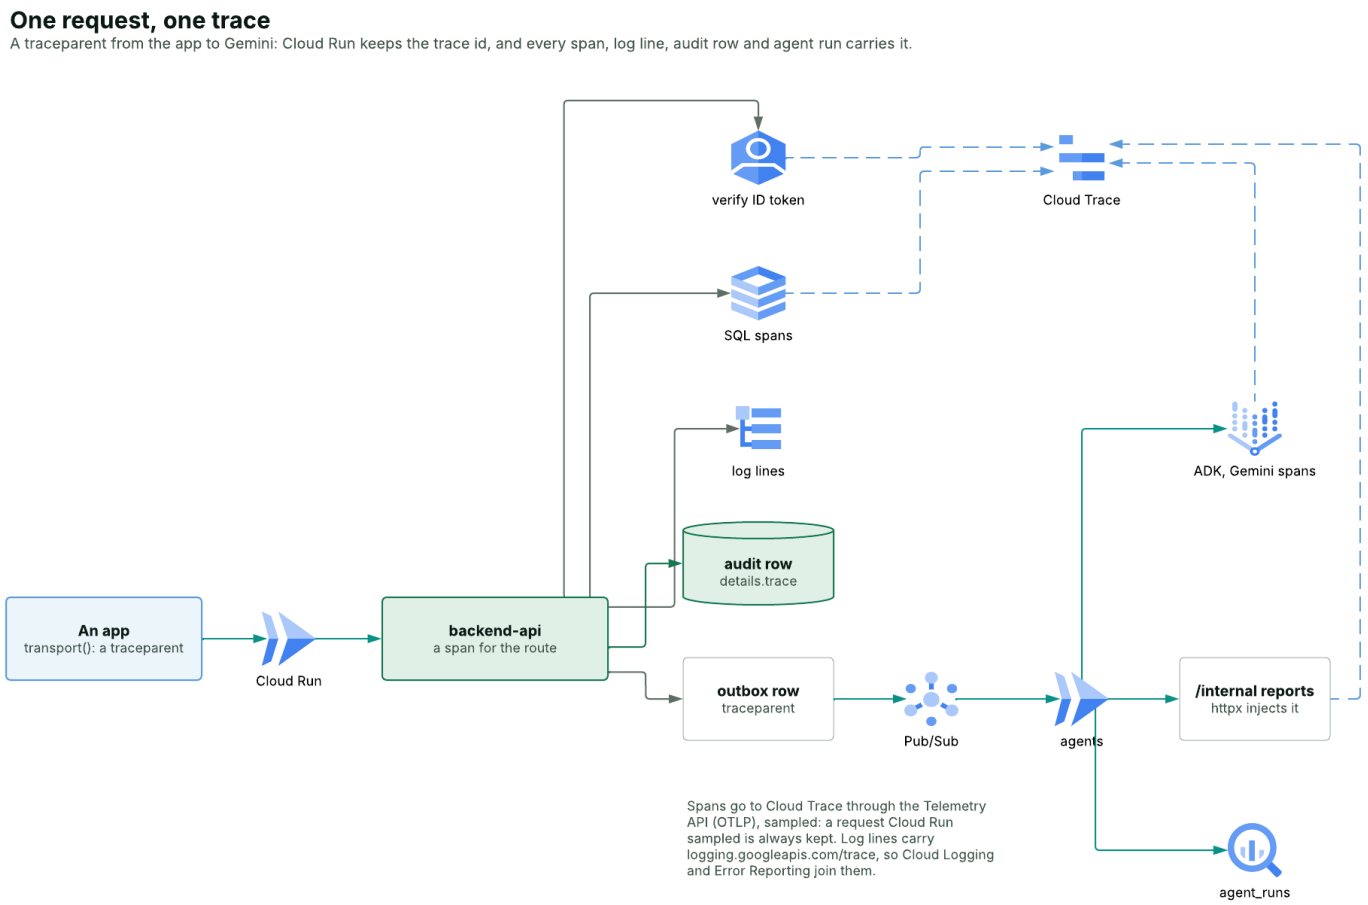
Local development: the apps' dev servers, backend-api on Postgres in Docker and the agents worker on the developer's machine, reaching Firebase Auth, Secret Manager, Pub/Sub, the local buckets, BigQuery and Gemini on Agent Platform



Lucid: Smart-Clearance · Local development · https://lucid.app/lucidchart/a4d19edf-22c8-4f09-9754-5e0510a9a242/view

A.23 One request, one trace

One request, one trace: an app's traceparent through Cloud Run into backend-api's spans, log lines, audit row and outbox, through Pub/Sub to the agents, their Gemini spans, /internal reports and agent_runs in BigQuery, all joined in Cloud Trace



Lucid: Smart-Clearance · One request, one trace · https://lucid.app/lucidchart/d33fb84b-b1ec-4745-bd55-8ed880e80e91/view

B APPENDIX B

Glossary

Agent Platform
Gemini Enterprise Agent Platform, Google's name since April 2026 for Vertex AI; Gemini is reached at aiplatform.googleapis.com.

ADK
Google's Agent Development Kit. Each agent pipeline is an ADK workflow of steps and bounded LlmAgent writers.

Batch
One lot of a SKU at a distributor's godown, with its best-before date; the unit the platform watches.

Case
A batch at risk worked through the journey: its phase, plan, approval, lines, papers and ledger.

Journey
The nine stops a case moves through: Connect, Detect, Verify, Value, Decide, Approve, Execute, Settle, Report.

Exit
One way out for packs at risk: the kiranas' scheme, the marketplace, a staff sale, a food bank, or destruction.

Distributor
The brand's stockist in a city, who holds the batch in a godown and serves the kiranas.

Godown
A warehouse, in Indian usage; where a distributor keeps stock.

Kirana
A neighbourhood grocery shop. The scheme offers it packs at a margin, in its own language.

ExpireSoon
The fictional near-expiry marketplace the Lister lists a lot on, and the Negotiator bargains with buyers on.

Staff sale
Packs sold to the distributor's staff, recorded on one consolidated receipt of UPI payments.

Destruction at the godown
Packs that sold nowhere, destroyed by an authorised agency on photo evidence, a signed certificate and a second yes.

Expiry day
Best-before, when open lines close as they stand and the report falls due; Report now brings it forward.

Price support
Credit from the brand to the distributor for the gap between the dealer price and what the exits fetched.

GST ITC
Input tax credit under India's GST; reversed under s.17(5)(h) for goods destroyed or given away.

IRN
Invoice reference number of an e-invoice; worked out here as the portal does, with an illustrative signed QR.

GSTIN, PAN, FSSAI
India's tax, income-tax and food-safety identifiers. Every one in Smart-Clearance is fictional.

BRSR
SEBI's Business Responsibility and Sustainability Report; Principle 6 carries food waste and packaging.

Prototype time
The journey's own clock: the story's calendar from day 0, at a client's chosen day length.

Outbox
Rows written in a step's transaction and published to Pub/Sub after it commits, so no event is lost.

SSE, FCM
Server-sent events keep open screens live; Firebase Cloud Messaging pushes to closed ones.

Workspace
A client's own app at <client>-smartclearance.web.app; the console is Smart-Clearance's staff app.

PWA
Progressive web app: the workspace installs on a laptop, a tablet or a phone, and takes pushes.

Stand-in
A console control that acts for a demo world's partners, such as filling a kirana scheme.

Smart-Clearance

Every near-expiry carton gets a second chance.

Source · github.com/aibuilder2026/smart-clearance

Landing page · smartclearance.web.app

Guided demo · smartclearance-demo.web.app

Workspace · munchly-smartclearance.web.app

Console · smartclearance-console.web.app

SC-TDD-001 · version 1.0 · 15 October 2026. Every company, person and figure in Smart-Clearance is fictional.

